

XBOOT: Free-XOR Gates for CKKS with Applications to Transciphering

Chao Niu , Zhicong Huang , Zhaomin Yang , Yi Chen , Liang Kong ,
Cheng Hong*  and Tao Wei 

Ant Group

{niuchao.niu,zhicong.hzc,yangzhaomin.ym,cy459642,kongliang.kong,vince.hc,lenx.wei}@antgroup.com

Abstract. The CKKS scheme is traditionally recognized for approximate homomorphic encryption of real numbers, but BLEACH (Drucker et al., JoC 2024) extends its capabilities to handle exact computations on binary or small integer numbers. Despite this advancement, BLEACH’s approach of simulating XOR gates via $(a - b)^2$ incurs one multiplication per gate, which is computationally expensive in homomorphic encryption. To this end, we introduce XBOOT, a new framework built upon BLEACH’s blueprint but allows for almost free evaluation of XOR gates. The core concept of XBOOT involves lazy reduction, where XOR operations are simulated with the less costly addition operation, $a + b$, leaving the management of potential overflows to later stages. We carefully handle the modulus chain and scale factors to ensure that the overflows are managed during the CKKS bootstrapping phase, preserving the correct XOR result without extra cost. We use AES-CKKS transciphering as a benchmark to test the capability of XBOOT, and achieve a throughput exceeding one kilobyte per second, which represents a $2.5\times$ improvement over the state-of-the-art (Aharoni et al., HES 2023). Moreover, XBOOT enables the practical execution of tasks with extensive XOR operations that were previously challenging for CKKS. For example, we can do Rasta-CKKS transciphering at over two kilobytes per second, more than $10\times$ faster than the baseline without XBOOT.

Keywords: homomorphic encryption · bootstrapping · advanced encryption standard · hybrid homomorphic encryption · transciphering

1 Introduction

Fully Homomorphic Encryption (FHE) [Gen09] represents a groundbreaking paradigm in cryptography, allowing computations to be performed directly on encrypted data without decryption. Initially regarded as purely theoretical and impractical, FHE has seen significant advancements in recent years, leading to the proposal of numerous modern FHE schemes [BV11, FV12, GSW13, CCK⁺13, BGV14, BV14, DM15, CKKS17, CGGI17, CGGI20] with markedly improved efficiency. Among the various FHE schemes, the Cheon-Kim-Kim-Song (CKKS) scheme [CKKS17] stands out as one of the most efficient, and has been extensively utilized in works such as privacy-preserving machine learning [LLL⁺22, LLKN23, JPK⁺24].

However, CKKS is inherently an approximate homomorphic encryption scheme; it introduces noise in the computation, which accumulates with each operation, potentially leading to loss of precision over multiple computational layers. As a result, it was believed that CKKS could only be used for applications where exact precision is not always

*Corresponding author: Cheng Hong (vince.hc@antgroup.com)

necessary, and is unsuitable for applications sensitive to even minor variations. To mitigate the limitations posed by CKKS’s approximation, the BLEACH methodology [DMPS24] was introduced. BLEACH provides an efficient CKKS framework for bit operations by representing Boolean values with real numbers close to 0 or 1 (e.g., 0.002 or 0.997). It implements an AND gate via multiplication $a \cdot b$, and an XOR gate via $(a - b)^2$. It employs a “cleanup” function $h_1(x) = 3x^2 - 2x^3$ to reduce accumulated errors, enabling CKKS to maintain a high accuracy for Boolean circuits. This innovation extends the usage of CKKS into areas previously dominated by exact FHE schemes like BFV/BGV or TFHE.

One significant application in these areas is **transciphering**. In this methodology, clients can utilize inexpensive symmetric encryption (such as AES) instead of FHE, transmitting the symmetric ciphertext to the server, who then converts this ciphertext into an FHE ciphertext to enable subsequent computations. Transciphering significantly lowers both the computational and communication demands on the client, thereby offering considerable advantages for resource-limited devices like embedded systems. The core principle of transciphering involves the homomorphic evaluation of the symmetric decryption function, which is generally treated as a Boolean circuit. Due to the exceptional efficiency of the CKKS scheme, the throughput of AES-CKKS transciphering [ADE⁺23] using BLEACH is often higher compared to AES transciphering using other FHE schemes [GHS12, SMK22, TCBS23, WLW⁺24, BPR24, SAP24].

Despite the advancements, a significant drawback in BLEACH is the cost associated with the XOR operation. Specifically, since CKKS only supports addition and multiplication, BLEACH has to implement the XOR gate using the formula $(a - b)^2$, which requires one multiplication operation per XOR gate. Given that ciphertext-ciphertext multiplications in FHE are expensive, the number of XOR gates can lead to performance bottlenecks.

What’s worse, cryptography researchers often prioritize optimizing algorithms by minimizing non-XOR gates, and pay less attention to XOR gates because they are cheap in secure multi-party computation (MPC) or exact FHE schemes like BGV. For example, the optimized AES circuit in [KSS12] comprises 9,100 non-XOR gates and 21,628 XOR gates. Similarly, the complexity of LowMC [ARS⁺15] and Rasta [DEG⁺18] ciphers is estimated solely based on AND gates without consideration of XOR gates. However, in many transciphering schemes [NWH⁺25], XOR gates are not free of cost and must be considered in the overall computational burden. This discrepancy underscores the need for efficient XOR handling in CKKS, prompting us to pose the following question:

Can we homomorphically evaluate Boolean circuits in a way that leverages the high efficiency of CKKS while reducing the cost of XOR gates?

1.1 Our Contributions

In this work, we propose XBOOT, a new CKKS framework for computation on binary values. XBOOT addresses the above challenges by enabling almost “free” XOR gates. Our contributions are as follows:

Free-XOR framework for CKKS. Building upon the work of BLEACH, we introduce a lazy reduction technique where XOR gates are substituted with homomorphic additions. This method allows intermediate results to temporarily exceed their intended range. Each result consists of two parts: the Least Significant Bit (LSB) part, representing the actual XOR result, and a redundant part. Then we carefully design the CKKS modulus chain and scale factors so that the redundant part is effectively removed at the SlotsToCoeffs step of CKKS bootstrapping process. Thus, the LSB is automatically extracted, achieving the XOR functionality with minimal additional cost.

Highly efficient transciphering. We use transciphering as a representative application to demonstrate the efficiency of XBOOT. The state-of-the-art on AES-CKKS transciphering [ADE⁺23] consumes 9 multiplication depths per round, while XBOOT requires only 3 multiplication depths per round. With the reduced multiplication depth, and combining the optimizations from CKKS bootstrapping for bits [BCKS24], the bit length of CKKS ciphertext modulus can be efficiently reduced to 830 bits from 1555 bits, and the cyclotomic ring degree N can be reduced to 2^{15} from 2^{16} while still maintaining 128-bit security. Experiments show that our AES transciphering achieves a throughput of 256KB per 236 seconds, reducing latency by a factor of 5 and increasing throughput by a factor of 2.5 compared to the state-of-the-art [ADE⁺23].

Furthermore, we highlight that numerous applications involve an extensive number of XOR gates, rendering them impractical for evaluation under previous CKKS frameworks. XBOOT address this limitation, unlocking new possibilities for such tasks. For instance, the Rasta cipher [DEG⁺18], which incorporates nearly a million XOR gates, serves as a compelling example. By applying XBOOT to Rasta-CKKS transciphering, we achieve a remarkable throughput of over two kilobytes per second. This outperforms the other Rasta transciphering methods by more than one order of magnitude. We open-sourced the experiment code in <https://github.com/AntCPLab/xboot-ckks-transciphering>.

1.2 Related Work

CKKS bootstrapping. Bootstrapping for CKKS was first introduced in [CHK⁺18]. Since then, continuous advancements have been made. Notable progress includes optimizations in homomorphic modular reduction [CHK⁺18, LLL⁺21, JM22, LLK⁺22], as well as complexity reductions for linear transformations, with FFT-like decomposition applied in [CCS19, HK20, BMTH21]. More recently, [BCKS24, BKSS24] proposed optimizations for bootstrapping bits or small integers.

Transciphering. Since Gentry et al. introduced the first homomorphic computation for AES, transciphering for AES serves as a benchmark for evaluating the efficiency of FHE.

Besides CKKS, the TFHE scheme, especially with functional bootstrapping [CGGI20], has been extensively studied for its application in transciphering. Stracovsky et al. [SMK22] evaluated an AES block in 4 minutes using 16 threads. More recently, Trama et al. [TCBS23] proposed a faster AES evaluation using functional and multi-value bootstrapping techniques. Recent works, such as [WWL⁺23, WWL⁺24], have reported even better results based on the rich features of the TFHE scheme. [CLW⁺24] optimized the computation of XOR gates in TFHE, and performed transciphering experiments on AES and ASCON.

In general, TFHE transciphering provides lower latency than CKKS transciphering but has worse throughput due to its lack of SIMD capabilities.

2 Preliminaries

In this section, we outline the essential background knowledge of FHE and transciphering and establish the notations used throughout the paper. For a power-of-two integer N , we define the N -degree polynomial ring as $\mathcal{R}_N = \mathbb{Z}[X]/(X^N + 1)$ and the residual ring $\mathcal{R}_{Q,N} = \mathcal{R}_N/Q\mathcal{R}_N$, which is formed by taking $\mathcal{R}_N$ modulo an integer Q .

2.1 Approximate Homomorphic Encryption (CKKS)

We revisit the CKKS scheme, which is defined within this polynomial space.

Notations. We denote individual elements—such as numbers or polynomials—in italics, e.g. a and b , and their vectors in bold, e.g., $\mathbf{a}$ and $\mathbf{b}$. The notation $\langle \mathbf{a}, \mathbf{b} \rangle$ represents the inner product between the two vectors. We denote $\|a\|_\infty$ as the infinity norm of the vector (or polynomial) a in the power basis and $\text{hw}(a)$ as the Hamming weight of the vector (or polynomial) a . We denote $x \bmod Q$ as the remainder of x modulo Q , $\lfloor x \rfloor$, $\lceil x \rceil$, and $\llbracket x \rrbracket$ as the rounding of x to the nearest previous, next and closest integer, respectively. If x is a polynomial, these operations are performed on each coefficient.

Given an integer a , its residue modulo q is denoted as $[a]_q$. If $\mathcal{C} = \{q_0, q_1, \dots, q_L\}$ is a set of pairwise co-prime positive integers (modulus chain), and $a \in \mathbb{Z}_{Q_L}$, where $Q_L = \prod_{i=0}^L q_i$ is the top modulus. The modulus at the level $\ell \in \{0, 1, \dots, L\}$ is denoted as $Q_\ell = \prod_{i=0}^\ell q_i$. Then the RNS representation of a with respect to $\mathcal{C}$ is denoted by $[a]_{\mathcal{C}} = ([a]_{q_0}, [a]_{q_1}, \dots, [a]_{q_L}) \in \mathbb{Z}_{q_0} \times \dots \times \mathbb{Z}_{q_L}$. The base of the logarithm in this paper is two.

Let $\text{DFT} : \mathbb{R}[X]/X^N + 1 \rightarrow \mathbb{C}^{N/2}$ be a variant of the discrete Fourier transform defined as evaluation of $N/2$ points:

$$\forall p \in \mathcal{R}_N : \text{DFT}(p) = (p(\zeta^{5^i}))_{0 \leq i < N/2},$$

where ζ is a primitive complex $(2N)$ -th root of unity. The inverse of this process is denoted as $\text{iDFT} : \mathbb{C}^{N/2} \rightarrow \mathbb{R}[X]/X^N + 1$.

We also define the transformation $f : \mathbb{F}_2^n \rightarrow \mathbb{F}_2$, which involves intermediate coefficients in $\mathbb{Z}$, as a Boolean function as follows:

Notation 1. Boolean Function: Let $f : \mathbb{F}_2^n \rightarrow \mathbb{F}_2$ be a Boolean function with intermediate coefficients in $\mathbb{Z}$:

$$f(\mathbf{x}) = f(x_0, x_1, \dots, x_{n-1}) = \sum_{\mathbf{u} \in \mathbb{F}_2^n} c_{\mathbf{u}} \prod_{i=0}^{n-1} x_i^{u_i},$$

where $c_{\mathbf{u}} \in \mathbb{Z}$, and

$$\mathbf{x}^{\mathbf{u}} = \prod_{i=0}^{n-1} x_i^{u_i} \quad \text{with} \quad x_i^{u_i} = \begin{cases} x_i, & \text{if } u_i = 1, \\ 1, & \text{if } u_i = 0. \end{cases}$$

denotes the monomials that comprise the polynomial of the output bit.

Basic operations. We first recall some necessary operations defined in the CKKS scheme [CKKS17, CHK⁺18].

- $\text{Ecd}(\mathbf{m}, \Delta, N)$ (*coefficient $\rightarrow$ slots*): Given a message $\mathbf{m} \in \mathbb{C}^{N/2}$, the canonical embedding $\mathbb{C}^{N/2} \rightarrow \mathbb{R}[X]/X^N + 1 \rightarrow \mathcal{R}_N$ is defined as:

$$\forall \mathbf{m} \in \mathbb{C}^{N/2} : \text{Ecd}(\mathbf{m}) = \lfloor \Delta \cdot \text{iDFT}(\mathbf{m}) \rfloor$$

where Δ represents the scaling factor that controls the encoding precision. After this process, the data $\mathbf{m} \in \mathbb{C}^{N/2}$ to be operated on is stored in the slots, with the plaintext polynomial rounded to $\text{Ecd}(\mathbf{m})$. In this form, the multiplication of polynomials corresponds to slot-wise multiplication of scaled $\mathbf{m}$.

- $\text{Dcd}(p, \Delta, N)$ (*slots $\rightarrow$ coefficient*): Given polynomial $p \in \mathcal{R}_N$, apply the inverse process of Ecd , defined as:

$$\forall p \in \mathcal{R}_N : \text{Dcd}(p) = \frac{1}{\Delta} \cdot \text{DFT}(p).$$

Recall that the ciphertext $\text{ct} = (b, a) \in \mathcal{R}_{Q,N}^2$ decrypts to the plaintext polynomial under a secret key $\text{sk} = (1, s)$ as $\langle \text{ct}, \text{sk} \rangle = b + as \approx m$, where $m \in \mathcal{R}_N$.

Key switching (KS) is required after homomorphic multiplication and automorphism. It is defined as $\text{KS}(\text{ct}, \text{swk})$, where swk is generated from the previous secret key sk to facilitate multiplication and rotations. We briefly recall the homomorphic operations as follows.

- **Add**(ct_1, ct_2): For $\text{ct}_1, \text{ct}_2 \in \mathcal{R}_{Q_\ell}^2$, output $\text{ct}_{\text{add}} = \text{ct}_1 + \text{ct}_2 \pmod{Q_\ell}$.
- **Mult**(ct_1, ct_2): Given $\text{ct}_1, \text{ct}_2 \in \mathcal{R}_{Q_\ell}^2$, output a level-downed ciphertext $\text{ct}_{\text{mult}} \in \mathcal{R}_{Q_{\ell-1}}^2$. This operation applies tensor product of $\text{ct}_{\text{new}} = \text{ct}_1 \otimes \text{ct}_2 \in \mathcal{R}_{Q_\ell}^3$, relinearization $\text{ct}_{\text{rel}} = \text{KS}(\text{ct}_{\text{new}}, \text{swk}_{\text{rel}}) \in \mathcal{R}_{Q_\ell}^2$ and rescale $\text{ct}_{\text{mult}} = \text{RS}(\text{ct}_{\text{rel}}) \in \mathcal{R}_{Q_{\ell-1}}^2$.
- **cMult**(ct, a): For $a \in \mathbb{Z}$, output $\text{ct}_{\text{cmult}} = a \cdot \text{ct}$ without applying $\text{Ecd}(\cdot)$, relinearization and rescale $\text{RS}(\text{ct})$ after **Mult**.
- **RS**(ct): For $\text{ct} \in \mathcal{R}_{Q_\ell}^2$, output $\text{ct}_{\text{RS}} = \lfloor q_\ell^{-1} \cdot \text{ct} \rfloor \pmod{Q_{\ell-1}}$. The rescaling process serves two purposes: it reduces the error size and maintains the scaling factor for each slot. Rescale is alongside the reduction of the modulus chain.

Bootstrapping. The bootstrapping procedure of the CKKS scheme [CHK⁺18] seeks to elevate the ciphertext to a higher modulus, enabling further homomorphic evaluation. The CKKS bootstrapping produces a ciphertext that increases the modulus back to certain point, with $Q_{\text{rem}} \gg q_0$, where Q_{rem} is the modulus after bootstrapping. The procedure consists of four steps: **SlotsToCoeffs**, **ModRaise**, **CoeffsToSlots** and **EvalMod**. We borrow the notations in [BCKS24] and briefly explain these steps for clarity.

$$\mathbf{z} \xrightarrow{\text{StC}} z(x) \xrightarrow{\text{ModRaise}} z(x) + q_0 I(x) \xrightarrow{\text{CtS}} \mathbf{z} + q_0 \mathbf{I} \xrightarrow{\text{EvalMod}} \mathbf{z}$$

- **SlotsToCoeffs:** The inverse of the canonical embedding evaluates DFT matrix multiplication homomorphically on the encrypted ciphertext that can decrypt to the vector $\mathbf{z}$. This operation converts it into the polynomial form $z(x)$.
- **ModRaise:** The ciphertext at modulus q_0 is expressed in the larger modulus $Q \gg q_0$. This results in a ciphertext that decrypts to $[\langle \text{ct}, \text{sk} \rangle]_Q = z(x) + q_0 I(x)$, where $q_0 I(x)$ is an integer polynomial with coefficients that are small multiples of the base modulus q_0 .

The remaining steps of the bootstrapping remove this redundant $q_0 I(x)$ polynomial by homomorphically evaluating a modular reduction by q_0 on $z(x)$.

- **CoeffsToSlots:** The canonical embedding evaluates the iDFT matrix multiplication homomorphically on $z(x) + q_0 I(x)$. To enable parallel (slot-wise) evaluation, the polynomial must be encoded in the slot domain, converting it to a ciphertext that decrypts to $\mathbf{z} + q_0 \mathbf{I}$.
- **EvalMod:** A polynomial approximation of the modular reduction by q_0 is homomorphically evaluated, thus removing the redundant $q_0 \mathbf{I}$ and obtaining the result ciphertext decrypted to $\mathbf{z}$.

2.2 Symmetric-Cipher-to-CKKS Transciphering Methods

Advanced transciphering methods for CKKS often function like a stream cipher [CCF⁺16], where the plaintext is masked with a keystream (e.g., XORed). The server pre-generates an encrypted keystream homomorphically using the FHE-encrypted secret key, enabling it

to decrypt incoming symmetric ciphertext by attaching the keystream to cancel the mask and recover the plaintext under FHE encryption.

The Real-to-Finite-field (RtF) framework [CHK⁺21] proposed the method of transciphering to BFV and then bootstrapping the BFV ciphertext into CKKS ciphertext. This approach is used with the specially designed symmetric cipher HERA and Rubato, where the ciphertext domain is modular p , matching the plaintext domain of BFV. This alignment is particularly suitable for applications involving numerical computing tasks. However, it is not compatible with standard symmetric ciphers like AES, which may raise additional security concerns on the client side. For example, the 192-bit security version of HERA was broken in [LKSM24], and the non-prime $\mathbb{F}_p$ parameter Rubato was broken by [GAH⁺23]. In this paper, we focus on the case of Boolean circuit, and thus do not compare with this line of works.

BLEACH methods. We recall the strategy introduced in [DMPS24], which allows Boolean operations using CKKS. In this approach, true and false are represented by 1 and 0, respectively. Utilizing addition, subtraction, and multiplication over real (or complex) numbers, any Boolean gate can be emulated, e.g., XOR, AND, OR, and NOT gates are defined as follows:

$$x \oplus y = (x - y)^2, \quad x \wedge y = x \cdot y, \quad x \vee y = x + y - x \cdot y, \quad \neg x = 1 - x.$$

Except for the free NOT gate, all other gates consume one multiplication depth. Besides, due to the approximate inputs $x + \epsilon_1$ and $y + \epsilon_2$, where $|\epsilon_1|, |\epsilon_2| < 1/4$, the output error is upper bounded by $5 \max(|\epsilon_1|, |\epsilon_2|)$ according to [DMPS24, Le. 2]. After multiple sequential operations, this emulated error can grow and must be reduced by the *cleanup* functions, which move values closer to 0 or 1. Specifically, they use the function $h_1(x) = 3x^2 - 2x^3$ from [CKK20] for this cleanup functionality, i.e., $h_1(0 \pm \epsilon) \simeq 0$ and $h_1(1 \pm \epsilon) \simeq 1$.

BLEACH for AES-CKKS transciphering. [ADE⁺23] designs AES-CKKS transciphering for the first time. They employ the BLEACH method to evaluate the AES circuit homomorphically. We briefly recall the AES circuit for better understanding. To begin with encryption, the plaintext block is XORed with the 128-bit key in the whitening step. Then, the AES round function is applied to update the resulting 128-bit state. Each round of AES consists of the following operations:

$$\text{AddRoundKey} \circ \underset{1}{\text{MixColumns}} \circ \underset{3}{\text{ShiftRows}} \circ \underset{0}{\text{SubBytes}}(\cdot).$$

This round function is repeated 9, 11, or 13 times for AES-128, AES-192, and AES-256, respectively. The MixColumn operation is excluded in the final round. After operating the remaining update function, the resulting state is output as the ciphertext.

AES-CKKS transciphering is ideally suited for the counter (AES-CTR) mode, where multiple AES blocks—each employing a 32-bit nonce and a 96-bit counter as input—are allocated to distinct slots in the CKKS plaintext, thereby enabling Single Instruction Multiple Data (SIMD) operations. This strategy allows each AES block to be processed independently, eliminating the need for costly cross-slot operations. The homomorphic look-up table (LUT) is employed for the AES 8-bit SubBytes operation, requiring 255 multiplications with a depth of three in their implementation, although the theoretical lower bound is $n - \log n - 1 = 247$. It is important to note that in the BLEACH method, the XOR operation, defined as $(x - y)^2$, is not free. The MixColumn and AddRoundKey operations, which involve XOR, consume multiplication depths of 3 and 1, respectively. Additionally, following several Boolean gate operations, the clean function $h_1(x) = 3x^2 - 2x^3$, which has a depth of two, is executed in every AES round. Overall, each AES round requires a multiplication depth of $3 + 3 + 1 + 2 = 9$. Consequently, ten rounds of AES require ten

repetitions of bootstrapping on each ciphertext (typically processed 128 in parallel in each round) to refresh the available depth.

2.3 Revisiting the Details of AES S-Box Monomial Construction

SubBytes is a key step in AES, but the authors of [ADE⁺23] just mentioned an indicator-based SubBytes homomorphic LUT method, and neither the algorithm nor the source code was available. To fill this gap, we introduced a homomorphic LUT algorithm that achieved the theoretical lower bound of 247 multiplications.

Let $x_0, x_1, \dots, x_{\lceil \log n \rceil - 1}$ represent the input bits of a LUT with n entries. The index of a single entry in this table can be expressed as $x = \sum_{i=0}^{\lceil \log n \rceil - 1} x_i \cdot 2^i$. Each bit x_i is an encrypted $\{0, 1\}$ value, with its inverse given by $\bar{x}_i = 1 - x_i$. We have the LUT result written as:

$$\begin{aligned} f(x) = & (\bar{x}_{\log n-1} \cdots \bar{x}_1 x_0) \cdot f(0) + (\bar{x}_{\log n-1} \cdots \bar{x}_1 x_0) \cdot f(1) \\ & + (\bar{x}_{\log n-1} \cdots \bar{x}_2 x_1 \bar{x}_0) \cdot f(2) + (\bar{x}_{\log n-1} \cdots \bar{x}_2 x_1 x_0) \cdot f(3) \\ & + \cdots + (x_{\log n-1} \cdots x_1 x_0) \cdot f(n-1), \end{aligned}$$

where $f(x)$ denotes the result of this bijective look-up table. It can be observed that iff the binary representation of the input index x (denoted as $x \Leftrightarrow (x_0, x_1, \dots, x_{\lceil \log n \rceil - 1})$) matches, the cumulative product within the parentheses equals one. The LUT result can be obtained without revealing the indices by summing all these values.

Given the definition $\bar{x}_i = 1 - x_i$, we can observe that each output bit of the homomorphic LUT for the S-boxes is the sum of multiple monomials of the input bits. Therefore, once all the monomials are computed, the output bit can be obtained through homomorphic additions. The monomials corresponding to the input bits are as follows:

$$x_0, x_1, \dots, x_{\log n-1}, x_0 x_1, x_0 x_2, \dots, \prod_{i=0}^{\log n-1} x_i.$$

We present a recursive algorithm in Algorithm 1 to achieve this lower bound for an 8-bit input with 247 multiplications and a depth of three. This algorithm is later adapted to operate on real FHE ciphertexts. To clarify, we illustrate a toy example of 4 input variables in Fig. 1.




Figure 1: Toy example of four input variables for EvalRecurse algorithm.

As depicted in Fig. 1, the input variables are divided into 2 sets, each containing two variables. This division corresponds to the deepest layer of the EvalRecurse algorithm. After multiplying each element in the set, the degree-two monomials are merged with the

Algorithm 1: EvalRecurse

Input: Monomial set $\mathcal{M}_{\text{in}}\{x_0, x_1, \dots, x_{\log n-1}\}$
Output: Monomial set $\mathcal{M}_{\text{out}}\{x_1, \dots, x_{\log n}, x_1x_2, x_1x_3, \dots, \prod_{i=0}^{\log n-1} x_i\}$

```

1 if  $|\mathcal{M}_{\text{in}}| = 2$  then
2   return  $\{\mathcal{M}_{\text{in}}[0], \mathcal{M}_{\text{in}}[1], \text{Mult}(\mathcal{M}_{\text{in}}[0], \mathcal{M}_{\text{in}}[1])\}$ 
3  $idx \leftarrow \frac{|\mathcal{M}_{\text{in}}|}{2}$ 
4  $\mathcal{M}_L \leftarrow \text{EvalRecurse}(\mathcal{M}_{\text{in}}[:idx])$ 
5  $\mathcal{M}_R \leftarrow \text{EvalRecurse}(\mathcal{M}_{\text{in}}[idx:])$ 
6  $\mathcal{M}_{\text{out}} \leftarrow \{\}$ 
7 for  $\text{mon}_i \in \mathcal{M}_L$  do
8   for  $\text{mon}_j \in \mathcal{M}_R$  do
9      $\text{mon}_{\text{new}} \leftarrow \text{Mult}(\text{mon}_i, \text{mon}_j)$ 
10     $\mathcal{M}_{\text{out}} \leftarrow \mathcal{M}_{\text{out}} \cup \text{mon}_{\text{new}}$ 
11  $\mathcal{M}_{\text{out}} \leftarrow \mathcal{M}_{\text{out}} \cup \mathcal{M}_L \cup \mathcal{M}_R$ 
12 return  $\mathcal{M}_{\text{out}}$ 

```

original variables and sent to the next recursion. Finally, each element in the set obtained at level one is multiplied, merging with the set obtained in the last recursion. The total multiplications consumed are $n - \log n - 1 = 11$ with depth 2.

To combine all the monomials into the output bit of the AES S-box, we use the notation defined in Notation 1 for Boolean functions and provide a detailed explanation of the AES SubBytes operation in Sect. 4.

3 The Free-XOR Technique

At a high level, our idea is based on two key observations: (i) The XOR of multiple binary elements can be translated into the least significant bit (LSB) extraction, which is equivalent to performing a modular reduction by 2 on their aggregated sum. (ii) Since CKKS operations are defined over the ciphertext modulus q , any overflow beyond q is automatically handled by reducing the values modulo q .

In this section, we introduce our strategy that automatically handles the overflows and achieves Free-XOR.

3.1 Automated Overflow in CKKS

We recall the bootstrapping procedure defined in Sect. 2.1. In this process, the ciphertext, which decrypts to a scaled message $\mathbf{z} + q_0\mathbf{I}$, is processed during the EvalMod step. A key aspect of bootstrapping is the homomorphic removal of the redundant term $q_0\mathbf{I}$, which involves the computationally intensive task of performing homomorphic modular reduction by q_0 .

By observing overflow behavior when the encrypted message exceeds the modulus q , we propose a framework that initially accumulates Boolean values into an integer and then applies modular reduction by retaining only the LSB through the overflow. This approach integrates seamlessly into the bootstrapping procedure, occurring naturally during the SlotsToCoeffs operations. The overflow mechanism leverages the inherent modular q operation, incurring minimal additional cost to the original bootstrapping process and thus providing a free functional enhancement.

Modulus and scale management of SlotsToCoeffs for automated overflow. Consider a scenario where τ consecutive XOR operations are performed, resulting in a large integer within the range $[0, \tau]$. To handle values that exceed the modulus capacity, we scale the value so the LSB aligns with the MSB of the modulus. This operation involves multiplying ciphertexts by constant plaintext values, which is most efficiently achieved by merging these values and pre-multiplying the resulting constants with the DFT matrices, as suggested by [BMTH21].

Specifically, we employ the SlotsToCoeffs-first bootstrapping approach, which requires starting the process with a modulus larger than q_0 . Assume the transformation matrix for DFT is factorized into deg matrices and the difference in the scaling factor between the shifted ciphertexts and the current one is given by $\delta_{\text{diff}} = q_0/2\Delta_{\text{StC}}$, where the ciphertext at the StC level has a scale of Δ_{StC} . Consequently, each factorized matrix is scaled by $\mu_{\text{StC}} = (\delta_{\text{diff}})^{\frac{1}{\text{deg}}}$. Then, we integrate the automated overflow into SlotsToCoeffs as follows.

Starting with the input ciphertext at the StC level, which decrypts to

$$\mathbf{z} = \Delta_{\text{StC}} \mathbf{m} \mod Q_{\text{StC}}, \quad \text{where } Q_{\text{StC}} = q_0 \cdots q_{\text{deg}},$$

we scale each factorized matrix by the factor μ_{StC} , such that $M'_i \leftarrow \mu_{\text{StC}} \times M_i$. The encoded DFT submatrices $\widetilde{M}_1, \dots, \widetilde{M}_{\text{deg}}$, with $\widetilde{M}_i = \text{Ecd}(M'_i, q_{i+1}, N)$, are then applied as follows:

$$\begin{aligned} \text{ct}_{\text{deg}-1} &\leftarrow \text{RS}_{\Delta_{\text{deg}}}(\widetilde{M}_{\text{deg}} \cdot \text{ct}_{\text{deg}}) \quad \text{with} \quad \text{ct}_{\text{deg}-1} = \text{Enc}_{\text{sk}}(\Delta_{\text{deg}-1}(\mu_{\text{StC}} \cdot \mathbf{m}')), \\ &\vdots \\ \text{ct}_0 &\leftarrow \text{RS}_{\Delta_1}(\widetilde{M}_1 \cdot \text{ct}_1) \quad \text{with} \quad \text{ct}_0 = \text{Enc}_{\text{sk}}(\Delta_0 \cdot m(x)). \end{aligned}$$

This scale management enables homomorphic decoding without sacrificing the precision of SlotsToCoeffs, as summarized by the following formula:

$$\mathbf{z} = \Delta_{\text{StC}} \mathbf{m} \mod Q_{\text{StC}} \xrightarrow{\text{SlotsToCoeffs}} \Delta_0 b(x) = \Delta_0 m(x) \mod q_0.$$

After this operation, the message is transformed into its coefficient form, and the modulus is reduced to the lowest level. Each coefficient is observed to decompose into two parts: the binary part and the even-number part, as defined by

$$\Delta_0 m(x) \mod q_0 = \Delta_0(b(x) + 2I'(x)) \mod q_0,$$

where $I'(x)$ is a polynomial with integer coefficients bounded by $\lceil \frac{\tau-1}{2} \rceil$, and $b(x)$ is a polynomial with coefficients in $\{0, 1\}^N$. Since the LSB has been shifted to the MSB (i.e., $\Delta_0 = q_0/2$), we obtain

$$\frac{q_0}{2} \cdot b(x) \mod q_0 \leftarrow \frac{q_0}{2} \cdot b(x) + \frac{q_0}{2} \cdot 2I'(x) \mod q_0.$$

This reduction occurs automatically because the redundant polynomial $I'(x)$ has coefficients that are multiples of the base modulus q_0 . In other words, this overflow happens during the SlotsToCoeffs operation by utilizing the scale management technique. Subsequently, the bootstrapping procedures are applied in a manner similar to Boolean bootstrapping [BCKS24], as follows:

$$\frac{q_0}{2} b(x) \xrightarrow{\text{ModRaise}} \frac{q_0}{2} b(x) + q_0 I(x) \xrightarrow{\text{CoeffsToSlots}} q_0 \left(\frac{1}{2} \mathbf{b} + \mathbf{I} \right).$$

After applying the approximated trigonometric function $\frac{1}{2}(1 - \cos(2\pi x))$ with the clean functionality, we can remove the redundant term $\mathbf{I}$. The base prime q_0 serves as the scale factor during the EvalMod phase and is often adjusted to match the modulus at this level.

Other methods on modular reduction in CKKS. A very recent work by Kim et al. [KN24] explores the general case of modular reduction in CKKS. Since their work focuses on computation over integers, they suggest to multiply an encrypted value of scale q_0/t^l by t^{l-1} after `SlotsToCoeffs` to extract its LSB. However, this approach increases the noise by t^{l-1} times and might not be suitable for Boolean circuits. Take transciphering as an example, if an AES-encrypted bit serves as the Most Significant Bit (MSB) of an input of the subsequent privacy-preserving application, even minor noise on that bit could result in large errors in the function’s output.

3.2 XBOOT: Automated Overflow Combined with Bootstrapping

Algorithm 2: XBOOT

Input: $\text{ct} = \text{Enc}_{\text{sk}}(\Delta_{\text{StC}} \mathbf{m} + \epsilon)$ with $\|\epsilon\|_\infty \ll 1$, where $\text{ct} \in \mathcal{R}_{Q_{\text{StC}}}^2$

Output: $\text{ct}_{\text{out}} \in \mathcal{R}_{Q_{\text{rem}}}^2$

- 1 Setting: $\delta_{\text{diff}} = q_0/2\Delta_{\text{StC}}$
 - 2 $\text{ct}_{\text{bin}} \leftarrow \text{SlotsToCoeffs}(\text{ct})$ // Automated overflow.
 - 3 $\text{ct}' \leftarrow \text{CoeffsToSlots} \circ \text{ModRaise}(\text{ct}_{\text{bin}})$
 - 4 $\text{ct}_{\text{real}} \leftarrow (\text{conj}(\text{ct}') + \text{ct}')/2$
 - 5 $\text{ct}_{\text{out}} \leftarrow \text{EvalMod}_{f_{\text{Bin}}}(\text{ct}_{\text{real}})$ // $f_{\text{Bin}} = \frac{1}{2}(1 - \cos(2\pi x))$.
 - 6 **return** ct_{out}
-

We describe a combination of automated overflow and binary bootstrapping in Algorithm 2. It takes as input a ciphertext $\text{ct} = \text{Enc}_{\text{sk}}(\Delta_{\text{StC}} \mathbf{m} + \epsilon)$, which decrypts to a vector of integers. After applying homomorphic decoding with automated overflow in line 2, we obtain $\text{ct}_{\text{bin}} = \text{Enc}_{\text{sk}}\left(\frac{b(x)}{2} + \epsilon'(x)\right)$, where the message’s LSB is automatically shifted to the MSB of the base modulus such that

$$\frac{q_0}{2}b(x) \leftarrow \frac{q_0}{2}b(x) + q_0I'(x) \pmod{q_0}.$$

After line 3, we have $\text{ct}' = \text{Enc}_{\text{sk}}\left(\frac{b+\epsilon_c}{2} + \mathbf{I}\right)$, where $\mathbf{I} \in \mathbb{Z}^{N/2}$ is an integer vector of small magnitude and $\epsilon_c \in \mathbb{C}^{N/2}$ represents noise related to the original ϵ' and the homomorphic operations. In line 4, the real part of ct' is extracted, yielding $\text{ct}_{\text{real}} = \text{Enc}_{\text{sk}}\left(\frac{b+\epsilon_r}{2} + \mathbf{I}\right)$, where $\epsilon_r \in \mathbb{R}^{N/2}$.

Finally, after the homomorphic evaluation of the polynomial approximation of $f_{\text{Bin}} = \frac{1}{2}(1 - \cos(2\pi x))$, we obtain $\text{ct}_{\text{out}} = \text{Enc}_{\text{sk}}(\mathbf{b} + \epsilon_{\text{out}})$ at a higher modulus Q_{rem} . We omit the noise analysis as it is similar to the noise estimation in the original binary bootstrapping [BCKS24, Thm 1].

Reduced interpolation interval. In XBOOT, a sparser secret key and a lower ring degree are utilized, allowing us to reduce the interpolation interval without compromising the failure probability of bootstrapping. We describe the rationale behind as follows.

Recall that each coefficient of the redundant polynomial after the `ModRaise` step follows the shifted Irwin–Hall distribution [LLL⁺21] over the interval $[-1/2, 1/2]$. The failure probability can be calculated using the adapted cumulative distribution function (CDF) corresponding to the binomial distribution for N trials, as specified in Bossuat et al. [BMTH21, Eq. 1].

A lower ring degree could lead to a reduced interpolation interval due to fewer trials in the binomial distribution. We revisit the distribution to compute the adapted K values, as shown in Table 1.

Table 1: Adapted K of $\Pr[\|I(x)\|_\infty > K] \approx 2^{-16}$ with variable h .

$\log N$	15									
$\log(h)$	5	6	7	8	9	10	11	12	13	14
K	10	14	20	28	40	56	80	113	160	226
$\log(\Pr[\ I(x)\ _\infty > K])$	-17.8	-15.6	-15.6	-14.7	-15.2	-14.5	-15.1	-15.0	-15.0	-15.0
$K/\sqrt{h}$	1.77	1.75	1.77	1.75	1.77	1.75	1.77	1.77	1.77	1.77

$\log N$	14									
$\log(h)$	4	5	6	7	8	9	10	11	12	13
K	7	10	14	20	28	39	56	79	111	158
$\log(\Pr[\ I(x)\ _\infty > K])$	-23.4	-18.8	-16.6	-16.6	-15.7	-14.8	-15.5	-15.4	-15.0	-15.3
$K/\sqrt{h}$	1.75	1.77	1.75	1.77	1.75	1.72	1.75	1.75	1.73	1.75

Additionally, [BTH22] suggests adopting a sparse secret key to reduce the complexity of the EvalMod step. This approach does not compromise security, as the sparse operates at a smaller modulus. XBOOT can significantly benefit from this technique, especially given its smaller base modulus. When combined with a smaller ring degree ($\log N \leq 15$), it results in a reduced interpolation interval K , thereby enhancing bootstrapping efficiency.

3.3 Possible optimizations and limitations

Lazier reduction. The previous sections focused on the specific scenario where bootstrapping is applied immediately after the XOR gates. However, it is important to note that XBOOT is also applicable in cases where XOR and AND gates are interleaved before bootstrapping. Consider a toy example of two values x, y , each results from k instances of dense XOR operations on Boolean values. Then an AND operation is applied to these two values. We can deduce the following:

$$\left(x = \bigoplus_{i=0}^k x_i\right) \wedge \left(y = \bigoplus_{i=0}^k y_i\right) \equiv \text{LSB} \left(\left(x = \sum_{i=0}^k x_i\right) \times \left(y = \sum_{i=0}^k y_i\right) \right).$$

This equivalence holds because

$$x \wedge y = 1 \iff \text{LSB} \left(\sum_{i=0}^k x_i \right) = 1 \text{ and } \text{LSB} \left(\sum_{i=0}^k y_i \right) = 1,$$

Therefore, we do not need to insert XBOOT immediately after the XOR gates; instead, we can be lazier and insert XBOOT after the AND gate. Similar formulas would apply to any Boolean circuit, where we can replace all XOR gates in the circuit with additions, analyze the dependencies of the computation graph, and strategically adjust the placement of XBOOT without changing the original circuit's topology.

Limitations. The BLEACH method of cleaning noise using $h_1(x) = 3x^2 - 2x^3$ only works with Boolean values and is incompatible with integer values. Consequently, we can only rely on bootstrapping for noise cleaning. In very rare cases, such as circuits primarily composed of AND gates with only a few XOR gates, our framework may not outperform BLEACH. However, for most practical scenarios where the number of XOR gates is higher or comparable to that of AND gates, our framework consistently performs better.

4 Application to AES Transciphering

XBOOT immediately increases the efficiency for CKKS computations on Boolean circuits. Consider the well-known AES transciphering as an example: the AES round function consists of four key steps—SubBytes, ShiftRows, MixColumns, and AddRoundKey. The best previous work on AES-CKKS transciphering [ADE⁺23] requires 9 multiplication depths plus one bootstrapping per round. We can transform the XOR operations in MixColumns and AddRoundKey into inexpensive homomorphic additions, followed by XBOOT. This transformation reduces the cost per round to 3 multiplication depths plus one bootstrapping. We describe the details in the next subsections.

4.1 Details of SubBytes Operation

Algorithm 3: SubBytes (Homomorphic LUT)

Input: Encrypted bits $\mathbf{x}_0, \mathbf{x}_1, \dots, \mathbf{x}_7$ with each bit $\mathbf{x}_i \in \{0 + \epsilon, 1 + \epsilon\}^{N/2}$
Output: $\mathbf{sbox}$ as the LUT output such that $\{\mathbf{y}_0, \mathbf{y}_1, \dots, \mathbf{y}_7\} = \text{S-box}(\mathbf{x}_0, \mathbf{x}_1, \dots, \mathbf{x}_7)$

```

// Obtain the monomial table for each AES S-box output bit.
1 for  $i = 0$  to 7 do
2    $T_i \leftarrow \text{AES S-box LUT}$ 

// Compute all required monomials homomorphically.
3  $\mathcal{M} \leftarrow \text{EvalRecurse}(\mathbf{x}_0, \dots, \mathbf{x}_7) // |\mathcal{M}| = 255$ 

// Initialize each bit with the constant term of the monomial.
4  $\mathbf{sbox}[i] \leftarrow \{\text{Ecd}(\text{const}_i)\}$  for  $i \in \{0, 1, \dots, 7\}$ 
5 for  $i = 0$  to 7 do
6   // Gather coefficients for each output bit
7   for  $j = 1$  to 255 do
8     if  $T_i[j] \neq 0$  then
9        $\text{mon} \leftarrow \text{cMult}(\mathcal{M}[j], T_i[j]) // \text{mon is of the form } c_{\mathbf{u}} \cdot (\prod_{i=0}^7 \mathbf{x}_i^{u_i})$ 
10       $\mathbf{sbox}[i] \leftarrow \text{Add}(\mathbf{sbox}[i], \text{mon})$ 
10 return  $\mathbf{sbox}$ 

```

As described in Section 2.3, we employ a recursive method in Algorithm 1 to obtain all the monomials required for the AES S-box LUT using 247 multiplications with a depth of 3. To achieve the final result of the SubBytes, we utilize a strategy that combines all the single terms of monomials into the Boolean function in Notation 1 for each S-box’s output bit. The procedure is detailed in Algorithm 3 and described as follows.

The input to the SubBytes operation consists of encrypted 8-bit ciphertexts. Each ciphertext decrypts to a Boolean vector $\mathbf{x}_i$, where each element represents a specific bit of the state from $N/2$ different blocks. These blocks are processed using SIMD operations, resulting in the encrypted 8-bit output $\{\mathbf{y}_0, \mathbf{y}_1, \dots, \mathbf{y}_7\}$.

- **Preprocess phase:** As shown in lines 1–2, we precompute the coefficients of the Boolean function for each output bit of the AES S-box and store them in the plaintext table T_i . This procedure involves applying the AES S-box LUT to the input variable and storing the result in the coefficient table.
- **Monomial making:** This procedure is the computational bottleneck, requiring 247 multiplications with a depth consumption of 3. After applying the EvalRecurse function homomorphically, we obtain all monomials in the set $\mathcal{M}$.

- **Boolean function:** We consider each monomial in the obtained set $\mathbf{x}^{\mathbf{u}} \in \mathcal{M}$ as defined in Notation 1. This procedure involves combining these monomials with corresponding coefficients and summing them together. Specifically, we apply a scalar multiplication cMult to each monomial $\mathcal{M}[j]$ using the constants from the pre-computed coefficient table T_i , corresponding to the AES S-box. By summing all these coefficient-adjusted monomials mon , we derive the Boolean function for a specific output bit.

The **Boolean function** phase of the **SubBytes** operation are highly efficient through the use of cMult with integer inputs from the precomputed table. This design eliminates the need for repeated homomorphic additions, which would otherwise incur a higher computational cost.

4.2 Free XOR AES Evaluation

We compare our free XOR procedure with the original procedure in Fig. 2 and provide a detailed implementation in Algorithm 4, which is described below.




Figure 2: Original AES evaluation procedure (left) and our free XOR AES evaluation procedure (right).

Algorithm 4: Homomorphic AES Evaluation

Input: $\mathbf{S} = \{n||c\}^{N/2}$, where each block takes a 32-bit nonce and a 96-bit counter.
Input: $\mathbf{K}_i$, where $0 \leq i \leq 10$ are the round keys duplicated for $N/2$ times.
Output: $\text{ct}_{\text{out}} = \text{Enc}_{\text{sk}}(\text{keystream} + \epsilon_{tr})$, FHE encrypted keystreams.

```

1  $\mathbf{S} \leftarrow \text{Ecd}(n||c)$ 
2  $\text{state} \leftarrow \text{AddWhiteKey}(\mathbf{S}, \mathbf{K}_0)$  // Apply whitening key with XOR
3 for  $i = 1$  to 9 do
4   for  $j = 0$  to 15 do
5      $\text{state}[8 * j : 8 * (j + 1)] \leftarrow \text{SubBytes}(\text{state}[8 * j : 8 * (j + 1)])$ 
6    $\text{state} \leftarrow \text{ShiftRows}(\text{state})$ 
7   for  $j = 0$  to 3 do
8     // Apply homomorphic addition instead of XOR
9      $\text{col} \leftarrow \text{FMixColumns}(\text{state}[32 * j : 32 * (j + 1)])$ 
10     $\text{state}[32 * j : 32 * (j + 1)] \leftarrow \text{col}$ 
11   $\text{state} \leftarrow \text{FAddRoundKey}(\text{state}, \mathbf{K}_i)$ 
12   $\text{state} \leftarrow \text{XBOOT}(\text{state})$ 
13 // Last Round
14 for  $i = 0$  to 15 do
15    $\text{state}[8 * i : 8 * (i + 1)] \leftarrow \text{SubBytes}(\text{state}[8 * i : 8 * (i + 1)])$ 
16  $\text{state} \leftarrow \text{ShiftRows}(\text{state})$ 
17  $\text{state} \leftarrow \text{FAddRoundKey}(\text{state}, \mathbf{K}_{10})$ 
18  $\text{ct}_{\text{out}} \leftarrow \text{XBOOT}(\text{state})$ 
19 return  $\text{ct}_{\text{out}}$ 

```

As shown in Fig. 2, the newly proposed free XOR AES evaluation uses the same input as the original one. Specifically, each 128-bit AES block takes a 32-bit nonce and a 96-bit counter as input. The vector of these blocks is denoted as $\mathbf{S}$, where each element $\mathbf{S}[i] = \mathbf{B} = \{B_0, B_1, \dots, B_{\frac{N}{2}}\}$ for $0 \leq i < 128$. This notation can also be expressed by gathering $\frac{N}{2}$ nonces and counters as $\mathbf{S} = \{n||c\}^{N/2}$. During the encryption process, the currently processed cipher block is also called the cipher's state.

After applying $\text{Ecd}(\cdot)$ to the batched state bits, we obtain 128 plaintext polynomials in $\mathcal{R}_N$, storing all $N/2$ blocks of AES-CTR input values, with each slot containing a noised Boolean value. First, we perform the **AddWhiteKey** operation by using the XOR operation on the obtained plaintext polynomial $\mathbf{S}$ with the duplicated encrypted symmetric key $\mathbf{K}_0$ provided by the client. This symmetric key, encrypted under sk , is denoted as $\mathbf{K}_0 = \underbrace{\{k_0, k_0, \dots, k_0\}}_{N/2 \text{ times}}$, which represents $N/2$ repetitions of the symmetric key.

Afterwards, we apply the 9-round AES free XOR function to update the state.

- **SubBytes:** As shown in Algorithm 3, the input consists of encrypted bits, and the output is the encrypted 8-bit AES S-box look-up table (LUT) result.
- **ShiftRows:** Each bit of the state is encrypted in a different FHE ciphertext, which only involves changing the indices of the ciphertext.
- **FMixColumns:** Let each column of the state as $[b_0, b_1, b_2, b_3]^\top$. According to [FRL19],

we have the updated column as:

$$\begin{aligned} b'_0 &= x \cdot (b_0 + b_1) + b_1 + b_2 + b_3 \\ b'_1 &= x \cdot (b_1 + b_2) + b_2 + b_3 + b_0 \\ b'_2 &= x \cdot (b_2 + b_3) + b_3 + b_0 + b_1 \\ b'_3 &= x \cdot (b_3 + b_0) + b_0 + b_1 + b_2. \end{aligned}$$

Let $(a_7, a_6, \dots, a_0)$ denote the 8 bits of the byte. Multiply by x in $\text{GF}(2^8)$ is

$$(a_7, a_6, a_5, a_4, a_3, a_2, a_1, a_0) = (a_6, a_5, a_4, a_3 + a_7, a_2 + a_7, a_1, a_0 + a_7, a_7).$$

In **FMixColumns**, the $+$ operation is translated into homomorphic addition, resulting in a small integer in each slot of the FHE ciphertext.

- **FAddRoundKey**: 128 parallel homomorphic additions are applied to the state.
- **XBOOT**: During **StC**, the bounded small integer in each slot is transformed into a Boolean value through automated overflow. This operation includes bootstrapping, ensuring clean functionality.

By repeating the round function described above for 9 rounds and then applying a final round that performs the same operations except for **FMixColumns**, one can obtain $N/2$ blocks of the keystream, denoted as ct_{out} in Algorithm 4. It is important to note that this procedure is independent of the messages and is considered part of the offline phase. After homomorphically XORing with the encoded $N/2$ blocks of the symmetric ciphertext $\text{Ecd}(c)$, the FHE-encrypted messages can be obtained.

4.3 Further Optimization

In each AES round function, 128 FHE ciphertexts are updated with each operation defined above, resulting in 128 parallel **XBOOT** calls, which can be effectively multi-threaded.

Algorithm 5: BatchXBOOT

Input: $\text{ct}_0, \text{ct}_1 \in \mathcal{R}_{Q_{\text{StC}}}^2$ that decrypts to $(\mathbf{m}_0 + \epsilon_0), (\mathbf{m}_1 + \epsilon_1)$ with $\epsilon_0, \epsilon_1 \in \mathbb{R}^{N/2}$
Output: $\text{ct}_0^{\text{out}}, \text{ct}_1^{\text{out}} \in \mathcal{R}_{Q_{\text{rem}}}^2$

- 1 Setting: $\delta_{\text{diff}} = q_0/2\Delta_{\text{StC}}$
- 2 $\text{ct} = \text{ct}_0 + \text{cMult}(\text{ct}_1, 1i) \ // \ \text{ct} = \text{Enc}_{\text{sk}}((\mathbf{m}_0 + i\mathbf{m}_1) + \epsilon')$
- 3 $\text{ct}_{\text{bin}} \leftarrow \text{SlotsToCoeffs}(\text{ct})$
- 4 $\text{ct}' \leftarrow \text{CoeffsToSlots} \circ \text{ModRaise}(\text{ct}_{\text{bin}})$
- 5 $\text{ct}_0^{\text{real}} \leftarrow (\text{conj}(\text{ct}') + \text{ct}')/2$
- 6 $\text{ct}'' = \text{cMult}(\text{ct}', -1i) \ // \ \text{ct}'' = \text{Enc}_{\text{sk}}((\mathbf{b}_1 - i\mathbf{b}_0 + \epsilon'')/2 + \mathbf{I})$
- 7 $\text{ct}_1^{\text{real}} \leftarrow (\text{conj}(\text{ct}'') + \text{ct}'')/2$
- 8 $\text{ct}_0^{\text{out}} \leftarrow \text{EvalMod}_{f_{\text{Bin}}}(\text{ct}_0^{\text{real}})$
- 9 $\text{ct}_1^{\text{out}} \leftarrow \text{EvalMod}_{f_{\text{Bin}}}(\text{ct}_1^{\text{real}})$
- 10 **return** $\text{ct}_0^{\text{out}}, \text{ct}_1^{\text{out}}$

We pack two ciphertexts that encrypt real numbers into a single complex number ciphertext to reduce the **XBOOT** calls. After **CoeffsToSlots**, both the real and imaginary parts of the ciphertext are extracted. We then apply **EvalMod** to both parts, obtaining two ciphertexts from a single **BatchXBOOT** operation in Algorithm 5. As a result, we can decrease the 128 **XBOOT** operations required in one AES round to 64 **BatchXBOOT** operations.

5 Unlocking New Potentials: Application to Rasta

XBOOT also unlocks new possibilities for CKKS, particularly in scenarios where the number of XOR gates is prohibitively large. For example, several symmetric ciphers have been defined over $\mathbb{Z}_2$ including LowMC [DEM15], Rasta [DEG⁺18], Dasta [HL20], and Fasta [CIR22]. What these ciphers have in common is their extensive use of XOR gates, which presents significant challenges for their evaluation under previous CKKS-based methods. The designers of these ciphers did not regard this as an issue, as they were primarily designed with BGV or MPC transciphering in mind, where XOR gates are relatively inexpensive. After all, BLEACH’s idea of CKKS computation on Boolean circuits had not been proposed at that time.

As BLEACH’s concept is quite new, applications that leverage CKKS for Boolean circuits (other than transciphering) are relatively scarce. However, Boolean applications such as secure computations of Hamming distance for privacy-preserving biometric authentication [MPR20], and the Game of Life [DMPS24, Section 5] for exploring secure cellular automata simulations or cryptographic puzzle design, are among the most directly applicable cases.

In this section, we still focus on transciphering and use Rasta as an example to demonstrate how XBOOT could make the difference.

The incompatibility between Rasta and BLEACH. Rasta incorporates randomly generated linear matrices within its linear layers to resist statistical attacks. This approach is motivated by the observation that statistical attacks require numerous queries to the same encryption function to gather sufficient data for formulating a statistical model.

Let $A \in \mathbb{R}_{n \times n}$ represent the randomly generated $n \times n$ square matrices. When evaluating matrix multiplication, we consider i -th row of the matrix $[a_{(i,0)}, a_{(i,1)}, \dots, a_{(i,n-1)}]$. The updated column is then given by:

$$y_0 = \bigoplus_{j=0}^{n-1} (a_{(0,j)} \wedge x_j), \quad y_1 = \bigoplus_{j=0}^{n-1} (a_{(1,j)} \wedge x_j), \quad \dots, \quad y_{n-1} = \bigoplus_{j=0}^{n-1} (a_{(n-1,j)} \wedge x_j),$$

where y_j and x_j denote the updated element and input element of a column, respectively.

As shown, each update incurs a substantial number of AND and XOR operations. The AND operations are applied to ciphertexts x and plaintexts a , thus are inexpensive. In contrast, the XOR operations involve ciphertext-ciphertext interactions. In BLEACH, the XOR function is represented as $(x - y)^2$, which consumes one multiplication. A 351-bit state in Rasta would require $351 \times (351 - 1) = 122,850$ homomorphic multiplications for applying the matrix multiplication if we use BLEACH.

Application of XBOOT to Rasta. We briefly revisit Rasta’s keystream generation in Fig. 3. Note that other FHE-friendly $\mathbb{Z}_2$ ciphers that include matrices operations also benefit from XBOOT.

Rasta employs the randomly generated permutation $P_{nc,i}$ on the symmetric key K and XOR the same K with the output of the permutation to output keystream $K_{nc,i} = K \oplus P_{nc,i}$. The permutation is defined as:

$$P_{nc,i} = A_{r,nc,i} \circ S \circ A_{r-1,nc,i} \circ \dots \circ S \circ A_{1,nc,i} \circ S \circ A_{0,nc,i}(K),$$

where r denotes the number of rounds based on the cipher version. Each matrix in $A_{r,nc,i}$ is generated by an extendable-output function (XOF) [PF15], depending on the public nonce nc and the block counter i . The affine layer is defined as $A_{r,nc,i} = M_{r,nc,i} \cdot \mathbf{x} + c_{r,nc,i}$, which involves a binary multiplication of the random matrix $M_{r,nc,i}$ followed by the addition of a round constant $c_{r,nc,i}$. The non-linear χ function layer S is defined as




Figure 3: Rasta keystream generate function. $S, A_{r,nc,i}$ denote the nonlinear function and the randomized matrix, respectively.

$y_i = x_i + x_{i+2} + (x_{i+1} \wedge x_{i+2})$, where $0 \leq i < n$, with indices i taken modulo n . We refer [DEG⁺18] for details.

In our implementation, we also consider the same data packing method as we did in AES-CKKS transciphering, with each bit of the Rasta state encoded in different FHE ciphertexts. This enables the simultaneous evaluation of $\frac{N}{2}$ cipher blocks. Each bit of the state is updated with S layers, involving one multiplication with a depth of one and producing a small integer bounded by three. This efficiency gain benefits from the interleaved XOR and AND operations in free XOR.

When evaluating the affine layer $A_{r,nc,i} = M_{r,nc,i} \cdot \mathbf{x} + c_{r,nc,i}$, all $\frac{N}{2}$ random matrices are encoded into $(n \times n)$ FHE plaintexts, while the round constants are encoded into n plaintexts. The cipher state is updated as

$$\mathbf{y}_0 = \sum_{j=0}^{n-1} (\mathbf{a}_{(0,j)} \wedge \mathbf{x}_j), \quad \mathbf{y}_1 = \sum_{j=0}^{n-1} (\mathbf{a}_{(1,j)} \wedge \mathbf{x}_j), \quad \dots, \quad \mathbf{y}_{n-1} = \sum_{j=0}^{n-1} (\mathbf{a}_{(n-1,j)} \wedge \mathbf{x}_j),$$

where $[\mathbf{a}_{(i,0)}, \mathbf{a}_{(i,1)}, \dots, \mathbf{a}_{(i,n-1)}]$ represents the i -th row of the batched random matrices, with each element $\mathbf{a}_{(i,j)} = \text{Ecd}(a_{(i,j)}^0, a_{(i,j)}^1, \dots, a_{(i,j)}^k, \dots, a_{(i,j)}^{\frac{N}{2}-1})$ for $0 \leq k < N/2$. After the RS operation, with a level consumption of one, the batched round constant is added to the state. By applying XBOOT to each FHE ciphertext, the round function can be iterated until the keystream is output.

With XBOOT, XOR operations are transformed into homomorphic additions, resulting in a noised integer in each slot of the FHE ciphertext, bounded by $\tau = (n+1) \times 3$ (where the result of each S layer is bounded by three). After applying XBOOT, these integers are converted into Boolean values, allowing other operations within the round function to proceed efficiently. This approach significantly improves performance by leveraging the SIMD capabilities of the CKKS scheme, with a bootstrapping layer consumption of $r+1$. Each bootstrapping layer contains $\lceil \frac{n}{2} \rceil$ calls to BatchXBOOT.

6 Experimental Results

We implemented XBOOT using the open-source homomorphic encryption library Lattigo V6 [lat24] and conducted experiments based on this framework. All experiments were performed on a server equipped with an Intel Xeon Platinum 8369B 64-Core 2.50 GHz machine with 128 GB RAM.

In the experiments, N denotes the ring degree of the CKKS polynomial. The symbols h and $\tilde{h}$ represent the Hamming weights of the dense and sparse secret keys, respectively,

as described in the sparse secret encapsulation from [BTH22]. We denote $\log(QP)$ as the total bit length of the modulus, and *depth per round* as the number of multiplication levels available after bootstrapping. All parameters used in the experiments ensure 128-bit security, as referenced in [APS15].

All transciphering operations are performed in counter (CTR) mode, which allows $N/2$ blocks of symmetric encryption to be processed simultaneously. For AES-CTR, this means that $N/2$ blocks can handle data typically amounting to $128 \times 2^{14} \approx 256$ KB.

6.1 FHE Parameters

Similar to [BCKS24], our XBOOT benefits from modulus savings due to the reduced gap between the base modulus q_0 and the scale Δ_0 . Moreover, the novel proposal of its free XOR functionality further reduces the depths of multiplication. As a result, we can decrease the ring degree N to 2^{15} with a total modulus $\log(QP) \leq 830$ -bit, ensuring 128-bit security. And we switch between different Hamming weights for bootstrapping as suggested in [BTH22]. We name each parameter set in the form of "Param-Cipher-log(BatchSize)", and they are listed in the following table.

Table 2: Concrete parameters for transciphering are as follows: The $\log(QP)$ column denotes the total modulus size. The $\log(q)$ columns represent the consumption of primes, with Base, StC, Mult, EvalMod, and CtS indicating the prime size and the number of primes consumed in each step. The $\log(p)$ column denotes the bit-size and the number of temporary moduli used for key switching.

	N	$(h, \tilde{h})$	$\log(QP)$	<i>depth per round</i>
Param-AES-13*	2^{14}	256, 16	433	3
Param-AES-14	2^{15}	256, 32	829	3
Param-Rasta-14	2^{15}		830	2

	$\log(q)$					$\log(p)$
	Base	StC	Mult	EvalMod	CtS	
Param-AES-13	32	28	31×3	32×6	28×2	32
Param-AES-14	38	36×2	37×3	38×7	38×4	38×5
Param-Rasta-14	42	41×2	41×2	42×7	40×3	42×5

* The significantly smaller $\log(qp) \approx 64$ -bit at the lowest level permits a lower Hamming weight of $\tilde{h} = 16$ while still achieving more than 128-bit security, according to [APS15].

6.2 AES Transciphering Results

We compare our implementation with the state-of-the-art works on AES transciphering in Table 3. Their numbers are directly taken from the corresponding papers.

As shown, the three CKKS-based methods achieve the least time per block in an amortized sense, which translates to the highest throughput. However, in terms of latency, several TFHE-based methods outperform the CKKS ones. This is because TFHE-based methods have a batch size of 1. In other words, they only have to deal with one AES block

at a time, whereas CKKS-based methods leverage the SIMD feature¹ to handle multiple AES blocks at a time. Specifically, the most advanced AES-TFHE technique [BBB⁺25] requires only 32 seconds per block per thread.

Compared to BGV-based method. The performance of BGV-based method [GHS12] falls between that of CKKS ones and TFHE ones: it exhibits moderate latency and time cost per block. Note that [GHS12] would have similar throughput as [ADE⁺23] if we divide their time cost by 64. However, FHE is not only computationally intensive but also incurs a significant memory footprint, which makes full parallelization challenging. Indeed, our 64-threaded implementation can only achieve a 20-30 times performance increase compared to our own single-threaded implementation.

Table 3: Comparison of AES transciphering efficiency. Following [ADE⁺23], we report the latency (the time required to transcipher a smallest batch of symmetric cipher) and the amortized time cost per block (calculated as latency/batch size). Each AES block contains 128 bits.

Scheme	Work	Latency	Time per block	Batch size
TFHE	BGV [GHS12]	240 s	2 s	120 blocks
	[SMK22]	252 s	252 s	
	[TCBS23]	270 s	270 s	
	[BPR24]	211 s	211 s	
	[WWL ⁺ 23]	86 s	86 s	1 block
	[WLW ⁺ 24]	110 s	110 s	
	[WLW ⁺ 24]	46 s	46 s	
	[SAP24]	61.9 s	61.9 s	
	[SAP24]*	0.95 s	0.95 s	
CKKS	[BBB ⁺ 25]	32 s	32 s	
	[ADE ⁺ 23]*	1200 s	37 ms	2 ¹⁵ blocks
	Param-AES-13*	153 s	18.7 ms	2 ¹³ blocks
	Param-AES-14*	236 s	14.4 ms	2 ¹⁴ blocks

* 64-threading is applied. The multi-threaded version of [SAP24] uses 16 threads while we use 64 threads, so we improve their performance by 4× for comparison.

Works without star symbol report their results using single thread.

Compared to BLEACH-based method. Among the three BLEACH-based methods, **Param-AES-13** exhibits lower latency but a higher amortized cost compared to **Param-AES-14** since its smaller parameter can accommodate only half as many AES blocks. But **Param-AES-14** achieves a 5× lower latency and a 2.5× reduction in amortized time cost than [ADE⁺23] under smaller parameters. This demonstrates the effectiveness of XBOOT.

Compared to CGGI-based method. A recent advancement, Hippogryph [BBB⁺25], achieves impressive per-block latency (e.g., 1s multi-threaded, 32s single-threaded) by strategically reducing TFHE bootstrapping operations. Its core techniques include: the

¹Some works such as [MS18, LW23] enable SIMD for TFHE-like schemes, but they are not concretely efficient.

Tree-Based Method (TBM) [TCBS23] for efficient S-box evaluation; Multi-LUT Programmable Bootstrapping (PBS) for post-S-box bit extraction; and p-encoding (with $p = 2$) [BPR24] to convert XORs into cheaper homomorphic additions.

Hippogryph’s low latency makes it highly suitable for applications involving a small number of AES blocks. For instance, it can process up to approximately 236 blocks faster than the total 236s transciphering time of XBOOT for 2^{14} blocks. Conversely, XBOOT’s CKKS-based architecture excels in amortized throughput for large data volumes due to its use of packing. This highlights a critical trade-off: TFHE-based methods like Hippogryph offer superior latency for low-volume tasks, whereas schemes like XBOOT provide higher throughput for large-scale processing. The choice between them is therefore dictated by an application’s specific scale and performance requirements.

6.3 Rasta Transciphering Results

Table 4: Comparison of Rasta-6 transciphering efficiency. Each block contains 351 bits.

Scheme	Work	Latency	Time per block	Batch size
BGV	[DEG ⁺ 18]	815 s	815 s	1 block
	[DGH ⁺ 23]	207 s	207 s	1 block
	[CHK ⁺ 21]	11980 s	16.6 s	720 blocks
TFHE	[DGH ⁺ 23]	3275 s	3275 s	1 block
CKKS	BLEACH*	> 16124 s	> 0.49 s	2^{15} blocks
	Param-Rasta-14*	303 s	18.5 ms	2^{14} blocks

* 64-threading is applied. Works without star symbol use single thread.

Transciphering for Rasta has been studied in several previous papers. We choose the Rasta-6 cipher and compare our implementation with their results in Table 4.

As shown, the BGV and TFHE methods require tens to thousands of seconds to process one block of Rasta cipher, significantly slower than our result of 18.5 milliseconds. Even if we grant them the advantage of perfect parallelization and divide their time cost by 64, the best-performing [CHK⁺21] is still at least 14 times slower than ours when evaluated by amortized time cost per block.

Since Rasta-CKKS transciphering using BLEACH without free XOR is extremely time-consuming, we chose not to implement it and instead estimated its cost to highlight the importance of free XOR. The number $16124 \approx 351 \times 350 \times 7 \times 1.2/64$ is explained as follows: Each evaluation of a batched random matrix requires 351×350 XOR operations. For the seven random matrices in Rasta-6, this translates to at least $351 \times 350 \times 7$ multiply-then-relinearize operations when using BLEACH. We only consider the cost of the KS (Key Switching) step in relinearization, which takes approximately 1.2 second according to [MBTPH20, Table 5]. Then we divide the cost by 64 to simulate multi-threading. The amortized time cost block is at least $16124/2^{15} \approx 0.49$ s. This is at least $26\times$ slower than ours because we treat the XORs as cheap additions.

6.4 Transciphering Accuracy

We assess the accuracy of AES-CKKS transciphering by measuring the average bias between decrypted results and the original binary values: $\text{AVG}(|\text{pt} - \text{Dec}_{\text{sk}}(\text{ct})|)$. According to [ADE⁺23], their average bias is reported at $\approx 1.16 \times 2^{-10}$. Our work achieves higher accuracy, leveraging optimizations from [BCKS24] that exploit the limited Boolean message

domain. For instance, the **Param-AES-13** configuration yields an average bias of $\approx 1.2 \times 2^{-13}$. **Param-AES-14**, which has larger modulus sizes and scaling factors (therefore better accuracy), demonstrates an average bias of $\approx 1.1 \times 2^{-19}$. For those requiring even higher precision, the META-BTS bootstrapping method [BCC⁺22] can be employed to further enhance accuracy.

Suitability for subsequent applications. To demonstrate XBOOT’s compatibility with downstream tasks, we conducted an experiment on privacy-preserving genome-wide association studies (GWAS) [BGP20]. Specifically, we applied the binary-to-real conversion algorithm [ADE⁺23, Algorithm 1] after AES-CKKS transciphering with XBOOT, then securely computed the χ^2 statistic for genetic associations to perform Pearson’s test of independence. The dataset comprises genetic data from 64 individuals, each with 16,384 Single-Nucleotide Polymorphisms (SNPs). Each SNP is encoded as $\{0, 1, 2\}$, representing the number of minor alleles at a locus. For more details, please refer to Appendix A. The transciphering process requires 236 seconds. The computation of the χ^2 statistic takes 2.9 seconds, with an accuracy of 17-bits.

7 Conclusion

In this paper, we study the problem of homomorphic computation on Boolean circuits using CKKS. We propose a novel method XBOOT that consumes zero multiplication depth per XOR gate instead of one. By achieving a shallower multiplication depth, we are able to enhance the AES transciphering efficiency to 256KB per 236 seconds, which is a $5\times$ latency improvement and $2.5\times$ throughput improvement over the state-of-the-art method. We also adopt our method to Rasta cipher to show the generality of XBOOT. As a result, we can better leverage the SIMD feature of CKKS homomorphic encryption and improve the throughput of Rasta transciphering by more than one order of magnitude.

References

- [ADE⁺23] Ehud Aharoni, Nir Drucker, Gilad Ezov, Eyal Kushnir, Hayim Shaul, and Omri Soceanu. E2e near-standard and practical authenticated transciphering. Cryptology ePrint Archive, Paper 2023/1040, 2023.
- [APS15] Martin R. Albrecht, Rachel Player, and Sam Scott. On the concrete hardness of learning with errors. *J. Math. Cryptol.*, 9(3):169–203, 2015.
- [ARS⁺15] Martin R. Albrecht, Christian Rechberger, Thomas Schneider, Tyge Tiessen, and Michael Zohner. Ciphers for MPC and FHE. In *Advances in Cryptology - EUROCRYPT 2015*, volume 9056 of *Lecture Notes in Computer Science*, pages 430–454. Springer, 2015.
- [BBB⁺25] Sonia Belaïd, Nicolas Bon, Aymen Boudguiga, Renaud Sirdey, Daphné Trama, and Nicolas Ye. Further improvements in aes execution over tfhe: Towards breaking the 1 sec barrier. *Cryptology ePrint Archive*, 2025.
- [BCC⁺22] Youngjin Bae, Jung Hee Cheon, Wonhee Cho, Jaehyung Kim, and Taekyung Kim. META-BTS: bootstrapping precision beyond the limit. In *Proceedings of the 2022 ACM SIGSAC Conference on Computer and Communications Security, CCS 2022, Los Angeles, CA, USA, November 7-11, 2022*, pages 223–234. ACM, 2022.

- [BCKS24] Youngjin Bae, Jung Hee Cheon, Jaehyung Kim, and Damien Stehlé. Bootstrapping bits with CKKS. In *Advances in Cryptology - EUROCRYPT 2024*, volume 14652 of *Lecture Notes in Computer Science*, pages 94–123. Springer, 2024.
- [BGPG20] Marcelo Blatt, Alexander Gusev, Yuriy Polyakov, and Shafi Goldwasser. Secure large-scale genome-wide association studies using homomorphic encryption. *Proceedings of the National Academy of Sciences*, 117(21):11608–11613, 2020.
- [BGV14] Zvika Brakerski, Craig Gentry, and Vinod Vaikuntanathan. (leveled) fully homomorphic encryption without bootstrapping. *ACM Trans. Comput. Theory*, 6(3):13:1–13:36, 2014.
- [BKSS24] Youngjin Bae, Jaehyung Kim, Damien Stehlé, and Elias Suvanto. Bootstrapping small integers with CKKS. In Kai-Min Chung and Yu Sasaki, editors, *Advances in Cryptology - ASIACRYPT 2024*, volume 15484 of *Lecture Notes in Computer Science*, pages 330–360. Springer, 2024.
- [BMTH21] Jean-Philippe Bossuat, Christian Mouchet, Juan Ramón Troncoso-Pastoriza, and Jean-Pierre Hubaux. Efficient bootstrapping for approximate homomorphic encryption with non-sparse keys. In Anne Canteaut and François-Xavier Standaert, editors, *Advances in Cryptology - EUROCRYPT 2021*, volume 12696 of *Lecture Notes in Computer Science*, pages 587–617. Springer, 2021.
- [BPR24] Nicolas Bon, David Pointcheval, and Matthieu Rivain. Optimized homomorphic evaluation of boolean functions. *IACR Trans. Cryptogr. Hardw. Embed. Syst.*, 2024(3):302–341, 2024.
- [BTH22] Jean-Philippe Bossuat, Juan Ramón Troncoso-Pastoriza, and Jean-Pierre Hubaux. Bootstrapping for approximate homomorphic encryption with negligible failure-probability by using sparse-secret encapsulation. In Giuseppe Ateniese and Daniele Venturi, editors, *Applied Cryptography and Network Security - 20th International Conference, ACNS 2022, Rome, Italy, June 20-23, 2022, Proceedings*, volume 13269 of *Lecture Notes in Computer Science*, pages 521–541. Springer, 2022.
- [BV11] Zvika Brakerski and Vinod Vaikuntanathan. Fully homomorphic encryption from ring-lwe and security for key dependent messages. In Phillip Rogaway, editor, *Advances in Cryptology - CRYPTO 2011*, volume 6841 of *Lecture Notes in Computer Science*, pages 505–524. Springer, 2011.
- [BV14] Zvika Brakerski and Vinod Vaikuntanathan. Efficient fully homomorphic encryption from (standard) lwe. *SIAM J. Comput.*, 43(2):831–871, 2014.
- [CCF⁺16] Anne Canteaut, Sergiu Carpov, Caroline Fontaine, Tancrede Lepoint, María Naya-Plasencia, Pascal Paillier, and Renaud Sirdey. Stream ciphers: A practical solution for efficient homomorphic-ciphertext compression. In Thomas Peyrin, editor, *Fast Software Encryption - 23rd International Conference, FSE 2016, Bochum, Germany, March 20-23, 2016, Revised Selected Papers*, volume 9783 of *Lecture Notes in Computer Science*, pages 313–333. Springer, 2016.
- [CCK⁺13] Jung Hee Cheon, Jean-Sébastien Coron, Jinsu Kim, Moon Sung Lee, Tancrede Lepoint, Mehdi Tibouchi, and Aaram Yun. Batch fully homomorphic encryption over the integers. In Thomas Johansson and Phong Q. Nguyen, editors,

- Advances in Cryptology - EUROCRYPT 2013, 32nd Annual International Conference on the Theory and Applications of Cryptographic Techniques, Athens, Greece, May 26-30, 2013. Proceedings*, volume 7881 of *Lecture Notes in Computer Science*, pages 315–335. Springer, 2013.
- [CCS19] Hao Chen, Ilaria Chillotti, and Yongsoo Song. Improved bootstrapping for approximate homomorphic encryption. In Yuval Ishai and Vincent Rijmen, editors, *Advances in Cryptology - EUROCRYPT 2019*, volume 11477 of *Lecture Notes in Computer Science*, pages 34–54. Springer, 2019.
- [CGGI17] Ilaria Chillotti, Nicolas Gama, Mariya Georgieva, and Malika Izabachène. Faster packed homomorphic operations and efficient circuit bootstrapping for TFHE. In Tsuyoshi Takagi and Thomas Peyrin, editors, *Advances in Cryptology - ASIACRYPT 2017*, volume 10624 of *Lecture Notes in Computer Science*, pages 377–408. Springer, 2017.
- [CGGI20] Ilaria Chillotti, Nicolas Gama, Mariya Georgieva, and Malika Izabachène. TFHE: fast fully homomorphic encryption over the torus. *J. Cryptol.*, 33(1):34–91, 2020.
- [CHK⁺18] Jung Hee Cheon, Kyoohyung Han, Andrey Kim, Miran Kim, and Yongsoo Song. Bootstrapping for approximate homomorphic encryption. In Jesper Buus Nielsen and Vincent Rijmen, editors, *Advances in Cryptology - EUROCRYPT 2018*, volume 10820 of *Lecture Notes in Computer Science*, pages 360–384. Springer, 2018.
- [CHK⁺21] Jihoon Cho, Jincheol Ha, Seongkwang Kim, ByeongHak Lee, Joohee Lee, Jooyoung Lee, Dukjae Moon, and Hyojin Yoon. Transciphering framework for approximate homomorphic encryption. In Mehdi Tibouchi and Huaxiong Wang, editors, *Advances in Cryptology - ASIACRYPT 2021*, volume 13092 of *Lecture Notes in Computer Science*, pages 640–669. Springer, 2021.
- [CIR22] Carlos Cid, John Petter Indrøy, and Håvard Raddum. FASTA - A stream cipher for fast FHE evaluation. In Steven D. Galbraith, editor, *Topics in Cryptology - CT-RSA 2022 - Cryptographers’ Track at the RSA Conference 2022, Virtual Event, March 1-2, 2022, Proceedings*, volume 13161 of *Lecture Notes in Computer Science*, pages 451–483. Springer, 2022.
- [CKK20] Jung Hee Cheon, Dongwoo Kim, and Duhyeong Kim. Efficient homomorphic comparison methods with optimal complexity. In Shiho Moriai and Huaxiong Wang, editors, *Advances in Cryptology - ASIACRYPT 2020*, volume 12492 of *Lecture Notes in Computer Science*, pages 221–256. Springer, 2020.
- [CKKS17] Jung Hee Cheon, Andrey Kim, Miran Kim, and Yong Soo Song. Homomorphic encryption for arithmetic of approximate numbers. In Tsuyoshi Takagi and Thomas Peyrin, editors, *Advances in Cryptology - ASIACRYPT 2017*, volume 10624 of *Lecture Notes in Computer Science*, pages 409–437. Springer, 2017.
- [CLW⁺24] Chunling Chen, Xianhui Lu, Ruida Wang, Zhihao Li, Xuan Shen, and Benqiang Wei. Free-xor gate bootstrapping. *IACR Cryptol. ePrint Arch.*, page 1703, 2024.
- [DEG⁺18] Christoph Dobraunig, Maria Eichlseder, Lorenzo Grassi, Virginie Lallemand, Gregor Leander, Eik List, Florian Mendel, and Christian Rechberger. Rasta: A cipher with low anddepth and few ands per bit. In Hovav Shacham and Alexandra Boldyreva, editors, *Advances in Cryptology - CRYPTO 2018*,

- volume 10991 of *Lecture Notes in Computer Science*, pages 662–692. Springer, 2018.
- [DEM15] Christoph Dobraunig, Maria Eichlseder, and Florian Mendel. Higher-order cryptanalysis of lowmc. In Soonhak Kwon and Aaram Yun, editors, *Information Security and Cryptology - ICISC 2015*, volume 9558 of *Lecture Notes in Computer Science*, pages 87–101. Springer, 2015.
 - [DGH⁺23] Christoph Dobraunig, Lorenzo Grassi, Lukas Helming, Christian Rechberger, Markus Schafnegg, and Roman Walch. Pasta: A case for hybrid homomorphic encryption. *IACR Trans. Cryptogr. Hardw. Embed. Syst.*, 2023(3):30–73, 2023.
 - [DM15] Léo Ducas and Daniele Micciancio. FHEW: bootstrapping homomorphic encryption in less than a second. In Elisabeth Oswald and Marc Fischlin, editors, *Advances in Cryptology - EUROCRYPT 2015*, volume 9056 of *Lecture Notes in Computer Science*, pages 617–640. Springer, 2015.
 - [DMPS24] Nir Drucker, Guy Moshkovich, Tomer Pelleg, and Hayim Shaul. BLEACH: cleaning errors in discrete computations over CKKS. *J. Cryptol.*, 37(1):3, 2024.
 - [FRL19] Hayato Fujii, Félix Carvalho Rodrigues, and Julio López. Fast AES implementation using armv8 ASIMD without cryptography extension. In Jae Hong Seo, editor, *Information Security and Cryptology - ICISC 2019*, volume 11975 of *Lecture Notes in Computer Science*, pages 84–101. Springer, 2019.
 - [FV12] Junfeng Fan and Frederik Vercauteren. Somewhat practical fully homomorphic encryption. *IACR Cryptol. ePrint Arch.*, page 144, 2012.
 - [GAH⁺23] Lorenzo Grassi, Irati Manterola Ayala, Martha Norberg Hovd, Morten Øygarden, Håvard Raddum, and Qingju Wang. Cryptanalysis of symmetric primitives over rings and a key recovery attack on rubato. In Helena Handschuh and Anna Lysyanskaya, editors, *Advances in Cryptology - CRYPTO 2023*, volume 14083 of *Lecture Notes in Computer Science*, pages 305–339. Springer, 2023.
 - [Gen09] Craig Gentry. Fully homomorphic encryption using ideal lattices. In Michael Mitzenmacher, editor, *Proceedings of the 41st Annual ACM Symposium on Theory of Computing, STOC 2009, Bethesda, MD, USA, May 31 - June 2, 2009*, pages 169–178. ACM, 2009.
 - [GHS12] Craig Gentry, Shai Halevi, and Nigel P. Smart. Homomorphic evaluation of the AES circuit. In Reihaneh Safavi-Naini and Ran Canetti, editors, *Advances in Cryptology - CRYPTO 2012*, volume 7417 of *Lecture Notes in Computer Science*, pages 850–867. Springer, 2012.
 - [GSW13] Craig Gentry, Amit Sahai, and Brent Waters. Homomorphic encryption from learning with errors: Conceptually-simpler, asymptotically-faster, attribute-based. In Ran Canetti and Juan A. Garay, editors, *Advances in Cryptology - CRYPTO 2013*, volume 8042 of *Lecture Notes in Computer Science*, pages 75–92. Springer, 2013.
 - [HK20] Kyoohyung Han and Dohyeong Ki. Better bootstrapping for approximate homomorphic encryption. In Stanislaw Jarecki, editor, *Topics in Cryptology - CT-RSA 2020 - The Cryptographers’ Track at the RSA Conference 2020, San Francisco, CA, USA, February 24–28, 2020, Proceedings*, volume 12006 of *Lecture Notes in Computer Science*, pages 364–390. Springer, 2020.

- [HL20] Phil Hebborn and Gregor Leander. Dasta - alternative linear layer for rasta. *IACR Trans. Symmetric Cryptol.*, 2020(3):46–86, 2020.
- [JM22] Charanjit S. Jutla and Nathan Manohar. Sine series approximation of the mod function for bootstrapping of approximate HE. In Orr Dunkelman and Stefan Dziembowski, editors, *Advances in Cryptology - EUROCRYPT 2022*, volume 13275 of *Lecture Notes in Computer Science*, pages 491–520. Springer, 2022.
- [JPK⁺24] Jae Hyung Ju, Jaiyoung Park, Jongmin Kim, Minsik Kang, Donghwan Kim, Jung Hee Cheon, and Jung Ho Ahn. Neujeans: Private neural network inference with joint optimization of convolution and FHE bootstrapping. In Bo Luo, Xiaojing Liao, Jun Xu, Engin Kirda, and David Lie, editors, *Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security, CCS 2024, Salt Lake City, UT, USA, October 14-18, 2024*, pages 4361–4375. ACM, 2024.
- [KN24] Jaehyung Kim and Taeyeong Noh. Modular reduction in CKKS. *IACR Cryptol. ePrint Arch.*, page 1638, 2024.
- [KSS12] Benjamin Kreuter, Abhi Shelat, and Chih-Hao Shen. Billion-gate secure computation with malicious adversaries. In Tadayoshi Kohno, editor, *Proceedings of the 21th USENIX Security Symposium, Bellevue, WA, USA, August 8-10, 2012*, pages 285–300. USENIX Association, 2012.
- [lat24] Lattigo v6. Online: <https://github.com/tuneinsight/lattigo>, August 2024. EPFL-LDS, Tune Insight SA.
- [LKSM24] Fukang Liu, Abul Kalam, Santanu Sarkar, and Willi Meier. Algebraic attack on fhe-friendly cipher HERA using multiple collisions. *IACR Trans. Symmetric Cryptol.*, 2024(1):214–233, 2024.
- [LLK⁺22] Yongwoo Lee, Joon-Woo Lee, Young-Sik Kim, Yongjune Kim, Jong-Seon No, and HyungChul Kang. High-precision bootstrapping for approximate homomorphic encryption by error variance minimization. In Orr Dunkelman and Stefan Dziembowski, editors, *Advances in Cryptology - EUROCRYPT 2022*, volume 13275 of *Lecture Notes in Computer Science*, pages 551–580. Springer, 2022.
- [LLKN23] Joon-Woo Lee, Eunsang Lee, Young-Sik Kim, and Jong-Seon No. Rotation key reduction for client-server systems of deep neural network on fully homomorphic encryption. In Jian Guo and Ron Steinfeld, editors, *Advances in Cryptology - ASIACRYPT 2023*, volume 14443 of *Lecture Notes in Computer Science*, pages 36–68. Springer, 2023.
- [LLL⁺21] Joon-Woo Lee, Eunsang Lee, Yongwoo Lee, Young-Sik Kim, and Jong-Seon No. High-precision bootstrapping of RNS-CKKS homomorphic encryption using optimal minimax polynomial approximation and inverse sine function. In Anne Canteaut and François-Xavier Standaert, editors, *Advances in Cryptology - EUROCRYPT 2021*, volume 12696 of *Lecture Notes in Computer Science*, pages 618–647. Springer, 2021.
- [LLL⁺22] Eunsang Lee, Joon-Woo Lee, Junghyun Lee, Young-Sik Kim, Yongjune Kim, Jong-Seon No, and Woosuk Choi. Low-complexity deep convolutional neural networks on fully homomorphic encryption using multiplexed parallel convolutions. In Kamalika Chaudhuri, Stefanie Jegelka, Le Song, Csaba

- Szepesvári, Gang Niu, and Sivan Sabato, editors, *International Conference on Machine Learning, ICML 2022, 17-23 July 2022, Baltimore, Maryland, USA*, volume 162 of *Proceedings of Machine Learning Research*, pages 12403–12422. PMLR, 2022.
- [LW23] Feng-Hao Liu and Han Wang. Batch bootstrapping I: - A new framework for SIMD bootstrapping in polynomial modulus. In Carmit Hazay and Martijn Stam, editors, *Advances in Cryptology - EUROCRYPT 2023 - 42nd Annual International Conference on the Theory and Applications of Cryptographic Techniques, Lyon, France, April 23-27, 2023, Proceedings, Part III*, volume 14006 of *Lecture Notes in Computer Science*, pages 321–352. Springer, 2023.
- [MBTPH20] Christian Vincent Mouchet, Jean-Philippe Bossuat, Juan Ramón Troncoso-Pastoriza, and Jean-Pierre Hubaux. Lattigo: A multiparty homomorphic encryption library in go. In *Proceedings of the 8th Workshop on Encrypted Computing and Applied Homomorphic Cryptography*, pages 64–70, 2020.
- [MPR20] Mahesh Kumar Morampudi, Munaga V. N. K. Prasad, and U. S. N. Raju. Privacy-preserving iris authentication using fully homomorphic encryption. *Multim. Tools Appl.*, 79(27-28):19215–19237, 2020.
- [MS18] Daniele Micciancio and Jessica Sorrell. Ring packing and amortized FHEW bootstrapping. In Ioannis Chatzigiannakis, Christos Kaklamanis, Daniel Marx, and Donald Sannella, editors, *45th International Colloquium on Automata, Languages, and Programming, ICALP 2018, July 9-13, 2018, Prague, Czech Republic*, volume 107 of *LIPIcs*, pages 100:1–100:14. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 2018.
- [NWH⁺25] Chao Niu, Benqiang Wei, Zhicong Huang, Zhaomin Yang, Cheng Hong, Meiqin Wang, and Tao Wei. Sok: Fhe-friendly symmetric ciphers and transciphering. *IACR Transactions on Cryptographic Hardware and Embedded Systems*, 2025(3):583–613, Jun. 2025.
- [PF15] NIST FIPS PUB and PUB FIPS. 202: Sha-3 standard: Permutation-based hash and extendable-output functions. *Federal Information Processing Standards Publication*, 202, 2015.
- [SAP24] Leonard Schild, Aysajan Abidin, and Bart Preneel. Fast transciphering via batched and reconfigurable LUT evaluation. *IACR Trans. Cryptogr. Hardw. Embed. Syst.*, 2024(4):205–230, 2024.
- [SMK22] Roy Stracovsky, Rasoul Akhavan Mahdavi, and Florian Kerschbaum. Faster evaluation of aes using tfhe. Poster Session, FHE.Org - 2022, 2022. Available at <https://rasoulam.github.io/data/poster-aes-tfhe.pdf>.
- [TCBS23] Daphné Trama, Pierre-Emmanuel Clet, Aymen Boudguiga, and Renaud Sirdey. At last! A homomorphic AES evaluation in less than 30 seconds by means of TFHE. *IACR Cryptol. ePrint Arch.*, page 1020, 2023.
- [WLW⁺24] Benqiang Wei, Xianhui Lu, Ruida Wang, Kun Liu, Zhihao Li, and Kunpeng Wang. Thunderbird: Efficient homomorphic evaluation of symmetric ciphers in 3gpp by combining two modes of TFHE. *IACR Trans. Cryptogr. Hardw. Embed. Syst.*, 2024(3):530–573, 2024.
- [WWL⁺23] Benqiang Wei, Ruida Wang, Zhihao Li, Qingu Liu, and Xianhui Lu. Fregata: Faster homomorphic evaluation of AES via TFHE. In Elias Athanasopoulos and Bart Mennink, editors, *Information Security - 26th International*

Conference, ISC 2023, Groningen, The Netherlands, November 15-17, 2023, Proceedings, volume 14411 of *Lecture Notes in Computer Science*, pages 392–412. Springer, 2023.

- [WWL⁺24] Ruida Wang, Yundi Wen, Zhihao Li, Xianhui Lu, Benqiang Wei, Kun Liu, and Kunpeng Wang. Circuit bootstrapping: Faster and smaller. In Marc Joye and Gregor Leander, editors, *Advances in Cryptology - EUROCRYPT 2024*, volume 14652 of *Lecture Notes in Computer Science*, pages 342–372. Springer, 2024.

A Pearson’s Test of Independence

First, the individual disease status ($y_i \in \{0, 1\}$) and allele counts ($s_i \in \{0, 1, 2\}$) are encrypted bit by bit using AES and transcribed to CKKS using XBOOT. Then the binary-to-real conversion in Algorithm 6 from [ADE⁺23] is applied to obtain y_i and s_i . The server then homomorphically computes $A = 2N(2n_{11}N - c_1r_1)^2$ and $B = c_1(2N - c_1)r_1 \cdot (2N - r_1)$, and return A, B to the client, where ($n_{11} = \mathbf{y}^\top \cdot \mathbf{s}$, $r_1 = 2\mathbf{y}^\top \cdot \mathbf{y}$, $c_1 = \sum_{i=1}^N s_i$), and N is cohort size. Finally, the client decrypts and compute $\chi^2 = \frac{A}{B}$. Then, we have the equation of the Pearson’s test of independence:

$$\chi^2 = \frac{2N(2n_{11}N - c_1r_1)^2}{c_1(2N - c_1)r_1 \cdot (2N - r_1)}.$$

Additional information on privacy-preserving genome-wide association studies is available in [BGP20].

Algorithm 6: Constructing numbers from bits

Input : An array in of n encrypted bits, $\{in_0, \dots, in_{n-1}\}$.
A cryptographic noise bound e .

Output : An integer $out \approx \sum_{i=0}^{n-1} 2^i \cdot in_i$, with total noise below e .

```

1 procedure ConstructInt
2    $out \leftarrow in_0 + 2 \cdot in_1$ ;
3   for  $i \leftarrow 2$  to  $n - 1$  do
4     // The term  $2^i \cdot in_i$  is computed below.
5     /* The choice of method depends on the index  $i$  to keep the cryptographic
6        noise from exceeding the bound  $e$ . Both branches compute the same
7        value. */
8     if  $i < -\frac{\log_2(e)}{2}$  then
9        $b \leftarrow (2^{\lfloor i/2 \rfloor} \cdot in_i)^2$ ;
10      if  $i \pmod 2 = 1$  then
11         $b \leftarrow 2b$ ;
12      else
13         $b \leftarrow (2^{\lfloor i/4 \rfloor} \cdot in_i)^4$ ;
14         $b \leftarrow 2^{i - (4 \lfloor i/4 \rfloor)} \cdot b$ ;
15       $out \leftarrow out + b$ ;
16 return  $out$ ;
```
