

Forking Lemma in EasyCrypt

Denis Firsov
Tallinn University of Technology
Input Output
Tallinn, Estonia
denis.firsov@taltech.ee

Jakub Janků
Masaryk University
Brno, Czech Republic
jjanku@mail.muni.cz

Abstract—Formal methods are becoming an important tool for ensuring correctness and security of cryptographic constructions. However, the support for certain advanced proof techniques, namely rewinding, is scarce among existing verification frameworks, which hinders their application to complex schemes such as multi-party signatures and zero-knowledge proofs.

We expand the support for rewinding in EasyCrypt by implementing a version of the general forking lemma by Bellare and Neven. We demonstrate its usability by proving EUF-CMA security of Schnorr signatures.

Index Terms—computer-aided cryptography, formal verification, EasyCrypt, rewinding, forking, Schnorr

I. INTRODUCTION

Although provable security has gradually become a de facto standard in cryptography, it is long known that it is not a silver bullet that can rule out all attacks [2]. One of the leading problems is that security proofs are difficult to read and verify. As a result, they do not receive sufficient attention during the review process and subtle errors may stay unnoticed.

Several methods of structuring security proofs, such as the game-based technique [3], have been developed to reduce the effort required to both create and check such proofs, yet human verification will always be error-prone. This idea led to the development of *computer-aided cryptography*, which strives to apply formal verification techniques to security [4].

Despite the recent successes in this field — e.g., fixing and mechanizing security proofs for new post-quantum signature schemes, Dilithium [5] and XMSS [6] — there remain some classical results that seem to evade formal verification. Unforgeability of Schnorr signatures is one of them. This contrasts with the renewed interest in this construction and the number of novel Schnorr-based schemes (such as EdDSA, MuSig [7], and FROST [8]), particularly in the active area (as witnessed by the NIST call for proposals [9]) of multi-party threshold cryptography.

The likely reason for this discrepancy is that security reductions for Schnorr-like schemes commonly rely on *forking* of the hypothetical adversary — a special case of *rewinding* which is a technique that was not implemented by any cryptographic verification framework until recently [10].

In this work, we build on the recent advancements in formalization of rewinding and address the described gap by implementing a *general forking lemma* in EasyCrypt and reproducing the well-known result on the security of simple Schnorr signatures. One of the main challenges we solve in the

process is how to rewind an adversary at an oracle query given the restrictions imposed by the EasyCrypt model of programs and oracle access. We hope that our work will pave the way for applying formal methods to more advanced constructions that previously appeared beyond reach — including some threshold schemes and zero-knowledge proofs (see Section V-B for more details). All EasyCrypt scripts are publicly available on GitHub¹.

II. RELATED WORK

Zain formally verified three signature schemes related to Schnorr using EasyCrypt [11] — EDL, CM, and GJKW. As the author explains, one of the reasons behind this choice was that all constructions feature a tight security reduction and do not require the forking lemma.

Several groups of authors formalized properties of zero-knowledge Σ -protocols and applied the results to the *interactive* Schnorr identification protocol; various verification frameworks were employed, e.g., CryptHOL [12], CertiCrypt [13], or EasyCrypt [14]. As far as we know, none of these works covered the properties of *non-interactive* protocols, or signature schemes obtained from interactive ones via the Fiat-Shamir transformation.

The rewinding technique was first implemented in EasyCrypt by Firsov and Unruh [10]. As this is the starting point of our work, we comment on it in more detail in Section III-B.

A partially overlapping work on forking is being carried out by Tuma and Hopper [15]. They develop an entirely new cryptographic verification framework, called VCVio, which they embed in the Lean programming language / theorem prover. Its focus is on design-level computational security (see the survey of computer-aided cryptography [4]) with emphasis on reasoning about oracles. To demonstrate its power, the authors use it to prove a variant of the forking lemma by Bellare and Neven [1]. This result was the subject of our formalization as well, however, our approaches differ in multiple ways.

The main distinction lies in the modeling of probabilistic programs and rewinding. In VCVio, computations consist of a sequence of syntactic oracle calls and pure continuation functions which handle the oracle responses (technically, this is captured by a free monad). Non-determinism is achieved by using a coin flipping oracle. This construction allows one

¹<https://github.com/jjanku/fsec>.

to very easily fix the randomness of a program, simply by swapping the oracle for one that answers using a pre-generated random tape (seed, in VCVio terminology). To rewind a program, one can run it again with the same random tape. However, this forces one to bound the maximum length of the tape.

In contrast, EasyCrypt programs are stateful and probabilistic by nature and no such randomness extraction is possible. Instead, we implement rewinding by state serialization. This model is more general, as the amount of randomness does not need to be bounded. However, the trade-off is that programs must adhere to a specific interface — we essentially require the user to split the computation into continuations manually — so that we can save the state at any oracle query. More detailed discussion of our rewinding technique and program model can be found in Sections III-B and IV-B.

In both projects, it was crucial to choose a suitable model of the interaction of a probabilistic program with an oracle that is flexible and comfortable enough to express and prove the forking lemma in formal terms. Tackling this problem seems essential for expanding the support for forking to more verification tools.

Another difference is the proof technique. VCVio’s proof relies on the denotational semantics of programs and the ability to reason about the probability distributions that each computation induces. Our proof is partly game-based and utilizes Hoare logics, as this is the primary mechanism of EasyCrypt to derive facts about programs. Moreover, we reduce the forking lemma to the rewinding lemma, which makes the proof more modular.

Regarding finer details, we note that our result fully reproduces the probability bound by Bellare and Neven [1] and is tighter roughly by a factor of Q , the number of random oracle queries the forked algorithm makes. This stems from the fact that the authors of VCVio guess the forking point before running the algorithm for the first time, whereas we determine the forking point based on the result of the first execution².

The applicability of the results is difficult to compare. We use our forking lemma to prove unforgeability of Schnorr signatures. VCVio claims to show unforgeability of signatures constructed by the Fiat-Shamir transformation of Σ -protocols with the necessary properties (Schnorr is an example thereof). However, upon inspection of the code base, we found no such proof at the time of writing.

Finally, let us comment on the general applicability of the two used frameworks. While the defining features of VCVio could prove useful in some settings, at the moment, EasyCrypt undoubtedly has an advantage in terms of the number of generic modular definitions and results available, which makes it easier to apply to new schemes. Moreover, the Hoare logics in EasyCrypt (discussed in Section III-A) seem to be a key ingredient for taming more complex proofs; given the novelty,

it is again unclear how VCVio will fare in this direction without them. Also importantly, existing infrastructure allows us to transfer security guarantees from EasyCrypt onto an optimized implementation in Jasmin (see [16] and Section V-B); VCVio lacks a similar feature.

To the best of our knowledge, there is no other formalization of the forking lemma and the security analysis of Schnorr signatures (in EasyCrypt or in other verification frameworks).

III. PRELIMINARIES

First, we introduce the basics of EasyCrypt. Then, we recall the rewinding technique and summarize the existing formalization this work builds upon.

A. EasyCrypt

EasyCrypt is an interactive proof assistant tailored for formal verification of cryptographic proofs. It consists of three main components: a general-purpose *expression language*, a *module language* for specifying security games, and a set of *logics* for reasoning about them. We briefly illustrate the key concepts of each component on the following code snippet:

```

type log_t = bool list.

op dbool : bool distr.

module type Adv = {
  proc guess() : bool
}.

module B : Adv = {
  var log : log_t

  proc guess() : bool = {
    var b : bool;
    b ← $ {0,1}; (* uniform *)
    log ← log || [b];
    return b;
  }
}.

module Game(A : Adv) = {
  var won : bool

  proc play() = {
    var b0, b1 : bool;
    b0 ← dbool;
    b1 ← A.guess();
    won ← b0 = b1;
  }
}.

```

a) *Expression Language*: This is a typed functional language. Types and operators are built-in (`bool`, `=`) or user-defined (`log_t`, `dbool`) and can be abstract (`dbool`) or concrete (`log_t`). From the built-in type constructors, we point out `distr` which creates a type of discrete probability (sub)distributions over the supplied type (`dbool` is a distribution over `bool`). For readability, operators that take no arguments or return a Boolean can be defined as *constants* or *predicates*, respectively. Evaluation of all EasyCrypt expressions is guaranteed to terminate, operators thus correspond to total functions.

²However, there is a newer *unfinished* proof that obtains the same bound as we do.

b) *Module Language*: Unlike the expression language, the module language is imperative. There are standard statements such as assignments, procedure calls, or `while` loops. Additionally, one can assign a value at random from a given distribution using $\leftarrow \$$. Statements must be grouped into procedures, which must belong to modules.

Modules can be standalone (B) or parametrized (Game) by an abstract module (A), typically an oracle or an adversary, of a specified type (Adv). Module types are akin to interfaces that declare the procedures that the module must implement. A trivial method of interface inheritance is available using the `include` keyword.

Apart from procedures, modules may include global variables. These are not private and may be accessed by other modules. Accordingly, we define the set `glob M` of global variables³ of a module M, as the union of global variables declared inside M and the set of variables M may read / write during a call to one of its procedures (i.e., `glob Game(B)` contains `B.log`).

An execution of a module's procedure is defined relative to a memory *m*, which defines the contents of all module's global variables. The state of a variable in a memory can be queried using curly braces (e.g., `Game.won{m}` or `(glob B){m}`). To compare states across two memories 1 and 2, a shorthand is available (`={B.log}` is equivalent to `B.log{1} = B.log{2}`).

c) *Logics*: EasyCrypt features three basic logics to reason about module properties: *classical* Hoare Logic (HL), *probabilistic* HL (pHL), and *probabilistic relational* HL (pRHL). Judgments in the given logic are denoted by the keyword `hoare`, `phoare`, and `equiv` respectively and take the following common form: `procedure(s) : precondition  $\Rightarrow$  postcondition`. Here, the conditions may refer to global variables; moreover, the keywords `arg` and `res` reference arguments and return values of the procedure(s) in question. (For pRHL, a side / memory specifier, {1} or {2}, is necessary to determine the procedure execution we are interested in. For pHL, a probability bound is additionally needed.)

To prove general mathematical facts (about operators) and connect judgments from the specialized logics (about modules), the so-called *ambient* logic is provided. Using it, we can state `axioms` (which are assumed to be true) and prove `lemmas`. Both can be parametrized (also by memories).

Formulas in the ambient logic can also contain *probability expressions*: `Pr[procedure(args) @ m : event]`. The semantics is the probability of event when executing procedure with the provided args from initial memory *m*. In event, the `res` keyword is again permitted.

Note that for clarity, we slightly deviate from the EasyCrypt syntax. In particular, we omit some types and typecasts (%r), use ligatures ($\wedge$ instead of /\), replace some operators ($\leftarrow$ instead of <@), simplify memory notation (italicized *m* instead

of &m), etc. Readers familiar with EasyCrypt should have no trouble mapping our notation to the original one.

More information on EasyCrypt can be found in the EasyCrypt reference manual [17]. All limitations of EasyCrypt that are discussed in this paper apply to version 2026.02 (the latest stable release available to us at the time of writing).

B. Rewinding

Rewinding is a common cryptographic proof technique that generally proceeds as follows:

- 1) run a given program / adversary $\mathcal{A}$ for some time (typically until it requires some input from the environment),
- 2) save $\mathcal{A}$'s state,
- 3) finish the execution of $\mathcal{A}$,
- 4) restore $\mathcal{A}$'s state,
- 5) run $\mathcal{A}$ again till completion;
- 6) (optionally repeat steps 4-5).

The idea behind rewinding is that we can extract some useful information from the multiple related outputs of $\mathcal{A}$ or (eventually) obtain an output that is otherwise suitable for us. Importantly, since the adversary has no knowledge of being rewound, its success probability can be constant over all individual runs.

The above paradigm can be applied to the forking lemma as well. For this reason, we based our work on the existing EasyCrypt formalization of rewinding by Firsov and Unruh [10]. Its short description follows.

To start, the authors first had to define what it means for a program to be *rewindable*. In pen-and-paper proofs, this capability is almost never introduced formally. In EasyCrypt, however, it is necessary to use the module types to explicitly require that a given module provides a rewinding interface:

```
type state_t.

module type Rewindable = {
  proc getState() : state_t
  proc setState(st : state_t) : unit
}.
```

Moreover, one has to axiomatize that an implementation of the above interface behaves as the procedure names suggest. Formally, a Rewindable module R should satisfy the following properties:

- 1) there exists some injective function *f* from `glob R` to `state_t`;
- 2) the procedure `getState` terminates, has no side-effects, and returns *f gR* when invoked in state *gR* : `glob R`;
- 3) the procedure `setState` terminates and sets the state of R to *gR* when called with an argument *st* such that *st* = *f gR*.

In the rest of this paper, all Rewindable modules are expected to satisfy the listed conditions; we omit them for brevity in the formulation of our lemmas. We also note that this assumption does not weaken our results; see the discussion in the original work [10].

³Technically, this is a type.

The final step is to fix the point at which the module R should be rewound. To model this, the authors additionally assume that the execution of R is divided into two procedures, `init` and `run`. This is necessary because in EasyCrypt procedures are always executed in one go. Technically, another module type, here called `FixedRewindable`, is introduced that inherits from `Rewindable`.

Once the rewinding interface is specified, it is possible to straightforwardly translate the initial informal description of rewinding into an EasyCrypt module:

```

module Rewinder(R : FixedRewindable) = {
  (* Two runs of R, with rewinding. *)
  proc run2(i) = {
    var st, r0, r1, r2;
    r0 ← R.init(i);
    st ← R.getState();
    r1 ← R.run(r0);
    R.setState(st);
    r2 ← R.run(r0);
    return ((r0, r1), (r0, r2));
  }

  (* Single run of R. *)
  proc run(i) = {
    var r0, r1;
    r0 ← R.init(i);
    r1 ← R.run(r0);
    return (r0, r1);
  }
}

```

Listing 1. Rewinder module.

A key result of Firsov and Unruh about the Rewinder module that we will utilize in Section IV-D intuitively states the following: if R succeeds once (its output satisfies some predicate P), then it should also succeed the second time after rewinding⁴. Formally:

```

lemma rewinding_lemma m P i :
  Pr[Rewinder(R).run2(i) @ m : P res.1 ∧ P res.2] ≥
  Pr[Rewinder(R).run(i) @ m : P res] ^ 2

```

Listing 2. Rewinding lemma.

The proof of this lemma relies on Jensen’s inequality and *probabilistic reflection*, the ability to reason about denotational semantics of EasyCrypt modules [10].

IV. FORKING LEMMA

In this section, we start by presenting the pen-and-paper formulation of the general forking lemma by which this work is inspired. In the following subsections, we describe our formalization and notable differences compared to the original work. At the end, we show how we specialize the general result for the common use-case with a random oracle.

⁴This result is less trivial than it might seem at first. Notice that the right-hand side of the inequality talks about probability of a full execution (i.e., initialization followed by the `run` procedure) while the probability on the left-hand side is with respect to a single initialization and two independent runs from the initialized state.

A. Pen-and-Paper Formulation

The original forking lemma was formulated by Pointcheval and Stern [18]. However, we have decided to focus on a later variant of the lemma, the General Forking Lemma by Bellare and Neven [1]. The reason for this choice is two-fold. Firstly, the original statement deals specifically with digital signatures in the random oracle model. By opting for the generalized lemma, we hope to broaden the applicability of our work, potentially beyond signatures. Secondly, the modular nature of the latter lemma makes it arguably more amenable to formal verification.

The essence of the forking lemma, which Bellare and Neven capture in the generalized version, can be described as follows. We consider an algorithm $\mathcal{A}$ that solves a certain problem. Importantly, it receives from the environment a sequence of values sampled at random — whether as a result of interactions with an oracle or as part of its input — and the solution it outputs (if any) is based on one of the obtained values, say at index j . The goal is to obtain from $\mathcal{A}$ two solutions that would be *related* but based on two *distinct* values. To achieve it, we run $\mathcal{A}$ once, generate part of the sequence starting at index j anew (preserving the rest) and run $\mathcal{A}$ for a second time, hoping $\mathcal{A}$ derives its second solution again from the j -th value which differs from the first one. The lemma then bounds the probability that we succeed:

Lemma 1 (General Forking Lemma [1]). *Fix an integer $q \geq 1$ and let $\mathcal{A}$ be a randomized algorithm that takes an input $i \in I$ together with a sequence of values $h_0, \dots, h_{q-1} \in H$, uses random coins $\omega \in \Omega$, and returns a number j from the range $0 \leq j < q$ on success or -1 on error and an auxiliary output $a \in A$.*

Next, consider the algorithm below:

```

Forker $\mathcal{A}$ (i)
 $\omega \leftarrow \$\Omega$  / random coins
 $h_0, \dots, h_{q-1} \leftarrow \$H$  / oracle responses
 $(j, a) \leftarrow \mathcal{A}(i, h_0, \dots, h_{q-1}; \omega)$ 
if  $j = -1$  then
  return  $(-1, \epsilon, \epsilon)$ 
 $h'_j, \dots, h'_{q-1} \leftarrow \$H$ 
 $(j', a') \leftarrow \mathcal{A}(i, h_0, \dots, h_{j-1}, h'_j, \dots, h'_{q-1}; \omega)$ 
if  $j = j' \wedge h_j \neq h'_j$  then
  return  $(j, a, a')$ 
else
  return  $(-1, \epsilon, \epsilon)$ 

```

Finally, for an input $i \in I$, let us define the accepting probability $\text{acc}(i)$ of $\mathcal{A}$ and the success probability $\text{frk}(i)$ of $\text{Forker}_{\mathcal{A}}$:

$$\begin{aligned} \text{acc}(i) &= \Pr[0 \leq j < q : (j, a) \leftarrow \$\mathcal{A}(i, h_0, \dots, h_{q-1}; \omega)], \\ \text{frk}(i) &= \Pr[0 \leq j < q : (j, a, a') \leftarrow \$\text{Forker}_{\mathcal{A}}(i)]; \end{aligned}$$

where in the former, the probability is taken over the choice of $h_0, \dots, h_{q-1}$ and ω .

Then for every input $i \in I$, we have:

$$frk(i) \geq acc(i) \cdot \left(\frac{acc(i)}{q} - \frac{1}{|H|} \right).$$

Intuitively, $\text{Forker}_{\mathcal{A}}$ succeeds only if both runs of $\mathcal{A}$ are accepting, thus the two terms $acc(i)$ are multiplied. The factor $\frac{1}{q}$ appears because the output indices j and j' from the range $\{0, \dots, q-1\}$ must match. Finally, the term $-\frac{1}{|H|}$ accounts for the possible collision of h_j with h'_j sampled from the set H .

Readers familiar with the original work may notice that an important part of the statement, averaging over different inputs i , is omitted here. For presentation purposes, we address this issue later in Section IV-F.

B. Adversary Model

One may observe that the algorithm $\text{Forker}_{\mathcal{A}}$ does not quite match the definition of rewinding that we gave earlier in Section III-B — indeed, the state of $\mathcal{A}$ is never saved (nor restored). Instead, the algorithm $\mathcal{A}$ is run twice from the beginning on related inputs and with the same random tape ω .

To unify the two views on rewinding, suppose that $\mathcal{A}$ uses the provided values $h_0, \dots, h_{q-1}$ one by one. In that case, since the inputs i , the first j values h , and the random tapes are equal, the second execution of $\mathcal{A}$ can diverge from the first one only after $\mathcal{A}$ consumes the value h'_j . Therefore, running $\mathcal{A}$ as $\text{Forker}_{\mathcal{A}}$ does is equivalent to rewinding $\mathcal{A}$, in the sense of Section III-B, to the state before it asks for h_j .

Closer inspection still reveals a minor discrepancy — when the random tape ω is fixed, $\mathcal{A}$ makes the same random choices in both runs, but when $\mathcal{A}$ is rewound, it may sample different values compared to the first run. However, this difference has little impact in practice. Unless we place additional restrictions on $\mathcal{A}$, we must assume that its behavior changes completely once it receives $h'_j \neq h_j$. In particular, it may access different parts of the random tape. In that case, the two runs of $\mathcal{A}$ may diverge arbitrarily, as is the case when we rewind $\mathcal{A}$.

Let us also comment on the assumption that $\mathcal{A}$ accesses the values h sequentially. Firstly, this is the typical setting which appears in the proof for simple Schnorr signatures (see Section V-A) as well as more advanced schemes of interest such as MuSig [7] or FROST [8]. Secondly, an algorithm $\mathcal{A}$ with random access to the provided values can be replaced by an algorithm $\mathcal{B}$ with sequential access such that the outputs are indistinguishable since all values h are sampled uniformly and independently of each other. Thus, random access is meaningful only in the context of $\text{Forker}_{\mathcal{A}}$. However, if we allow this setting, then the two executions of $\mathcal{A}$ may diverge as soon as they encounter values $h'_k \neq h_k$ where $k \geq j$, which can generally occur at any time. The lemma is then arguably less useful as we can no longer claim that the two runs proceeded equally before reading h_j (h'_j). Yet, if this property is not needed, we can enable random access in the state-saving model if we modify $\mathcal{A}$ so that it first accumulates all the values from the environment which it can subsequently use in any order.

Although it should be possible to formalize both approaches to forking using EasyCrypt⁵, we argue that the state-saving model is considerably more practical than the other model that fixes randomness. This is because the EasyCrypt *module language* is inherently stateful and probabilistic. For example, we cannot declare a deterministic abstract module.

One solution could be to limit ourselves to the functional *expression language* and model the rewound programs as operators (`op`; a total function). However, by that, we would give up on a significant proving power of EasyCrypt, such as the relational Hoare logic. Alternatively, we could let users of our forking theory prove that the provided modules are deterministic. But this poses a significant burden and is currently again a use-case out of the scope of EasyCrypt.

Moreover, both discussed workarounds would likely lead to loss of generality as distributions in EasyCrypt are discrete and pre-sampling the random coins for $\mathcal{A}$ would therefore force us to bound the amount of randomness $\mathcal{A}$ can use. Consequently, some almost-surely terminating programs (e.g., relying on rejection sampling), would be inexpressible in this model.

For the reasons above, we decided to further explore the state-saving model. This introduces a new challenge — how do we save the state at the right time? The method outlined in Section III-B cannot record the state of a module mid a procedure call, it merely serializes the global variables of a module (not local). Hence if we wish to obtain the state of $\mathcal{A}$ before it consumes the j -th input value h , we must decompose the computation of $\mathcal{A}$ such that it yields control back to the caller before accessing h_j — similarly to how we split the computation of `module R` in Section III-B into procedures `init` and `run`. But since we learn the value of j only after $\mathcal{A}$ finishes, we require $\mathcal{A}$ to stop before processing *each* next value h . This leads us to the following interface for $\mathcal{A}$:

```

type query_t, resp_t.

type in_t, out_t.

module type Stoppable = {
  proc init(i : in_t) : query_t
  proc continue(r : resp_t) : query_t
  proc finish(r : resp_t) : out_t
}.

```

To mentally map $\mathcal{A}$ to `Stoppable`, think of `query_t` as the unit type⁶, `resp_t` as H , `in_t` as I , and `out_t` as $(\{0, \dots, q-1\} \cup \{-1\}) \times A$.

To generate the responses (values of type `resp_t`) for a `Stoppable` module, we will use a module of the type below (the concrete module will be listed in a later section):

```

module type Oracle = {
  proc get(q : query_t) : resp_t
}.

```

⁵Though there may be further obstacles that we are currently unaware of.

⁶Note that we keep the definition intentionally abstract here so that it can be later reused in the random oracle setting where the query influences the response, see Section IV-G.

The computation of any Stoppable module can then be naturally expressed using a Runner module:

```
(* Number of queries. *)
const Q : {int | 1 ≤ Q} as Q_pos.

module Runner(S : Stoppable, O : Oracle) = {
  proc run(i : in_t) : out_t = {
    var o : out_t;
    var q : query_t;
    var r : resp_t;
    var c : int;

    q ← S.init(i);
    c ← 1;

    while (c < Q) {
      r ← O.get(q);
      q ← S.continue(r);
      c ← c + 1;
    }

    r ← O.get(q);
    o ← S.finish(r);

    return o;
  }
}.
```

It can be shown by structural induction at the meta-level⁷ that any EasyCrypt module (with a run procedure) making a bounded number of oracle queries can be equivalently written as a composition of some Stoppable module and the Runner module.

We also wish to emphasize here that since the Runner module expresses the decomposed computation of a Stoppable module as a single procedure, it enables us to use the Stoppable module as an ordinary module; in particular, we do not have to modify the definitions of standard security notions in any way, as we shall see in Section V-A.

To conclude the discussion on the two different algorithm models, we can observe that both imply certain restrictions to the user in EasyCrypt. Despite that, we believe that the model we chose is superior to the other one, as it imposes only a syntactic restriction (implemented interface) as opposed to a semantic one (determinism and extracted randomness).

C. Forking Module

With the issue of rewinding being resolved in the previous section, we are now ready to formalize the forking algorithm in EasyCrypt.

We begin by specifying which modules can be forked. A Forkable module must implement the Rewindable as well as Stoppable interface so that its state can be queried at the appropriate times:

```
module type Forkable = {
  include Rewindable
  include Stoppable
}.
```

Here, the Stoppable type is a clone of the one defined in Section IV-B where `type out_t = int * aux_t` (cf. the output of $\mathcal{A}$ in Lemma 1).

⁷That is, the theorem cannot be proven in EasyCrypt itself.

We will consider an execution of a Forkable module successful iff the index it outputs is in range:

```
op success (j : int) : bool = 0 ≤ j < Q.
```

As alluded to before, we also have to define the response generator, or the oracle, for the Stoppable module. Following the General Forking Lemma 1, which generates values h independently of $\mathcal{A}$, we use a module that disregards the output (query) of the Stoppable module and always samples a fresh response. We call this oracle a *Forgetful Random Oracle* (FRO):

```
op [lossless uniform] dresp : resp_t distr.

module FRO : Oracle = {
  proc get(q : query_t) : resp_t = {
    var r : resp_t;
    r ← dresp;
    return r;
  }
}.
```

Finally, we need a way to log the responses produced by this oracle since we want to compare them ($h_j \neq h'_j$). For this purpose, we use a generic oracle transformer Log that behaves exactly as the oracle it wraps but additionally records all invocations in its global variable:

```
type log_t = query_t * resp_t.

module Log(O : Oracle) : Oracle = {
  var log : log_t list

  proc get(q : query_t) : resp_t = {
    var r;
    r ← O.get(q);
    log ← log || [(q, r)];
    return r;
  }
}.
```

At this point, the definition of the forking module is straightforward, the main structure copies that of `Forker $\mathcal{A}$` :

```
module Forker(F : Forkable) = {
  var j1, j2 : int
  var log1, log2 : log_t list
  var r1, r2 : resp_t

  (* First full run of F. *)
  proc fst(i : in_t) : out_t * (log_t list) *
    (state_t list) = {
    var sts, st, i, o, q, r, c;
    sts ← [];
    Log.log ← [];
    q ← F.init(i);
    c ← 1;

    while (c < Q) {
      st ← F.getState();
      sts ← sts || [st];
      r ← Log(FRO).get(q);
      q ← F.continue(r);
      c ← c + 1;
    }

    st ← F.getState();
    sts ← sts || [st];
    r ← Log(FRO).get(q);
```

```

    o ← F.finish(r);

    return (o, Log.log, sts);
}

(* Second partial run of F. *)
proc snd(q : query_t, c : int) : out_t * (log_t
list) = {
    var o, r;
    Log.log ← [];

    while (c < Q) {
        r ← Log(FRO).get(q);
        q ← F.continue(r);
        c ← c + 1;
    }

    r ← Log(FRO).get(q);
    o ← F.finish(r);

    return (o, Log.log);
}

proc run(i : in_t) : int * aux_t * aux_t = {
    var sts : state_t list;
    var st : state_t;
    var j : int;
    var a1, a2 : aux_t;
    var q : query_t;

    ((j1, a1), log1, sts) ← fst(i);
    (q, r1) ← nth witness log1 j1;

    (* Rewind. *)
    st ← nth witness sts j1;
    F.setState(st);

    ((j2, a2), log2) ← snd(q, j1 + 1);
    log2 ← (take j1 log1) || log2;
    r2 ← (nth witness log2 j1).2;

    j ← if success j1 ∧ j1 = j2 ∧ r1 ≠ r2
        then j1 else -1;

    return (j, a1, a2);
}
}

```

In the above, the procedure `fst` is a slight modification of `Runner(F, FRO).run` where we record the states `F` goes through and the responses it receives during its execution.

The procedure `snd` runs the module `F` as well with the difference that it responds to `F` just $Q - j$ times, i.e., it can finish the execution of `F` after it is rewound to the j -th query.

D. Probability Analysis

For the forking algorithm presented in the preceding section, we prove a direct equivalent of the probability bound from Lemma 1:

```

const pr_collision =
    1 / (size (to_seq (support dresp))).

lemma pr_fork_success m i :
    let pr_runner_succ =
        Pr[Runner(F, FRO).run(i) @ m : success res.1] in
    let pr_fork_succ =
        Pr[Forker(F).run(i) @ m : success res.1] in
    pr_fork_succ ≥ pr_runner_succ ^ 2 / Q -
        pr_runner_succ * pr_collision.

```

Proof Sketch. In spite of the change in the algorithm model, our formal proof of the forking lemma rather closely copies the original pen-and-paper proof [1]. The main distinction is that we invoke the Rewinding lemma (see Listing 2) at an appropriate time instead of proving an intermediate goal anew. We therefore only outline the main steps of the proof.

First, we observe that conditioned on success $j1$, the success event of the `Forker` is a conjunction of $j1 = j2$ and $r1 \neq r2$. We can thus bound the `Forker`'s success probability as shown below:

```

pr_fork_success ≥
    Pr[Forker(F).run(i) @ m :
        success Forker.j1 ∧ Forker.j1 = Forker.j2] -
    Pr[Forker(F).run(i) @ m :
        success Forker.j1 ∧ Forker.r1 = Forker.r2]

```

The latter probability in the difference is easy to analyze and leads to the term `pr_runner_succ * pr_collision` in the lemma. We perform a case analysis of the former probability:

```

Pr[Forker(F).run(i) @ m :
    success Forker.j1 ∧ Forker.j1 = Forker.j2] =
bigi predT (fun j ⇒
    Pr[Forker(F).run(i) @ m :
        Forker.j1 = j ∧ Forker.j2 = j]
) 0 Q

```

At this point, we can focus on `Pr[Forker(F).run(i) @ m : Forker.j1 = j ∧ Forker.j2 = j]` for a *fixed* j . A simple, but crucial step of our proof is showing that this probability remains the same if we modify `Forker` so that it always rewinds `F` to the j -th query — indeed, the modification only manifests itself if `Forker.j1 ≠ j`, but this event is mutually exclusive with the one we are interested in:

```

module SplitForker(F : Forkable) = {

    (* Forker.fst that runs F only until the
    * first C queries. *)
    proc fst_partial(C : int) : query_t * (log_t
list) * (state_t list) = {
        (* ... *)
    }

    (* Forker.snd, but with state recording. *)
    proc snd(q : query_t, c : int) : out_t * (log_t
list) * (state_t list) = {
        (* ... *)
    }

    proc fst(C : int) : out_t * (log_t list) *
(state_t list) = {
        (* ... *)
        (q, log1, sts1) ← fst_partial(C);
        (o, log2, sts2) ← snd(q, C);
        (* ... *)
    }

    proc run(i : in_t, j : int) : int * int * aux_t *
aux_t * (log_t list) * (log_t list) = {
        (* ... *)
        ((j1, a1), log1, sts1) ← fst(j + 1);
        (* ... *)
        (* Rewind to the j-th query. *)
        q ← (nth witness log1 j).1;

```

```

st ← nth witness sts1 j;
F.setState(st);

((j2, a2), log2, sts2) ← snd(q, j + 1);
(* ... *)
}
}.

```

Notice that after this change, the `SplitForker(F).run` procedure roughly corresponds to `Rewinder(R).run2` (see Listing 1) if we consider `fst_partial` and `snd` to be the `init` and `run` procedures of the module `R` (respectively). This allows us, after some additional refactoring, to apply the Rewinding Lemma (Listing 2) and eventually obtain the inequality below:

```

Pr[Forker(F).run(i) @ m :
  Forker.j1 = j ∧ Forker.j2 = j] ≥
Pr[Runner(F, FRO).run(i) @ m : res.1 = j] ^ 2

```

The final step is to undo the case analysis and express the bound on the sum of probabilities in terms of a single event again. For that, we use the following lemma:

```

lemma square_sum (n : int) (f : int → real) :
  (1 ≤ n) ⇒
  (forall j, 0 ≤ j < n ⇒ 0 ≤ f j) ⇒
  bigi predT (fun j ⇒ square (f j)) 0 n ≥
  square (bigi predT f 0 n) / n.

```

The proof is an easy consequence of Jensen’s inequality, which is present in the standard `EasyCrypt Distr` theory⁸. □

E. Property Transfer

In a pen-and-paper setting, the previous section could likely conclude our discussion of the forking lemma. Here, we need to do additional work to make the result practically usable. The problem is that the probability analysis of the success event alone gives us no information about the side outputs that are produced.

In a typical use-case, we would reason the following way: “If the program `F` succeeds, its side output `a` satisfies some property `P`; hence if the forker succeeds, both runs of `F` must have been successful and thus both auxiliary outputs `a1` and `a2` satisfy `P`.” However, in our formal setting, this is not trivial as the two runs of `F` are essentially interleaved (due to the chosen rewinding approach). We therefore provide the `property_transfer` lemma that captures the implication above:

```

pred P_in : in_t * glob F.
pred P_out : out_t * (log_t list).

axiom run_prop : hoare[
  Runner(F, Log(FRO)).run :
  P_in (i, glob F) ∧ Log.log = [] ⇒
  P_out (res, Log.log)
].

```

⁸<https://github.com/EasyCrypt/easycrypt/blob/r2026.02/theories/distributions/Distr.ec>

```

hoare property_transfer :
  Forker(F).run :
  P_in (i, glob F) ⇒
  let (j, a1, a2) = res in
  success j ⇒
  P_out ((j, a1), Forker.log1) ∧
  P_out ((j, a2), Forker.log2).

```

Moreover, note that the predicate `P_out` is defined over `out_t * (log_t list)` so that we can express properties of the output with respect to the oracle log. Since the logs from the two executions share a prefix (see `Forker` in Section IV-C), this also allows us to relate the two side outputs — thereby replicate the reasoning that “because the values must have been computed before the forking point, they must be equal in both runs”. This will prove useful later in Section V-A.

Proof Sketch. Showing that the output from the first run of `F` satisfies `P_out` is trivial. To show the same about the second output, we must prove that the first run of `F` does not interfere with the second one. For that, we would want to use the following equivalence, which intuitively says that rewinding fully restores the state of `F` (which easily follows from the Axioms 1 to 3 of `Rewindable` modules):

```

equiv[
  <skip> ~ st ← F.getState(); ...; F.setState(st) :
  = {glob F} ⇒ = {glob F}
].

```

The challenge is that we do not a priori know which `getState` call is relevant as we record the state multiple times and choose one only after we obtain the output index `j1`.

To circumvent this problem, we again rely on case analysis (on `j1`) and the `SplitForker` module. The proof thus heavily utilizes the intermediate results of Section IV-D. Specifically, we use the fact that assuming the `Forker` rewinds to the `j`-th query, we can replace it by the `SplitForker` module, which can be further transformed into the shape of the `Rewinder` module. After this change, we can easily syntactically determine the critical `getState` call and apply the equivalence above, essentially proving the below:

```

equiv[
  Rewinder(R).run ~ Rewinder(R).run2 :
  = {i, glob R} ⇒ res{1} = res.2{2}
].

```

□

For convenience, we also provide a `phoare` judgment that combines the probability analysis with the property transfer and some simple assertions about the two produced oracle logs, in particular, that they *first* diverge at the returned index where the queries equal but the responses differ; see the full code for details.

F. Averaging over Inputs

Many security games are played by first randomly generating some parameter (say a key) and only then letting the adversary attack. As a consequence, if we try to analyze the

success probability of a reduction involving the Forker in such a game, the `pr_fork_success` lemma alone (Section IV-D) does not suffice — it is concerned with a fixed input whereas here, we need to average over all possibly generated inputs.

To enable application of the forking lemma in the described scenario, we first introduce a module type for input generators:

```
module type IGen = {
  proc gen() : in_t
}.
```

And then, we create modified versions of the Forker and Runner modules presented earlier:

```
module IForker(I : IGen, F : Forkable) = {
  (* ... *)
  proc fst() : out_t * (log_t list) * (state_t
    list) = {
    (* ... *)
    i ← I.gen();
    q ← F.init(i);
    (* run F *)
  }
  (* ... *)
}.
```

```
module IRunner(I : IGen, S : Stoppable, O : Oracle)
= {
  proc run() : out_t = {
    var i, o;
    i ← I.gen();
    o ← Runner(S, O).run(i);
    return o;
  }
}.
```

For these modules, we prove analogous lemmas as those presented in Sections IV-D and IV-E (we do not list them again for brevity).

Proof Sketch. While Bellare and Neven first prove the probability bound for a fixed input and then solve the average case via Jensen’s inequality and basic properties of expectation, we chose a different approach here.

We notice that the Rewinding Lemma (Section III-B, [10]) already involves some averaging, namely over all the different states the rewind module may end in after running the `init` procedure. The proof in Section IV-D therefore goes through with almost no modifications needed even if we add the `I.gen()` call as shown.

Practically, this has the implication that we prove all lemmas directly for `IForker` and recover the earlier results by using a constant input generator. $\square$

G. Random Oracle Model

Although the flexibility of the forking lemma in its general form can be essential to handle more complex cases, we also implement a specialized version of the lemma for the arguably more common random oracle setting.

Specifically, we add a new `IForkerRO` module that provides the executed module with a standard *Lazy Random Oracle* (LRO) — unlike the earlier `IForker` which provides

a *Forgetful Random Oracle* (FRO) — and prove similar results as before.

One notable accommodation we have to make for this setup is to change the output type of the forked module. It no longer makes sense to return an *index* of a query as repeated queries need to return the same response, i.e., we cannot freely generate a new one after the module is rewound. Instead, we require the forked module to return some queried *value*; `IForkerRO` then rewinds the module to the point where the value was *first queried*. This change is reflected in the definition of the `ForkableRO` module type — the return type of the `finish` procedure changes from `int * aux_t` to `query_t option * aux_t`. The module may return `None` or provide a value that was not queried to signal a failure (previously, this was done by returning an out-of-range index); this condition is checked using the `success_ro` operator (an analog of `success`).

Proof Sketch. Internally, `IForkerRO` executes `IForker` parametrized by the module `F : ForkableRO` wrapped in `Red_LRO_FRO`. This module creates an overlay above the FRO oracle that simulates LRO.

The simulation is simple; we need a single table to store previous responses and *at most one* new FRO query for each LRO query. However, technically, we mirror *all* F’s LRO queries, record the responses of FRO, and fix those that would create an inconsistency in the simulation before passing them to F^9 :

```
module Red_LRO_FRO(F : ForkableRO) : Forkable = {
  var q : query_t
  var m : log_t list

  (* proc getState(), setState(), finish() *)

  proc init(i : in_t) : query_t = {
    m ← [];
    q ← F.init(i);
    return q;
  }

  proc fix_resp(r : resp_t) : resp_t = {
    m ← m || [(q, r)];
    (* Return 1st response associated with q. *)
    r ← oget (assoc m q);
    return r;
  }

  proc continue(r : resp_t) : query_t = {
    r ← fix_resp(r);
    q ← F.continue(r);

    return q;
  }
}.
```

A key step is to show that the above simulation is indeed correct, formally:

⁹We could have mirrored only F’s *unique* queries, i.e., return from the `continue` function only when F makes a new query, but this approach appears to be more challenging as it leads to nested `while` loops once the `Runner` module is involved.

```

equiv red_log_fro_lro_equiv :
  IRunner(I, Red_LRO_FRO(F), Log(FRO)).run ~
  IRunnerRO(I, F, LRO).run :
  = {glob I, glob F} ∧ Log.log{1} = [] ∧
    LRO.m{2} = empty ⇒
  = {glob I, glob F} ∧ ofassoc Log.log{1} = LRO.m{2}.

```

The rest of the proof involves a simple application of the general forking lemma. $\square$

V. APPLICATIONS

To show that our formalization of the forking lemma is practical, we apply it in the security analysis of Schnorr signatures. We also discuss possible future applications.

A. Schnorr Signatures

The security proof of Schnorr signatures was among the first examples of usage of the original forking lemma by Pointcheval and Stern [18]. In addition, the signature algorithm is still highly relevant today. Hence, it presents a natural first test for our forking lemma.

The basis of the signature scheme is the interactive Schnorr identification scheme, which generally allows a prover (P) identified by a public key (pk) to authenticate itself using a secret key (sk) to a verifier (V). Here, the keypair consists of a group element (pk) and its discrete logarithm (sk). The scheme can also be described as a 3-move sigma protocol for the discrete logarithm relation, i.e., P first chooses a *commitment* and sends it to V; next, V responds with a random *challenge*; finally, P returns a *response*. At the end, V can decide based on the *transcript* of the communication with P to accept or reject P's claim. In this context, the public key is called a *statement* and the related secret key is its *witness*.

The signature scheme is then obtained by applying the Fiat-Shamir transformation which replaces the verifier's challenge by the hash of, or the random oracle (RO) output for, the statement, commitment, and message to be signed. This step removes interactivity. The final signature consists of the commitment and the response.

We start the formalization by defining an EasyCrypt module, called `Schnorr`, corresponding to the scheme in the RO model; the definition follows the high-level structure above. The module will be later instantiated by a LRO whose output distribution equals `dchal`.

```

(* Import definition of a cyclic "group" of prime
 * "order" with a generator "g" from the EasyCrypt
 * standard library. The type "exp" corresponds to
 * ZZ/"order". *)

type com_t = group. (* Commitment *)
type chal_t = exp. (* Challenge *)
type resp_t = exp. (* Response *)
type trans_t = (* Transcript *)
  com_t * chal_t * resp_t.

type pk_t = group.
type sk_t = exp.

type msg_t.
type sig_t = com_t * resp_t.

```

```

op [lossless uniform] dnonce : exp distr.
op [lossless uniform] dsk : sk_t distr.
op [lossless uniform] dchal : chal_t distr.

```

```

op verify (pk : pk_t) (t : trans_t) =
  let (com, chal, resp) = t in
  g ^ resp = com * (pk ^ chal).

```

```

module Schnorr (RO : Oracle) = {
  proc keygen() : pk_t * sk_t = {
    var sk, pk;
    sk ← $dsk;
    pk ← g ^ sk;
    return (pk, sk);
  }

  proc sign(sk : sk_t, m : msg_t) : sig_t = {
    var pk, nonce, com, chal, resp;
    pk ← g ^ sk;
    nonce ← $dnonce;
    com ← g ^ nonce;
    chal ← RO.get(pk, com, m);
    resp ← nonce + sk * chal;
    return (com, resp);
  }

  proc verify(pk : pk_t, m : msg_t, s : sig_t) :
    bool = {
    var com, resp, chal;
    (com, resp) ← s;
    chal ← RO.get(pk, com, m);
    return verify pk (com, chal, resp);
  }
}.

```

Listing 3. Schnorr signature scheme.

In the security analysis, we begin with existential unforgeability against key-only attacks (EUF-KOA) [19] and move on to security against chosen message attacks (EUF-CMA) [19] next; both results are relative to the hardness of the discrete logarithm (DL) problem. All mentioned notions are formalized in the EasyCrypt standard library (in the `DigitalSignaturesROM`¹⁰ and `DLog`¹¹ theories, respectively) and we use them without any change¹². Our adversary model includes *forkable* attackers only (in the sense of Section IV).

Let us now proceed to the formulation of the first lemma, which informally says that if an attacker A making at most QR random oracle queries has non-negligible probability of forging a signature in the EUF_KOA_ROM game, then the attacker `Red_KOA_DL(A)` has non-negligible probability of solving the DL problem in the `Exp_DL` experiment, assuming support `dchal` is large enough:

```

lemma schnorr_koa_secure m :
  Pr[Exp_DL(Red_KOA_DL(A)).main() @ m : res] ≥
  Pr[EUF_KOA_ROM(LRO, Schnorr,
    FAdv_KOA_Runner(A)).main() @ m : res] ^ 2
  / (QR + 1) -
  Pr[EUF_KOA_ROM(LRO, Schnorr,
    FAdv_KOA_Runner(A)).main() @ m : res]
  / (size (to_seq (support dchal))).

```

¹⁰<https://github.com/EasyCrypt/easycrypt/blob/r2026.02/theories/crypto/DigitalSignaturesROM.eca>

¹¹<https://github.com/EasyCrypt/easycrypt/blob/r2026.02/theories/crypto/DLog.ec>

¹²With the exception of renaming `DLogExperiment` to `Exp_DL`.

Note that to be able to use A in the `EUF_KOA_ROM` game which expects an adversary implementing the `forge` procedure, we wrap it in the `FAdv_KOA_Runner` module where `forge = Runner(A).run`.

Proof Sketch. The formal proof copies the classical argument. We utilize the *special soundness* [20] of the underlying sigma protocol. This property gives us an (efficient) algorithm that extracts a witness for a given statement from two accepting protocol runs that used the same commitment, but different challenge. Thus, all we have to do is to force the adversary to forge two signatures that are related as described. This is done by *forking* the adversary at the so-called *critical query*, which is the random oracle query used to obtain the challenge for the forgery¹³.

In more detail, we first build a new `ForkableRO` module `AdvWrapper(A)` that runs A , verifies the returned forgery, and, if valid, outputs the critical query along with the forgery itself as a side output; otherwise, it returns *none*. From the definition, it follows that the probability `AdvWrapper(A)` returns *some* query for a randomly sampled key is equal to the probability that A wins the `EUF_KOA_ROM` game.

Next, we define the DL adversary `Red_KOA_DL(A)` which utilizes `ForkerRO` to rewind `AdvWrapper(A)` to the critical query. If the forking algorithm succeeds, it employs the special soundness extractor to recover the secret key. It remains to apply the forking lemma to establish the probability bound on the success of the DL adversary; two challenges arise.

Firstly, `Red_KOA_DL(A)` uses `ForkerRO` (not `IForkerRO`) because it tries to solve *one* input DL challenge. However, in the experiment, the adversary’s success probability is *averaged* over different inputs. It is easy to show that the `Exp_DL` experiment with `Red_KOA_DL(A)` is equivalent to an auxiliary experiment that uses `IForkerRO` with a suitable input generator `KeyGen`. After that, the averaging lemma developed in Section IV-F can be applied.

Secondly, we have to prove that if the forking succeeds, the returned exponent is indeed a correct solution to the DL problem. It suffices to show that the preconditions of the extractor are satisfied. Here, we wish to highlight how the property transfer from Section IV-E crucially simplifies the task. We show that if `AdvWrapper(A)` succeeds, then the returned forgery is valid w.r.t. the used random oracle:

```
hoare success_impl_verify :
  IRunnerRO(KeyGen, AdvWrapper(A), LRO).run :
  true ==>
  let (qopt, forge) = res in
  success_ro LRO.m qopt ==>
  let q = oget qopt in (* Unwrap the option. *)
  let (_pk, com, _m) = q in
  let (_com, resp) = forge in
  verify (g ^ KeyGen.sk) (com, oget LRO.m.[q],
    resp).
```

¹³Note that there is no guarantee that once the adversary is rewound and given a different challenge, it will use this new challenge for its second forgery (in other words, that the same query will be critical in both runs). However, the forking lemma exactly bounds the probability of this event.

Crucially, notice that the assertion uses the commitment included in the critical query. Forking lemma with property transfer immediately gives us that if `IForkerRO` succeeds, the same query was critical in the two runs and both forgeries are valid, each w.r.t. its corresponding oracle. Moreover, we know that the two oracles differ at the critical query. Together, this implies the transcripts are related as required.

In summary, the property transfer allows us to focus on a *single run* of `AdvWrapper(A)`, treat the forker as a *black box*, and still relate the two produced outputs via the used oracles. $\square$

As we outlined earlier, our second result deals with EUF-CMA security. This implies that the adversary A is now additionally allowed to make up to QS signature queries to the signature oracle `BoundedSO` which stops responding after this limit is reached. It is also worth mentioning how we provide A with this oracle; we define the adversary type below:

```
module type FAdv_CMA (SO : SOracle_CMA_ROM) = {
  include Stoppable
  include Rewindable
}.
```

As can be seen, we combine the `Stoppable` interface with EasyCrypt’s built-in method of oracle access, the two approaches can be mixed and do not necessarily clash with each other.

```
lemma schnorr_cma_secure m :
  let pr_cma_succ =
    Pr[EUF_CMA_ROM(LRO, Schnorr, FAdv_CMA_Runner(A),
      BoundedSO).main() @ m : res] in
  let pr_dl_succ =
    Pr[Exp_DL(Red_KOA_DL(Red_CMA_KOA(A))).main()
      @ m : res] in
  let pr_bad = QS * (QS + QR) / order in
  pr_cma_succ >= pr_bad ==>
  pr_dl_succ >=
    (pr_cma_succ - pr_bad) ^ 2 / (QR + 1) -
    1 / (size (to_seq (support dchal))).
```

Proof Sketch. As the formulation of the lemma suggests, we reduce the EUF-CMA setting to the EUF-KOA case. This is in contrast to some pen-and-paper proofs, which reduce EUF-CMA directly to DL — our approach is more layered. It also implies that the proof is uninteresting from the point of view of forking as we have already applied the lemma of interest in the preceding proof. This proof does, however, further show how one can work with `Stoppable` modules.

We also note that the EUF-CMA-to-EUF-KOA reduction was formalized as part of a previous work on the Fiat-Shamir transformation with aborts [5]. Our problem is considerably simpler as the Schnorr identification scheme always terminates successfully (with honest parties).

To carry out the said reduction, we must simulate the signing oracle for the EUF-CMA adversary without knowledge of the secret key. This simulation relies on the (*special*) *honest verifier zero knowledge* (HVZK) [20] property of the associated sigma protocol. The property asserts the existence of an (efficient) algorithm that, given only the statement, produces

transcripts that are distributed identically to those generated by the interaction of the (honest) verifier with the prover (on the same statement and a valid witness). Such an algorithm can then be easily used to forge signatures. However, we must ensure that the random oracle behaves consistently, i.e., it returns the correct challenge for the tuple consisting of the key, commitment, and the message we wish to sign (recall the Fiat-Shamir transformation).

The formal definition of the reduction can be found in the `Red_CMA_KOA` module which resembles the `Red_LRO_FRO` module seen in Section IV-G. We again build a certain *overlay* over the provided random oracle by replaying A’s queries and modifying the received responses, if necessary, before passing them to A. This method enables us to gain control over the oracle and program it freely as required by the HVZK simulator.

It may, of course, happen that the overlay already contains an entry for the query whose response we need to program and the simulation fails as a result. We use the `fel` (Failure Event Lemma) tactic to bound the probability of this happening during a run of `Red_CMA_KOA(A)`. Conveniently, this tactic can be applied to a sequence of statements (as opposed to a single call to an abstract module); having the execution of A span across multiple calls (`init`, `continue` in a loop, and `finish`) therefore poses no significant hurdle.

To conclude the proof, it remains to show that if A wins `EUF_CMA_ROM`, then `Red_CMA_KOA(A)` wins `EUF_KOA_ROM` (provided the bad event does not occur). Here, the main challenge is that in the former game, the forgery is verified against a RO which contains queries made by the signature oracle whereas in the latter game, the *overlaid* RO is used which may not be defined for these queries (unlike the *overlay*). Ultimately, this difference does not matter as a forgery for an already signed message is invalid. However, expressing the relationship between the RO and its overlay forms a substantial part of the formal proof. $\square$

B. Future Work

Given the broad use of the forking lemma, there are multiple directions in which we could extend our current work.

The most obvious choice is to connect our results on the security of Schnorr signatures with a prior work on the *interactive* Schnorr protocol [21] and hash functions¹⁴ in Jasmin [22], a framework designed for “high-assurance and high-speed cryptography”. This way, by means of semantics-preserving extraction of Jasmin to EasyCrypt [16], we could create a fully verified implementation of the Schnorr signature scheme. In the same vein, we could also build on the ongoing Ed25519 Jasmin implementation effort¹⁵.

Another option is to test the limits of our formalization on one of the Schnorr-based multi-party signature schemes from the recent series of works in this area. Some schemes apply the forking lemma in intricate ways, such as nested forking at two

different points [7], while others rely on various modifications of this lemma [23].

Finally, we could also continue with the formalization of security properties of zero-knowledge protocols [14] and generalize the statements from the previous section to the setting where the Fiat-Shamir transformation is applied to an arbitrary sigma protocol satisfying special soundness and honest verifier zero knowledge.

VI. CONCLUSION

This paper introduced the forking technique into EasyCrypt. We first defined an appropriate model for forkable algorithms. Next, we formalized a version of the general forking lemma. Finally, we showed how our result can be used to derive the standard bound on security of Schnorr signatures.

To our knowledge, no previous work focused on forking in EasyCrypt. Our work therefore crucially expands what is in reach of this popular verification framework and we hope it will serve as a stepping stone towards formal analysis of new, more involved cryptographic constructions — notably threshold signatures and zero-knowledge proofs.

REFERENCES

- [1] M. Bellare and G. Neven, “Multi-signatures in the plain public-key model and a general forking lemma,” in *ACM CCS 2006: 13th Conference on Computer and Communications Security*, A. Juels, R. N. Wright, and S. De Capitani di Vimercati, Eds., Alexandria, Virginia, USA: ACM Press, Oct. 2006, pp. 390–399. DOI: 10.1145/1180405.1180453.
- [2] N. Kobitz and A. Menezes, “Critical perspectives on provable security: Fifteen years of ‘another look’ papers,” *Adv. Math. Commun.*, vol. 13, no. 4, pp. 517–558, 2019. DOI: 10.3934/AMC.2019034. [Online]. Available: <https://doi.org/10.3934/AMC.2019034>.
- [3] V. Shoup, *Sequences of games: A tool for taming complexity in security proofs*, Cryptology ePrint Archive, Report 2004/332, 2004. [Online]. Available: <https://eprint.iacr.org/2004/332>.
- [4] M. Barbosa, G. Barthe, K. Bhargavan, *et al.*, “SoK: Computer-aided cryptography,” in *2021 IEEE Symposium on Security and Privacy*, San Francisco, CA, USA: IEEE Computer Society Press, May 2021, pp. 777–795. DOI: 10.1109/SP40001.2021.00008.
- [5] M. Barbosa, G. Barthe, C. Doczkal, *et al.*, “Fixing and mechanizing the security proof of Fiat-Shamir with aborts and Dilithium,” in *Advances in Cryptology – CRYPTO 2023, Part V*, H. Handschuh and A. Lysyanskaya, Eds., ser. Lecture Notes in Computer Science, vol. 14085, Santa Barbara, CA, USA: Springer, Cham, Switzerland, Aug. 2023, pp. 358–389. DOI: 10.1007/978-3-031-38554-4_12.

¹⁴<https://github.com/formosa-crypto/libjade>

¹⁵<https://github.com/formosa-crypto/formosa-25519/pull/30>

- [6] M. Barbosa, F. Dupressoir, B. Grégoire, A. Hülsing, M. Meijers, and P.-Y. Strub, “Machine-checked security for XMSS as in RFC 8391 and SPHINCS⁺,” in *Advances in Cryptology – CRYPTO 2023, Part V*, H. Handschuh and A. Lysyanskaya, Eds., ser. Lecture Notes in Computer Science, vol. 14085, Santa Barbara, CA, USA: Springer, Cham, Switzerland, Aug. 2023, pp. 421–454. DOI: 10.1007/978-3-031-38554-4_14.
- [7] G. Maxwell, A. Poelstra, Y. Seurin, and P. Wuille, “Simple Schnorr multi-signatures with applications to Bitcoin,” *Designs, Codes and Cryptography*, vol. 87, no. 9, pp. 2139–2164, 2019. DOI: 10.1007/s10623-019-00608-x.
- [8] C. Komlo and I. Goldberg, “FROST: Flexible round-optimized Schnorr threshold signatures,” in *SAC 2020: 27th Annual International Workshop on Selected Areas in Cryptography*, O. Dunkelman, M. J. Jacobson Jr., and C. O’Flynn, Eds., ser. Lecture Notes in Computer Science, vol. 12804, Halifax, NS, Canada (Virtual Event): Springer, Cham, Switzerland, Oct. 2020, pp. 34–65. DOI: 10.1007/978-3-030-81652-0_2.
- [9] L. T. A. N. Brandão and R. Peralta, “NIST first call for multi-party threshold schemes,” National Institute of Standards and Technology, Tech. Rep., 2025, Second Public Draft. DOI: 10.6028/NIST.IR.8214C.2pd.
- [10] D. Firsov and D. Unruh, “Reflection, rewinding, and coin-toss in EasyCrypt,” in *CPP ’22: 11th ACM SIGPLAN International Conference on Certified Programs and Proofs, Philadelphia, PA, USA, January 17 - 18, 2022*, A. Popescu and S. Zdancewic, Eds., ACM, 2022, pp. 166–179. DOI: 10.1145/3497775.3503693. [Online]. Available: <https://doi.org/10.1145/3497775.3503693>.
- [11] S. Zain, “Machine-checked verification of digital signature schemes in EasyCrypt,” Ph.D. dissertation, University of Bristol, 2023. [Online]. Available: <https://hdl.handle.net/1983/0c762e35-d43e-4ce7-a7a6-780be08991d1>.
- [12] D. Butler, A. Lochbihler, D. Aspinall, and A. Gascón, “Formalising Σ -protocols and commitment schemes using CryptHOL,” *J. Autom. Reason.*, vol. 65, no. 4, pp. 521–567, 2021. DOI: 10.1007/S10817-020-09581-W. [Online]. Available: <https://doi.org/10.1007/s10817-020-09581-w>.
- [13] G. Barthe, D. Hedin, S. Zanella-Béguelin, B. Grégoire, and S. Heraud, “A machine-checked formalization of sigma-protocols,” in *CSF 2010: IEEE 23rd Computer Security Foundations Symposium*, A. Myers and M. Backes, Eds., Edinburgh, Scotland, UK: IEEE Computer Society Press, Jul. 2010, pp. 246–260. DOI: 10.1109/CSF.2010.24.
- [14] D. Firsov and D. Unruh, “Zero-knowledge in EasyCrypt,” in *CSF 2023: IEEE 36th Computer Security Foundations Symposium*, Dubrovnik, Croatia: IEEE Computer Society Press, Jul. 2023, pp. 1–16. DOI: 10.1109/CSF57540.2023.00015.
- [15] D. Tuma and N. Hopper, *VCVio: A formally verified forking lemma and Fiat-Shamir transform, via a flexible and expressive oracle representation*, Cryptology ePrint Archive, Report 2024/1819, 2024. [Online]. Available: <https://eprint.iacr.org/2024/1819>.
- [16] J. B. Almeida, M. Barbosa, G. Barthe, et al., “The last mile: High-assurance and high-speed cryptographic implementations,” in *2020 IEEE Symposium on Security and Privacy*, San Francisco, CA, USA: IEEE Computer Society Press, May 2020, pp. 965–982. DOI: 10.1109/SP40000.2020.00028.
- [17] *Easycrypt reference manual*, 2024. [Online]. Available: <https://www.easycrypt.info/easycrypt-doc/refman.pdf>.
- [18] D. Pointcheval and J. Stern, “Security proofs for signature schemes,” in *Advances in Cryptology – EUROCRYPT’96*, U. M. Maurer, Ed., ser. Lecture Notes in Computer Science, vol. 1070, Saragossa, Spain: Springer Berlin Heidelberg, Germany, May 1996, pp. 387–398. DOI: 10.1007/3-540-68339-9_33.
- [19] S. Goldwasser, S. Micali, and R. L. Rivest, “A digital signature scheme secure against adaptive chosen-message attacks,” *SIAM Journal on Computing*, vol. 17, no. 2, pp. 281–308, Apr. 1988. DOI: 10.1137/0217017. [Online]. Available: <https://doi.org/10.1137/0217017>.
- [20] R. Cramer, I. Damgård, and B. Schoenmakers, “Proofs of partial knowledge and simplified design of witness hiding protocols,” in *Advances in Cryptology – CRYPTO’94*, Y. Desmedt, Ed., ser. Lecture Notes in Computer Science, vol. 839, Santa Barbara, CA, USA: Springer Berlin Heidelberg, Germany, Aug. 1994, pp. 174–187. DOI: 10.1007/3-540-48658-5_19.
- [21] J. B. Almeida, D. Firsov, T. Oliveira, and D. Unruh, *Schnorr protocol in Jasmin*, Cryptology ePrint Archive, Report 2023/752, 2023. [Online]. Available: <https://eprint.iacr.org/2023/752>.
- [22] J. B. Almeida, M. Barbosa, G. Barthe, et al., “Jasmin: High-assurance and high-speed cryptography,” in *ACM CCS 2017: 24th Conference on Computer and Communications Security*, B. M. Thuraisingham, D. Evans, T. Malkin, and D. Xu, Eds., Dallas, TX, USA: ACM Press, Oct. 2017, pp. 1807–1823. DOI: 10.1145/3133956.3134078.
- [23] C. Komlo, *A note on various forking lemmas*, 2023. [Online]. Available: <https://www.chelseakomlo.com/assets/content/notes/Forking-Lemma-Variants.pdf>.