

Do not Mix Models: Revisiting Generic Transforms for Committing Authenticated Encryption

Kazuhiko Minematsu^{1,2} and Akiko Inoue¹

¹ NEC, Kawasaki, Japan

`k-minematsu@nec.com, a_inoue@nec.com`

² Osaka University, Suita, Japan

Abstract. Committing security for authenticated encryption (AE) captures the difficulty of constructing a distinct input tuple, including the key, that yields the same ciphertext. This notion is relatively new but has attracted significant attention due to its practical relevance. A promising direction is to design generic transforms that convert any AE scheme into a committing one. A common approach to generic transforms, initiated by the CTX transform (Chan and Rogaway, ESORICS 2022), is to add a hash function that uses part of the AE input/output, assuming the hash is ideal, i.e., a random oracle. Because the baseline AE is assumed to be secure in the standard model, this approach inherently mixes standard-model and idealized-model assumptions. We revisit this approach. We show that a number of state-of-the-art generic transforms relying on a mixed model (Chen and Karadžić, Eurocrypt 2025, and Bhattacharjee et al., ePrint 2024), proposed after CTX, are vulnerable once the hash function is instantiated, by presenting practical attacks against them. Our attacks exploit the fact that the baseline AE may depend on the instantiation of the generic transform, whereas the opposite is not true for the principle of the generic transform. In most cases, the attacks are effective with any instantiation, and the baseline AEs in the attacks have a natural structure, such as Enc-then-MAC with a counter mode encryption. We also demonstrate how to rectify these broken transforms with minimal algorithmic modifications, relying solely on the standard-model assumptions.

Keywords: Authenticated Encryption · Committing security · Generic Transform · Provable security · Idealized model · Random Oracle

1 Introduction

(Committing) authenticated encryption. Authenticated Encryption (AE) is a symmetric-key cryptographic function that combines encryption and message authentication [BN00, Rog02]. Traditional AE security notions, namely privacy and authenticity, treat the key as a secret. The recent studies [FOR17, GLR17,

DGRW18, LGR21, IIM21, ADG⁺22] have shown that these notions do not necessarily capture the security requirements for AE expected in real applications. These works introduce the notion of committing security, which captures the infeasibility of finding two distinct input tuples that yield the same ciphertext. The crucial difference from the traditional black-box notions is that the adversary chooses the keys as part of the input. Because this notion is new, most standard AE schemes, such as GCM [Dwo07a], do not satisfy committing security. The idea of committing security for AE was first proposed by Farshim *et al.* [FOR17], followed by formalization and classification of various committing security notions [BH22, CR22]. Now, committing AE is a very active and timely research topic for its practical relevance [GLR17, DGRW18, IIM21, ADG⁺22, LGR21].

A promising approach to achieving a committing AE is to design generic transforms that convert any AE scheme into a committing one. Chan and Rogaway [CR22] introduced an elegant transform, CTX, which augments the baseline AE scheme Π with a single hash function H . CTX computes $T^* = H(K, N, A, T)$, where K is the key of Π , N is a nonce, A is associated data, and T is the tag produced by encryption of Π . While the original output is (C, T) for the ciphertext C , the output of CTX is (C, T^*) , hence it saves T from the output (See Fig. 12). Chan and Rogaway proved that CTX transforms any black-box secure Π into a context-committing AE, assuming H is modeled as a Random Oracle (RO). Here, context committing [BH22, CR22] is the strongest committing security notion.

Despite its simplicity and efficiency, CTX has several limitations. It requires that the encryption algorithm of Π be decomposable into separate (plain) encryption and tag-generation routines. AE schemes with this structure are called Tag-based AEs [CR22], which include GCM and OCB [KR11]. Chen and Karadžić [CK25] (CK25) noted that CTX is not compliant with the AE standard API; specifically, the decryption routine must be modified. CK25 and Bhat-tacharjee *et al.* [BBD24] (BBD24) further noted that CTX does not support misuse-resistant AEs [RS06] and other non-tag-based AEs, which is a practical drawback. In summary, CTX is not strictly generic.³

Both CK25 and BBD24 proposed generic transforms that address these limitations while retaining the core idea of CTX, namely, hashing a portion of AE’s input/output. In particular, CK25 introduced five strictly generic transforms that offer specific usability advantages over existing approaches. Their security proofs follow the CTX methodology; they prove AE security of the transform assuming black-box AE security of Π and assuming that H is an RO. The proposals of BBD24 have strong structural similarities with CTX but incorporate mechanisms to maintain API compatibility while keeping small ciphertext overhead. Consequently, CK25 and BBD24 represent the current state of the art in strictly generic transforms.

³ A strictly generic transform is compatible with standard AE API and requires no additional security property on the baseline AE beyond the black-box AE security.

Our contributions. In this paper, we revisit the “hash-as-ideal” approach introduced and popularized by CTX. Although RO abstraction is extremely useful, real ROs do not exist, and they must be instantiated using concrete hash functions such as SHA-2 or SHA-3 [Dwo15]. Because a generic transform must work with any AE scheme, the baseline Π may also employ SHA-2 or SHA-3. This creates a setting that diverges from the assumptions made in the original proofs, yet is both practical and unavoidable, since SHA-2 and SHA-3 are widely used to construct standard-model-secure functions such as PRFs and nonce-based (authenticated) schemes [KBC97, Sar10, DHM⁺25, RVV15]. We show that, under such instantiations, most transforms proposed by CK25 and BBD24 become completely insecure. For each transform Ω , we construct an AE Π that uses the same instantiation of the hash function H employed by Ω . The resulting Π is natural in structure and provably secure with a reasonable standard model assumption on H . However, once transformed into $\Omega[\Pi, H]$, the scheme becomes vulnerable to simple privacy/authenticity attacks, and in some cases even plaintext recovery or key recovery attacks. These attacks demonstrate that a malicious party could embed a backdoor into a protocol that relies on black-box AE schemes combined with such generic transforms. Although our results do not contradict the original RO-based proofs, they reveal a fundamental gap in the provable-security guarantees of these transforms. Table 1 summarizes our attacks. We have informed the authors of CK25 and BBD24 of our results. Both acknowledged the validity of our attacks, though their assessments of the attacks’ practicality are not entirely consistent with ours.⁴

The root cause of our attacks lies in the use of idealized models, an issue recognized for decades [CGH04]. However, our attacks are based on a slightly different context from those of the previous “uninstantiable RO”-based ones. More specifically, the issue in our setting arises from the mixed use of standard-model and idealized-model assumptions. This mixture obscures the precise conditions under which the baseline AE Π and the hash function H can be instantiated securely. In contrast, generic transforms that rely solely on standard-model assumptions do not suffer from this problem. Guided by this principle, we repair the broken transforms using minimal algorithmic modifications together with standard-model assumptions such as collision resistance and pseudorandomness of the underlying primitives. For most transforms, we apply a pseudorandom generator (PRG) to the secret key to derive two independent keys, one for AE and the other for the transform, a technique we call Key Separation (KS). The computation of KS is lightweight and can be cached, making its online cost negligible (see Table 2). With additional collision resistance assumptions on the PRG and the hash, our fixed versions are provably secure under black-box standard-model AE security notions and context-committing security. By construction, they are immune to the instantiation-sharing attacks we showed against CK25 and BBD24, while preserving desirable features of the original transforms, such as avoiding nonce-based rekeying. We remark that nonce-based

⁴ The authors of BBD24 fully acknowledged practicality, while the authors of CK25 expressed some different views.

Table 1: Our attacks. $\mathbf{X}_e$: the attack which works for a class of instantiations. $\mathbf{X}$ denotes the attack which works for any instantiation. $(\mathbf{X})$: the attack based on slightly artificial, yet provably secure instantiation. **Model** denotes the security model: M = Mixed, S = Standard. Gray cells denote non-existence of attacks for entirely relying on the standard model.

Scheme	Ref.	Generic	Nonce-based Rekey	Model	Our Attack			
					Priv	Auth	Key	Recov.
CTX	[CR22]	$\mathbf{X}$		M	–	–	–	
CTY	[BH24]	$\mathbf{X}$		M	–	–	–	
HtE-UtC/RtC	[BH22]	✓	✓	S				
HtAE	[CK25]	✓	✓	M	–	–	–	
AEaH	[CK25]	✓		M	$\mathbf{X}$	$\mathbf{X}$	($\mathbf{X}$)	
EtH	[CK25]	✓		M	$\mathbf{X}$	$\mathbf{X}$	($\mathbf{X}$)	
AEtH	[CK25]	✓		M	$\mathbf{X}$	–	($\mathbf{X}$)	
chaSIV	[CK25]	✓	✓	M	$\mathbf{X}_e$	–	($\mathbf{X}_e$)	
PACT	[BBD24]	$\approx \checkmark^*$		M	$\mathbf{X}$	$\mathbf{X}$	($\mathbf{X}$)	
comPACT	[BBD24]	$\approx \checkmark^*$		M	$\mathbf{X}$	$\mathbf{X}$	($\mathbf{X}$)	

*: API is generic but requires a non-standard property for AE

rekeying is used by many committing AE designs, such as [BH22]. Notably, two of our fixed transforms (AEaH-KS and AEtH-KS) are the *first* transforms that are 1) strictly generic, 2) context committing, 3) free from nonce-based rekeying, and 4) standard-model secure. Although our transforms introduce little technical novelty, they are meaningful in light of these combined properties. Finally, we briefly analyze CTX and its refinement CTY by Bellare and Hoang [BH24]. We do not identify attacks against them despite their reliance on a mixed model, and we show that standard-model-secure variants can be obtained straightforwardly using the KS technique (See Sect. 5).

Related Work. Bellare and Rogaway [BR93] introduced the RO model and demonstrated its usefulness in provable security. Since the seminal work of Canetti, Goldreich, and Halevi [CGH04], uninstantiable constructions in the RO model have been extensively studied, including those by Bellare, Boldyreva, and Palacio [BBP04] and Brzuska, Farshim, and Mittelbach [BFM15]. Black [Bla06] presented an uninstantiable construction based on the ideal cipher model. More recently, Khovratovich, Rothblum and Soukhanov exploited the RO abstraction in the Fiat-Shamir transform [KRS25]. Committing attacks against widely deployed AE schemes have been reported in many works [GLR17, DGRW18, ADG⁺22, LGR21, CR22, MLGR23]. Bellare and Hoang presented a variety of (generic) transforms for committing AE, including CTY [BH24], and HtE, UtC, and RtC [BH22]. Transforms dedicated for Encode-then-Encipher (EtE) [BR00] style AE were proposed by [HM24b, NSS25]. Dedicated constructions based on popular AE and hash designs, such as GCM, CCM [Dwo07b], Ascon [TMK⁺25], and SHA-3, have been proposed [BH22, NSS24, NSS23, DHM⁺25]. Krämer, Struck, and

Table 2: Our fixes for the CK25 generic transforms. **Alg. chg**: Algorithmic change. **Nb rekey**: Nonce-based rekey. **Add. comp**: Additional cryptographic computation, represented by the total number of increased/decreased input bits given to the primitives, where κ denotes the key bit length.

Scheme	Fix	Alg. chg	Nb Rekey	Add. comp
HtAE-ST	H as RCR-PRF	no	✓	0
AEaH-KS	Key Separation	yes		$O(\kappa)^\dagger$
EtH-KS	Key Separation	yes		$O(\kappa)^\dagger$
AEtH-KS	Key Separation	yes		$O(\kappa)^\dagger$
chaSIV-NKS	N-b Key Sep. + remove nonce from H	yes	✓	0*

†: Computation can be cached per key.

*: Use of a PRF of domain $\mathcal{N}$ (nonce space) reduces input of H by $\log |\mathcal{N}|$ bits.

Weishäupl [KSW24] studied the committing security of the finalists of the NIST Lightweight Cryptography Project [NIS], showing positive results for those based on Sponge/Duplex [BDPV08, BDPV12]. Committing security of Duplex was analyzed by [DFG23]. Additional committing AE constructions have been proposed in works such as [FOR17, GLR17, DGRW18, ADG⁺22, HM24a, SW24, BCC⁺24].

Organization. In Sect. 2, we present preliminaries to define security notions and basic cryptographic primitives. Sect. 3 defines generic transform to committing AE and its security goals. In Sect. 4, we present our attacks against the instantiations of recently-proposed generic transforms by CK25 and BBD24. In Sect. 5, we present how to fix the CK25 transforms with concrete proposals and a brief analysis of CTX and CTY. Sect. 6 concludes the paper.

2 Preliminaries

Basic notations. For integers i and j with $0 \leq i < j$, let $[i..j] := \{i, i+1, \dots, j\}$. Let $\{0, 1\}^*$ be the set of all bit strings, including the empty string ε , the single element in $\{0, 1\}^0$. For $X \in \{0, 1\}^*$, $|X|$ denotes its bit length and $|X|_n$ denotes $\lceil |X|/n \rceil$. We define the parsing of a string into n -bit blocks for a positive integer n by $(X[1], X[2], \dots, X[m]) \stackrel{n}{\leftarrow} X$, where $X[1] \parallel X[2] \parallel \dots \parallel X[m] = X$, $|X[i]| = n$ for $1 \leq i < m$ and $0 < |X[m]| \leq n$ when $|X| > 0$. A concatenation of bit strings X and Y is denoted as $X \parallel Y$, or simply XY . Let 0^i be the string of i zero bits. Hence, we write 10^i for $1 \parallel 0^i$. For $X \in \{0, 1\}^*$ with $|X| \geq i$, $\text{msb}_i(X)$ is the first (left) i bits of X , and $\text{lsb}_i(X)$ is the last (right) i bits of X . Let $\langle X \rangle_c$ be the c -bit standard representation of $X \in [0..2^c - 1]$. For example, $\langle 5 \rangle_4 = 0101$. If c is clear from the context, we may omit the subscript. If X is uniformly chosen from the set $\mathcal{X}$, we write $X \leftarrow \$\mathcal{X}$.

Oracles and advantage. Suppose $\mathcal{A}$ is an adversary in a game. $\mathcal{A}$ can query to s oracles, $\mathcal{O}_1, \dots, \mathcal{O}_s$, in any order. In the case of a distinguishing game, $\mathcal{A}$ outputs a bit after the queries, under the probability space defined by $\mathcal{A}$ and

$\{\mathcal{O}_i\}_{i \in [1..s]}$. Let $[\mathcal{A}^{\mathcal{O}_1, \dots, \mathcal{O}_s} = 1]$ denote the event that this bit is 1. We call $\mathcal{O}_i$ the i -th oracle in the game. By writing $\text{Adv}_{\Pi}^{\text{xxx}}(\mathcal{A})$, we mean the *advantage* of the adversary $\mathcal{A}$ in the game xxx involving Π , and $\mathcal{A}$ may be called xxx-adversary. The time complexity of an adversary is denoted by tm . By writing tm' we mean $\text{tm}' = \text{tm} + O(\theta)$ for some adversarial resource parameter θ specified in the game. We say Π is xxx-secure if the corresponding advantage is negligible for all computationally bounded adversaries.

Primitives. Any keyed function $F : \mathcal{K} \times \mathcal{X} \rightarrow \mathcal{Y}$ with key space $\mathcal{K}$, we may write $F_K(x)$ to denote $F(K, x)$. Unless otherwise specified, we assume $K \leftarrow \$\mathcal{K}$ in evaluation of $F_K(x)$. A block cipher $E : \mathcal{K} \times \mathcal{M} \rightarrow \mathcal{M}$ is a keyed function such that for any $K \in \mathcal{K}$, $E(K, \cdot)$ is a permutation over the message space $\mathcal{M}$. We write $E_K^{-1}(\cdot)$ to denote the decryption function. Let $\text{Func}(\mathcal{X}, \mathcal{Y})$ be the set of all functions: $\mathcal{X} \rightarrow \mathcal{Y}$. A uniform random function (URF) $R : \mathcal{X} \rightarrow \mathcal{Y}$ is a random function uniformly chosen over $\text{Func}(\mathcal{X}, \mathcal{Y})$.

A random oracle (RO) is a public random function $\text{RO} : \{0, 1\}^* \rightarrow \{0, 1\}^m$ for some fixed m . The output is uniformly random for each input. An ideal cipher (IC) is a public function $\text{IC} : \mathcal{K} \times \mathcal{M} \rightarrow \mathcal{M}$ that uniformly distributes over all the block ciphers with key space $\mathcal{K}$ and message space $\mathcal{M}$. If the key space of IC is a singleton, it reduces to a public random permutation (RP), RP. RO, IC, and RP can be evaluated publicly, and IC/RP can be evaluated in both forward and backward directions.

PRG and PRF. For a function $F : \mathcal{K} \times \mathcal{X} \rightarrow \mathcal{Y}$ and a URF $R_F : \mathcal{X} \rightarrow \mathcal{Y}$, let

$$\text{Adv}_F^{\text{prf}}(\mathcal{A}) := |\Pr[\mathcal{A}^{F_K} = 1] - \Pr[\mathcal{A}^{R_F} = 1]|,$$

where $K \leftarrow \$\mathcal{K}$, denote the pseudorandom function (PRF) advantage of F against $\mathcal{A}$. If $\mathcal{A}$ is (q, σ, tm) -PRF adversary, it uses q queries with total σ blocks and tm time. When F takes fixed-length inputs, we omit σ . For a function $G : \mathcal{K} \rightarrow \mathcal{K}'$,

$$\text{Adv}_G^{\text{prg}}(\mathcal{A}) := |\Pr[\mathcal{A}^{G(K)} = 1] - \Pr[\mathcal{A}^{\$} = 1]|,$$

where $K \leftarrow \$\mathcal{K}$ and $\$$ returns uniformly random sample on $\mathcal{K}'$, denote the pseudorandom generator (PRG) advantage of G against $\mathcal{A}$. If $\mathcal{A}$ is a tm -PRG adversary, it has tm time complexity.

(Right-half) Collision resistance. Let $H_1 : \mathcal{X} \rightarrow \mathcal{Y}$ and $H_2 : \mathcal{X} \rightarrow (\mathcal{Y}_1 \times \mathcal{Y}_2)$ be functions. For $(Y_1, Y_2) = H_2(X)$, let $H_2^{\text{right}}(X) = Y_2$. We define Collision Resistance (CR) of H_1 against the adversary $\mathcal{A}_{H_1}$ and Right-half Collision Resistance (RCR) of H_2 against the adversary $\mathcal{A}_{H_2}$ as

$$\begin{aligned} \text{Adv}_{H_1}^{\text{cr}}(\mathcal{A}_{H_1}) &:= \Pr[(X, X') \leftarrow \mathcal{A}_{H_1} : X \neq X', H_1(X) = H_1(X')], \\ \text{Adv}_{H_2}^{\text{rcr}}(\mathcal{A}_{H_2}) &:= \Pr[(X, X') \leftarrow \mathcal{A}_{H_2} : X \neq X', H_2^{\text{right}}(X) = H_2^{\text{right}}(X')]. \end{aligned}$$

If $\mathcal{A}_{H_1}$ ($\mathcal{A}_{H_2}$) is tm -CR (tm -RCR) adversary, it has tm time complexity. RCR is also introduced and used by [CK25].

Almost XOR universal hash function. A function $H : \mathcal{K} \times \mathcal{X} \rightarrow \{0, 1\}^n$ is said to be ϵ -almost XOR universal (ϵ -AXU) if

$$\max_{X, X' \in \mathcal{X}, X \neq X', \delta \in \{0, 1\}^n} \Pr[H(K, X) \oplus H(K, X') = \delta] \leq \epsilon,$$

where K is uniformly random. A popular example of AXU hash function is GHASH inside GCM. It is a classical polynomial hash function defined over $\text{GF}(2^n)$. For $X = (X[1] \parallel X[2] \parallel \dots \parallel X[m])$ with $|X[i]| = n$, it is defined as $\text{PH}(K, X) = \sum_{i \in [1..m]} K^{m-i+1} \cdot X[i]$ (additional length encoding may be needed for variable-length input). For the algebraic structure, once $\text{GHASH}(K, (X_1, X_2))$ is known for some (X_1, X_2) with random K , we can learn K with a high probability. This property has been used by a classical nonce-reuse attack against GCM [Jou06].

2.1 Authenticated encryption

A nonce-based AE scheme $\Pi = (\Pi.\mathcal{E}, \Pi.\mathcal{D})$ is defined over a key space $\mathcal{K}$, a nonce space $\mathcal{N}$, an associated data (AD) space $\mathcal{AD}$, a message space $\mathcal{M}$. An AD is a variable that is not encrypted, but its integrity must be ensured. Formally, AE is a tuple of an encryption function $\Pi.\mathcal{E} : \mathcal{K} \times \mathcal{N} \times \mathcal{AD} \times \mathcal{M} \rightarrow \mathcal{M}$, and a decryption function $\Pi.\mathcal{D} : \mathcal{K} \times \mathcal{N} \times \mathcal{AD} \times \mathcal{M} \rightarrow \mathcal{M} \cup \{\perp\}$, where symbol $\perp$ indicates a verification failure. We require $\Pi.\mathcal{D}_K(N, A, C) = M$ if $\Pi.\mathcal{E}_K(N, A, M) = C$ for correctness. For authenticity, a ciphertext must expand. If Π is a baseline AE scheme before transform (*i.e.*, not committing secure), we assume $|C| = |M| + \tau$ for some fixed, predetermined $\tau > 0$ for any encryption of (N, A, M) . We also assume $\mathcal{K} = \{0, 1\}^\kappa$. If Π is said to be tidy, $\Pi.\mathcal{D}_K(N, A, C) = M$ holds for some $M \neq \perp$ implies $\Pi.\mathcal{E}_K(N, A, M) = C$. Basically, each nonce in encryption must be unique, which we call the nonce-respecting model. A nonce in decryption has no restrictions; thus, it may repeat or collide with those used in encryption. A nonce-respecting adversary may choose a nonce arbitrarily in encryption, except that this condition must be respected [Rog04]. For two AE schemes with the same nonce, AD, and message spaces, and the same ciphertext expansion factor, we call them *compatible*. Note that their key spaces may differ.

For an AE scheme $\Pi = (\Pi.\mathcal{E}, \Pi.\mathcal{D})$, let $\$$ denote the random-bit oracle that takes $(N, A, M) \in \mathcal{N} \times \mathcal{AD} \times \mathcal{M}$ and returns random $|\Pi.\mathcal{E}_K(N, A, M)|$ bits. Here, $\perp$ (as an oracle) denotes the decryption oracle that always returns $\perp$. The standard nonce-respecting security notions of Π are defined as follows.

Definition 1 (PRIV and AUTH). Let $\mathcal{A}_{\text{priv}}$ and $\mathcal{A}_{\text{auth}}$ be nonce-respecting adversaries against Π . PRIV and AUTH advantages are defined as follows.

$$\begin{aligned} \text{Adv}_{\Pi}^{\text{priv}}(\mathcal{A}_{\text{priv}}) &:= |\Pr[\mathcal{A}_{\text{priv}}^{\Pi.\mathcal{E}_K} = 1] - \Pr[\mathcal{A}_{\text{priv}}^{\$} = 1]|, \\ \text{Adv}_{\Pi}^{\text{auth}}(\mathcal{A}_{\text{auth}}) &:= \Pr[\mathcal{A}_{\text{auth}}^{\Pi.\mathcal{E}_K, \Pi.\mathcal{D}_K} \text{ forges}], \end{aligned}$$

The event $[\mathcal{A}_{\text{auth}}^{\Pi.\mathcal{E}_K, \Pi.\mathcal{D}_K} \text{ forges}]$ means that $\mathcal{A}_{\text{auth}}$ receives some $M' \neq \perp$ from $\Pi.\mathcal{D}_K$. The adversary $\mathcal{A}_{\text{auth}}$ is not allowed to make a trivial forgery, namely a decryption query (N, A, C) after obtaining C by an encryption query (N, A, M) .

We say Π is AE-secure if it is secure with respect to both PRIV and AUTH notions.

While nonces used by encryption queries are strictly unique under the PRIV and AUTH notions, it is also important to consider the nonce-misuse setting as introduced by Rogaway and Shrimpton [RS06], namely, encryption queries may repeat nonces, to reflect various practical situations that (erroneously) allow nonce repeat. The security notions capturing this setting are obtained by modifying PRIV and AUTH so that the adversary may repeat nonces in encryption queries but never repeat the entire encryption query (*i.e.*, an encryption query (N, A, M) is always unique). We define these nonce-misusing variations of PRIV and AUTH, written as MR-PRIV and MR-AUTH, as $\text{Adv}_{\Pi}^{\text{mr-priv}}(\mathcal{A}_{\text{priv}})$ and $\text{Adv}_{\Pi}^{\text{mr-auth}}(\mathcal{A}_{\text{auth}})$, assuming $\mathcal{A}_{\text{priv}}$ and $\mathcal{A}_{\text{auth}}$ may repeat nonces in encryption. A class of AE schemes that satisfy the MR-PRIV and MR-AUTH notions is called a Misuse-resistant AE (MRAE). SIV [RS06] is a representative MRAE scheme that combines CMAC-based MAC and a counter mode. We say Π is MRAE-secure if it is secure with respect to both MR-PRIV and MR-AUTH notions.

We write a $(q_e, \sigma_e, \text{tm})$ -PRIV adversary to mean a PRIV adversary of q_e encryption queries with σ_e total blocks in encryption queries and tm time. Similarly, we write a $(q_e, q_d, \sigma_e, \sigma_d, \text{tm})$ -AUTH adversary to mean an AUTH-adversary with q_e encryption queries, q_d decryption queries, σ_e total blocks in encryption queries, σ_d total blocks in decryption queries, and tm time. The same applies to MR-PRIV/MR-AUTH.

Nonce-based encryption. A nonce-based encryption scheme $\Gamma = (\Gamma.\mathcal{E}, \Gamma.\mathcal{D})$, where $\Gamma.\mathcal{E}, \Gamma.\mathcal{D} : \mathcal{K} \times \mathcal{N} \times \mathcal{M} \rightarrow \mathcal{M}$, is a plain encryption scheme without authenticity. For correctness, we require $\Gamma.\mathcal{D}_K(N, \Gamma.\mathcal{E}_K(N, M)) = M$ for any (K, N, M) . When Γ is tidy, $\Gamma.\mathcal{D}_K(N, C) = M$ implies $\Gamma.\mathcal{E}_K(N, M) = C$. We also define the PRIV notion for Γ and the PRIV advantage $\text{Adv}_{\Gamma}^{\text{priv}}(\mathcal{A}_{\text{priv}})$ analogously to the case of AE.

Committing security. Committing security of Π captures the intractability of finding a valid distinct input tuple $((K, N, A, M), (K', N', A', M'))$ that leads to the same ciphertext. Among many notions, we focus on the context-committing security [BH22, CR22], defined as follows.

Definition 2 (Context-committing security). Let $\Pi = (\Pi.\mathcal{E}, \Pi.\mathcal{D})$ be an AE scheme. The context committing advantage of $\mathcal{A}_{\text{cmt}}$ against Π is defined as

$$\begin{aligned} \text{Adv}_{\Pi}^{\text{cmt}}(\mathcal{A}_{\text{cmt}}) &:= \Pr[((K, N, A, M), (K', N', A', M')) \leftarrow \mathcal{A}_{\text{cmt}} : \\ &\quad (K, N, A, M) \neq (K', N', A', M'), \\ &\quad \Pi.\mathcal{E}(K, N, A, M) = \Pi.\mathcal{E}(K', N', A', M')]. \end{aligned}$$

Context committing security is the strongest among the committing security notions, as it does not restrict the input tuple beyond distinctness, $(K, N, A, M) \neq (K', N', A', M')$. It is also called CMT-4 [BH22] or full committing [CR22]. We write CMT to mean context committing.

2.2 Additional notations

All the proofs of our proposals are based on the hybrid arguments. For notational convenience, we introduce the following general indistinguishability notions.

Definition 3 (Indistinguishability of AE). *Let $\Pi_i = (\Pi_i.\mathcal{E}, \Pi_i.\mathcal{D})$ for $i = 1, 2$ be compatible AE schemes with a key space $\mathcal{K}_i$.*

$$\begin{aligned}\text{Adv}_{\Pi_1.\mathcal{E}, \Pi_2.\mathcal{E}}^{\text{ind}}(\mathcal{A}) &:= |\Pr[\mathcal{A}^{\Pi_1.\mathcal{E}_{K_1}} = 1] - \Pr[\mathcal{A}^{\Pi_2.\mathcal{E}_{K_2}} = 1]|, \\ \text{Adv}_{\Pi_1, \Pi_2}^{\text{ind}}(\mathcal{A}') &:= |\Pr[\mathcal{A}'^{\Pi_1.\mathcal{E}_{K_1}, \Pi_1.\mathcal{D}_{K_1}} = 1] - \Pr[\mathcal{A}'^{\Pi_2.\mathcal{E}_{K_2}, \Pi_2.\mathcal{D}_{K_2}} = 1]|,\end{aligned}$$

where $\mathcal{A}$ and $\mathcal{A}'$ are nonce-respecting, and the probabilities are defined over the adversary's randomness and $K_i \leftarrow \$\mathcal{K}_i$ for $i \in \{1, 2\}$. The adversary $\mathcal{A}'$ is not allowed to make a trivial forgery, namely a decryption query (N, A, C) after obtaining C by an encryption query (N, A, M) .

By definition, for any $\mathcal{A}_{\text{priv}}$ and $\mathcal{A}_{\text{auth}}$, we have

$$\begin{aligned}\text{Adv}_{\Pi}^{\text{priv}}(\mathcal{A}_{\text{priv}}) &= \text{Adv}_{\Pi.\mathcal{E}, \$}^{\text{ind}}(\mathcal{A}_{\text{priv}}), \\ \text{Adv}_{\Pi}^{\text{auth}}(\mathcal{A}_{\text{auth}}) &= \text{Adv}_{\Pi, (\Pi.\mathcal{E}, \perp)}^{\text{ind}}(\mathcal{A}_{\text{auth}}).\end{aligned}$$

For any Π_1, Π_2 and AUTH adversary $\mathcal{A}_{\text{auth}}$,

$$\text{Adv}_{\Pi_1}^{\text{auth}}(\mathcal{A}_{\text{auth}}) - \text{Adv}_{\Pi_2}^{\text{auth}}(\mathcal{A}_{\text{auth}}) \leq \text{Adv}_{\Pi_1, \Pi_2}^{\text{ind}}(\mathcal{A}')$$

holds for some distinguishing adversary $\mathcal{A}'$ using the same number of queries and the total queried blocks as $\mathcal{A}_{\text{auth}}$.

3 Defining Generic Transform

Let $\Omega[*, \mathbb{P}]$ be a generic transform for committing AEs. The first argument takes either an AE scheme or a nonce-based plain encryption scheme, and $\mathbb{P}$ denotes a list of cryptographic primitives specified by the transform Ω . Let $\mathbb{P} = (F_1, \dots, F_a, P_1, \dots, P_b)$ for some positive integers a and b , where F_i denotes a primitive under a standard model (*e.g.*, PRF) and P_j denotes a primitive under an idealized model (*e.g.*, RO). We say $\Omega[*, \mathbb{P}]$ is instantiated if all the primitives in $\mathbb{P}$ are instantiated, *i.e.*, their concrete algorithms are determined. For example, F_1 may be an 128-bit block PRP and an instantiation of F_1 by AES-128 means that we define $F_1(x)$ as $\text{AES}(K, x)$ with $K \leftarrow \$\{0, 1\}^{128}$. If P_i is an RO, its instantiation could be typically a popular cryptographic hash function such as SHA-2 or SHA-3 [Dwo15].

Our focus is *generic* transform; thus, we assume Ω can access baseline AE Π only in a black-box way, that is, it can invoke $\Pi.\mathcal{E}$ and $\Pi.\mathcal{D}$ but cannot access any internal functions. If Π uses $\mathbb{P}$ as its primitives, we write $\Pi[\mathbb{P}]$. The goal of a generic transform is to guarantee AE security and committing security after the transform.

Let $\mathbb{P}$ denote an instantiation of $\mathbb{P}$. The corresponding instantiation of $\Omega[*, \mathbb{P}]$ is denoted by $\Omega[*, \mathbb{P}]$. We consider the following security properties of Ω .

Definition 4. Let $\Omega[*, \overline{\mathbb{P}}]$ be an instantiation of Ω and let Π be an efficient (i.e., polynomial-time) AE scheme with AE security. The algorithm of Π may depend on $\Omega[*, \overline{\mathbb{P}}]$. We say $\Omega[\Pi, \overline{\mathbb{P}}]$ is broken if any of PRIV/AUTH/CMT security of $\Omega[\Pi, \overline{\mathbb{P}}]$ is broken with $O(1)$ queries and negligible time.

We say Ω is uninstantiable if there exists AE-secure $\Pi[\overline{\mathbb{P}}]$ such that $\Omega[\Pi[\overline{\mathbb{P}}], \overline{\mathbb{P}}]$ is broken for any $\overline{\mathbb{P}}$. Similarly, we say Ω is conditionally uninstantiable if there exists an AE-secure $\Pi[\overline{\mathbb{P}}]$ such that $\Omega[\Pi[\overline{\mathbb{P}}], \overline{\mathbb{P}}]$ is broken for some $\overline{\mathbb{P}}$.

If Ω is uninstantiable it is also conditionally uninstantiable, but not vice versa.

Remarks. Definition 4 reflects the natural motivation of generic transforms, i.e., a generic transform should maintain security irrespective of the algorithm of a baseline secure AE scheme. If the target transform is uninstantiable or conditionally uninstantiable, a potentially malicious party could build an AE scheme that is secure as a stand-alone scheme, however, it turns out to be totally broken when transformed into a committing AE. It would be generally non-trivial to tell if a generic transform involving idealized primitives is securely instantiable since an idealized primitive does not exist in the first place. Conversely, a provably secure transform that uses only standard-model-secure primitives is inherently immune to the attack against a specific instantiation. That is, there is no gap between the model and the instantiation.

4 Analysis of Generic Transforms

Definition 4 suggests an interesting object to investigate: a generic transform with provable security claims based on idealized-model assumptions on its primitives and standard-model assumptions on AE, which we call the mixed (security) model. A provable security claim based on the mixed model does not generally guarantee the security of its instantiation, and at worst, attacks on instantiations may occur.

This section shows that the worst case occurs with the state-of-the-art generic transforms proposed by CK25 and BBD24. All our attacks result in privacy and/or authenticity break, and even plaintext recovery or key recovery, with a few queries. Most of our attacks work against any instantiation of the internal primitives, thus showing they are uninstantiable. The “crafted” AE schemes used by our attacks are mostly quite simple and natural; for instance, some use a keyed hash function in the counter mode and build an AE via generic composition. We remark that hash-based encryption schemes have been discussed, e.g., by [HKR01, PRZ99], and have received attention by practitioners, as evidenced by multiple threads on Stack Exchange [STeA, STeB, STeC, STeD].

4.1 Transforms by Chen and Karadžić

At Eurocrypt 2025, Chen and Karadžić proposed a family of generic transforms [CK25] consisting of three basic transforms, HtAE, AEaH, and EtH, and

$\text{AEaH}.\mathcal{E}[\Pi, \mathbf{H}](K, N, A, M)$ 1. $C \leftarrow \Pi.\mathcal{E}(K, N, \varepsilon, M)$ 2. $T \leftarrow \mathbf{H}(K, N, A)$ 3. return (C, T)	$\text{AEaH}.\mathcal{D}[\Pi, \mathbf{H}](K, N, A, (C, T))$ 1. $M \leftarrow \Pi.\mathcal{D}(K, N, \varepsilon, C)$ 2. if $M = \perp \vee \mathbf{H}(K, N, A) \neq T$: 3. return $\perp$ 4. return M
$\text{EtH}.\mathcal{E}[\Gamma, \mathbf{H}](K, N, A, M)$ 1. $C \leftarrow \Gamma.\mathcal{E}(K, N, M)$ 2. $T \leftarrow \mathbf{H}(K, N, A, C)$ 3. return (C, T)	$\text{EtH}.\mathcal{D}[\Gamma, \mathbf{H}](K, N, A, (C, T))$ 1. $T^* \leftarrow \mathbf{H}(K, N, A, C)$ 2. if $T^* \neq T$: return $\perp$ 3. $M \leftarrow \Gamma.\mathcal{D}(K, N, C)$ 4. return M
$\text{HtAE}.\mathcal{E}[\Pi, \mathbf{H}](K, N, A, M)$ 1. $(K_e, T) \leftarrow \mathbf{H}(K, N, A)$ 2. $C \leftarrow \Pi.\mathcal{E}(K_e, N, \varepsilon, M)$ 3. return (C, T)	$\text{HtAE}.\mathcal{D}[\Pi, \mathbf{H}](K, N, A, (C, T))$ 1. $(K_e, T^*) \leftarrow \mathbf{H}(K, N, A)$ 2. if $T^* \neq T$: return $\perp$ 3. $M \leftarrow \Pi.\mathcal{D}(K_e, N, \varepsilon, C)$ 4. return M

Fig. 1: Basic transforms of [CK25]. (Top) AEaH, (Middle) EtH, and (Bottom) HtAE.

two advanced transforms, AEtH and chaSIV. Their proposals are very clean and simple: most of them invoke a hash function that takes a portion of baseline AE’s input/output. The simplicity of their schemes makes them similar to the aforementioned CTX [CR22], but unlike CTX, their schemes are strictly generic. The basic transforms achieve PRIV and AUTH (thus under nonce-respecting setting) and CMT security, and the advanced ones achieve MR-PRIV and MR-AUTH and CMT notions.⁵ Some schemes (AEaH, EtH, AEtH) avoid nonce-based rekeying, a common approach to achieving committing security, but this approach is costly. See the paper [CK25] for a more detailed comparison with the previous (generic) transforms.

We show practical PRIV/AUTH attacks against their instantiations, except HtAE. We also present full- and partial-plaintext recovery and key-recovery attacks for some schemes. All our attacks require a few queries. Throughout the section, we assume all the internal idealized primitives have been instantiated. For example, by writing $\mathbf{H}$, we mean a concrete hash function.

4.2 Attacks against AEaH

AE and Hash (AEaH) is a basic transform which uses a hash function $\mathbf{H} : \{0, 1\}^* \rightarrow \{0, 1\}^t$.⁶ See Figs. 1 and 2 for the pseudocode and the illustration. Let $\text{AEaH}[\Pi, \mathbf{H}]$ denote AEaH applied to an AE Π . CK25 proved PRIV and AUTH for $\text{AEaH}[\Pi, \mathbf{H}]$

⁵ They also achieve context discovery security [MLGR23], which is out of our focus.

⁶ To be more precise, AEaH assumes $\mathbf{H}$ to accept a tuple input in $\mathcal{K} \times \mathcal{N} \times \mathcal{AD}$, hence we assume an injective encoding here. Similar implicit encoding is assumed for other transforms as well.




Fig. 2: AEaH.




Fig. 3: EtH.




Fig. 4: HtAE.

assuming that H is an RO and Π is AE secure. Recall that we allow Π to use H as we consider attacks against instantiations of **AEaH** (thus H).

Privacy attack. We first describe a privacy attack. Let Π be an Enc-then-MAC AE using H as the primitive. For an mt -bit plaintext $M = (M[1] \parallel \dots \parallel M[m])$, $|M[i]| = t$, the tuple (N, A, M) is encrypted as $\Pi.\mathcal{E}(K, N, A, M) \rightarrow C = (\tilde{C}, \tilde{T})$, $\tilde{C} = (\tilde{C}[1] \parallel \dots \parallel \tilde{C}[m])$ such that

$$\tilde{C}[i] = H(K, N, \langle i \rangle) \oplus M[i] \text{ for } i \in [1..m].$$

$\tilde{T}$ is the tag of Π defined as $\tilde{T} = \text{MAC}_K(N, A, \tilde{C})$ for some MAC function MAC . MAC may be arbitrarily instantiated by H or any other additional primitive, using the key K , possibly with a domain separation from the encryption part. A secure AE can be easily realized assuming H_K is a PRF by provable security analysis of Enc-then-MAC [NRS14]. The PRF assumption on H_K would be reasonable if H is a reasonable instantiation of RO. The privacy attack works as follows.

1. Encrypt (N, A, M) with any N , $A = \langle 1 \rangle$, $|M| = mt$ for some $m \geq 1$. Receive $(C = (\tilde{C}, \tilde{T}), T) = \text{AEaH}[\Pi, H].\mathcal{E}(K, N, A, M)$.
2. Output 1 (decide as real) if $\tilde{C}[1] \oplus T = M[1]$ holds. Otherwise, output 0.

Clearly, because $T = H(K, N, A) = H(K, N, \langle 1 \rangle)$, $\tilde{C}[1] \oplus T = M[1]$ holds with probability 1 in the real world, while with a probability $\sim 1/2^t$ in the ideal world. The adversary's advantage is almost 1. Moreover, a (partial) plaintext recovery attack is possible; observing (C, T) for the aforementioned $(N, A = \langle 1 \rangle, M)$ with unknown M will reveal $M[1]$. Generally, if A is chosen among $\{\langle j \rangle\}_{j \in [1..m]}$, any

$\text{AEtH}.\mathcal{E}[\Pi, H](K, N, A, M)$ 1. $C \leftarrow \Pi.\mathcal{E}(K, N, A, M)$ 2. $T \leftarrow H(K, N, A, \text{lsb}_r(C))$ 3. return (C, T)	$\text{AEtH}.\mathcal{D}[\Pi, H](K, N, A, (C, T))$ 1. $T^* \leftarrow H(K, N, A, \text{lsb}_r(C))$ 2. if $T^* \neq T$: return $\perp$ 3. $M \leftarrow \Pi.\mathcal{D}(K, N, A, C)$ 4. return M
$\text{chaSIV}.\mathcal{E}[\Gamma, F, H](K, N, A, M)$ 1. $K_e \leftarrow F(K, N)$ 2. $T \leftarrow H(K, N, A, M)$ 3. $N_e \leftarrow \text{lsb}_r(T)$ 4. $C \leftarrow \Gamma.\mathcal{E}(K_e, N_e, M)$ 5. return (C, T)	$\text{chaSIV}.\mathcal{D}[\Gamma, F, H](K, N, A, (C, T))$ 1. $K_e \leftarrow F(K, N)$ 2. $N_e \leftarrow \text{lsb}_r(T)$ 3. $M \leftarrow \Gamma.\mathcal{D}(K_e, N_e, C)$ 4. if $M = \perp$: return $\perp$ 5. $T^* \leftarrow H(K, N, A, M)$ 6. if $T^* \neq T$: return $\perp$ 7. return M

Fig. 5: Advanced transforms of [CK25]. (Top) AEtH and (Bottom) chaSIV.




Fig. 6: AEtH.




Fig. 7: chaSIV.

single plaintext block could be recovered. When $m = 1$, it means full plaintext recovery.

Authenticity attack. Let Γ be a nonce-based encryption scheme. We instantiate Π as a variant of GCM combining Γ and GHASH and H . The encryption algorithm of Π is as follows. Suppose $\mathcal{N} = \{0, 1\}^n \setminus \{\text{cst}_1, \text{cst}_2\}$ for some distinct constants cst_1 and cst_2 . Let $\Pi.\mathcal{E}(K, N, A, M) = C = (\tilde{C}, \tilde{T})$, where

$$L_i = H(K, \text{cst}_i, \varepsilon) \text{ for } i = 1, 2, \quad \tilde{C} = \Gamma.\mathcal{E}(L_1, N, M), \\ \tilde{T} = \text{GHASH}(L_2, (A, \tilde{C})) \oplus H(K, N, A).$$

The above Π is provably AE-secure if Γ is PRIV secure and H_K is a PRF; the entire scheme can be seen as a variant of GCM with two keys, and the keys L_1, L_2 and the last masking value $H(K, N, A)$ will look independent and random thanks to the domain separation. Its provable security bounds will be easily derived using the known results on generic composition and GCM [NRS14, IOM12]. Recall that the baseline Π fixes AD to empty when composed. The transformed scheme

$\text{AEaH}[\Pi, H]$ is vulnerable to an authenticity attack. The procedure is similar to the classical nonce-reuse attack against GCM [Jou06] and enables almost universal forgery (arbitrary manipulation of the ciphertext). Assume AD $A = \varepsilon$.

1. Encrypt $(N, A(= \varepsilon), M)$ and receive $(C, T) = \text{AEaH}[\Pi, H].\mathcal{E}(K, N, \varepsilon, M)$, where $C = (\tilde{C}, \tilde{T}) = \Pi.\mathcal{E}(K, N, \varepsilon, M)$ and $T = H(K, N, \varepsilon)$.
2. Since $\tilde{T} \oplus T = \left(\text{GHASH}(L_2, (\varepsilon, \tilde{C})) \oplus H(K, N, \varepsilon) \right) \oplus H(K, N, \varepsilon)$ is equal to $\text{GHASH}(L_2, (\varepsilon, \tilde{C}))$, the adversary recovers L_2 .
3. Decrypt $(N, A(= \varepsilon), C', T)$ such that $C' = (\tilde{C}', \tilde{T}')$, $\tilde{C}' \neq \tilde{C}$, and $\tilde{T}' = \text{GHASH}(L_2, (\varepsilon, \tilde{C}')) \oplus H(K, N, \varepsilon)$.

The decryption query will be accepted with probability 1.

4.3 Attacks against EtH

Enc then Hash (EtH) is a basic transform which uses a hash function $H : \{0, 1\}^* \rightarrow \{0, 1\}^t$. The baseline scheme Γ may be a nonce-based encryption or an AE. See Figs. 1 and 3 for the pseudocode and the illustration. Let $\text{EtH}[\Gamma, H]$ denote EtH applied to Γ . CK25 proved PRIV and AUTH for $\text{EtH}[\Gamma, H]$ assuming that H is an RO and Γ is PRIV secure.

Privacy attack. Our privacy attack utilizes a H -based CTR mode with a certain trick in the argument. Let $\mathcal{N} = \{0, 1\}^*$, $\mathcal{M} = \{0, 1\}^{\leq m \cdot t}$ for some $m > 0$, and $\text{AD} = \{0, 1\}^*$. We define $\Gamma.\mathcal{E} : \mathcal{K} \times \mathcal{N} \times \mathcal{M} \rightarrow \mathcal{M}$ as a variant of CTR. Let $\text{cst} \in \{0, 1\}^*$ be a constant. Let $c = \log_2 m$.

For $(N, M) \in \mathcal{N} \times \mathcal{M}$ such that $(M[1], \dots, M[m]) \xleftarrow{t} M$, $C = \Gamma.\mathcal{E}(K, N, M)$ where $C = C[1] \parallel \dots \parallel C[m]$ is defined as follows.

$$C[i] = \text{msb}_{|M[i]|}(H(K, \text{cst}, N, \langle i-1 \rangle_c)) \oplus M[i].$$

This Γ is PRIV secure if H_K is a PRF; however, the placement of nonce in the argument of H is inconsistent with that of H call inside EtH. This inconsistency is exploited as follows.

1. Encrypt $(N = \text{cst}, A, M)$ for arbitrary $A \neq \text{cst}$ and M such that $|M| = \ell$ for some $\ell \in [1..c]$ (possibly 1). Receive (C, T) where $|C| = \ell$ and $T = H(K, \text{cst}, A, C)$.
2. Let $j \in [1..2^c]$ so that $\langle j-1 \rangle_c = C$. Encrypt $(N' = A, A', M')$ for arbitrary A' and M' such that $|M'| \geq j \cdot t$. Receive (C', T') .
3. Check if $C'[j] \oplus M'[j]$ equals to T . If it holds, decide it as real, and otherwise decide it as ideal.

Since

$$\begin{aligned} C'[j] &= H(K, \text{cst}, N', \langle j-1 \rangle_c) \oplus M'[j] \\ &= H(K, \text{cst}, A, C) \oplus M'[j] \\ &= T \oplus M'[j] \end{aligned}$$

holds, the check will always succeed in the real world, whereas it succeeds with a negligible probability in the ideal world. Thus, the advantage of the attack is almost 1.

Authenticity attack. Similar to the case of AEaH, an authenticity attack is possible when Γ is a counter mode using H : we first parse M into t -bit blocks, $M[1], M[2], \dots$, and encrypt as $C[i] = H(K, N, \mathbf{cst}, \langle i \rangle_c) \oplus M[i]$ for some constant $\mathbf{cst} \in \{0, 1\}^*$ and some sufficiently large $c > 0$. Then $H(K, N, \mathbf{cst}, \langle i \rangle_c)$ for $i = 1, 2, \dots$, can be learned by an encryption query, resulting in the following attack.

1. Encrypt (N, A, M) for any (N, A) and $|M| \geq m \cdot t$ for some positive integer m . Receive $(C, T) = \text{EtH}[\Pi, H].\mathcal{E}(K, N, A, M)$.
2. Pick any $i \in [1..m]$. Decrypt $(N, A' = \mathbf{cst}, C' = \langle i \rangle_c, T' = M[i] \oplus C[i])$.

The decryption query will be accepted as $M[i] \oplus C[i] = H(K, N, \mathbf{cst}, \langle i \rangle_c) = H(K, N, A', C')$.

4.4 Attack against AEtH

AE then Hash (AEtH) is an advanced transform which uses a hash function $H : \{0, 1\}^* \rightarrow \{0, 1\}^t$. See Figs. 5 and 6 for the pseudocode and the illustration. Let $\text{AEtH}[\Pi, H]$ denote AEtH applied to an AE Π . CK25 proved MR-PRIV and MR-AUTH for $\text{AEtH}[\Pi, H]$ assuming that H is an RO and Π is MRAE secure. AEtH is called MRAE-preserving for this security guarantee.

Let Π be an SIV-like MRAE scheme using some variable-input-length PRF $F : \mathcal{K} \times \{0, 1\}^* \rightarrow \{0, 1\}^\tau$ and H . For encryption of (N, A, M) for $|M| = mt$, $\Pi.\mathcal{E}$ first computes $V = F(K, N, A, M)$ (assuming implicit encoding). After t -bit parsing of M , each block is encrypted as $C[i] = M[i] \oplus H(K, N, A, V \oplus \langle i - 1 \rangle_\tau)$ for $i \in [1..m]$. The output is $C = (C[1] \parallel \dots \parallel C[m] \parallel V)$. The provable MRAE security of this Π is established if MAC part (F_K) and Encryption part (H_K) are computationally independent [RS06]. In practice, this could be ensured by using a different algorithm, *e.g.*, using SHA-3 for H and CMAC-AES for F_K , or instantiating F by H with an adequate domain separation.

When the above Π is transformed by AEtH, an encryption query (N, A, M) generates $(C, T) = \text{AEtH}[\Pi, H].\mathcal{E}(K, N, A, M)$, where $C = (C[1] \parallel \dots \parallel C[m] \parallel V)$ and $T = H(K, N, A, \text{lsb}_\tau(C))$ with $\text{lsb}_\tau(C) = V = F(K, N, A, M)$.

Privacy attack. The privacy attack is similar to the previous cases: since $C[1] \oplus M[1] = H(K, N, A, V \oplus \langle 0 \rangle_\tau (= 0^\tau)) = T$ holds, the adversary can distinguish (C, T) from random with high probability. In the context of a plaintext recovery attack, the adversary can recover $M[1]$.

We have not found any authenticity attack.

4.5 Attack against chaSIV

Committing hash-based SIV (chaSIV) is an advanced transform which uses a hash function $H : \{0, 1\}^* \rightarrow \{0, 1\}^t$ and a PRF $F : \mathcal{K} \times \mathcal{N} \rightarrow \mathcal{K}_e$. The baseline scheme Γ is either a nonce-based encryption or an AE, taking key space $\mathcal{K}_e$ and nonce space $\mathcal{N}_e = \{0, 1\}^r$. Following CK25, we assume $\mathcal{K} = \mathcal{K}_e = \{0, 1\}^\kappa$, which ensures that the key size does not change before and after the transform. See Figs. 5 and 7 for the pseudocode and the illustration. Let $\text{chaSIV}[\Gamma, F, H]$ denote chaSIV applied to Γ . CK25 proved MR-PRIV and MR-AUTH for $\text{chaSIV}[\Gamma, F, H]$ assuming that H is an RO, F is a PRF, and Γ is PRIV secure, thus chaSIV is *MRAE-lifting*.

Unlike other transforms, the primitives of the chaSIV transform already rely on the mixed model. We consider the case that F uses the instantiation of H as a primitive.

Privacy attack. Let $\kappa \leq t$ and $r \leq t$. Let $F(K, N)$ be $\text{msb}_\kappa(H(K, N, \text{cst}_1, \text{cst}_2))$ for some constants cst_1 and cst_2 (they may depend on N). Following other cases, the PRF assumption on F would be reasonable. To break privacy, the adversary encrypts (N, A, M) for $A = \text{cst}_1$ and $M = \text{cst}_2$. Then the output $(C, T) = \text{chaSIV}.\mathcal{E}[\Gamma, H, F](K, N, A, M)$ reveals K_e as $K_e = \text{msb}_\kappa(T)$. This attack implies a complete privacy break. Moreover, if the adversary does not know M , it implies full plaintext recovery. The attack works regardless of the algorithm used by Γ .

The above PRIV attack is devastating (full plaintext recovery for any Γ). On the other hand, we could not find an authenticity attack. Moreover, our PRIV attack needs specific conditions on the instantiation (*i.e.*, F uses H). The latter implies that chaSIV is conditionally uninstantiable. We have not found an authenticity attack.

4.6 PACT and comPACT transforms

Bhattacharjee *et al.* [BBD24] proposed PACT and comPACT. Their structures are similar to CTX. Both transforms use a hash function $H : \{0, 1\}^* \rightarrow \{0, 1\}^t$ and a block cipher $E : \mathcal{K}_b \times \mathcal{M}_b \rightarrow \mathcal{M}_b$, $\mathcal{K}_b = \{0, 1\}^t$ and $\mathcal{M}_b = \{0, 1\}^\tau$. We require $\tau = t$. The τ -bit tag part of the baseline AE is encrypted by E_{K^*} , where $K^* = H(K, N, A)$. On decryption, the tag for the baseline AE is first recovered by $E_{K^*}^{-1}$. Unlike CTX, this mechanism enables the use of the baseline AE in a black-box way. Fig. 9 illustrates PACT and comPACT. Fig. 8 shows their pseudocodes. comPACT is identical to PACT except comPACT fixes the AD input of Π to ε . The technique is similar to CTY [BH24] and improves efficiency as AD is only processed once by H . Let $\text{PACT}[\Pi, E, H]$ and $\text{comPACT}[\Pi, E, H]$ denote the PACT and comPACT transforms applied to Π . Bhattacharjee *et al.* proved (MR-)PRIV and (MR-)AUTH of $\text{PACT}[\Pi, E, H]$ assuming collision resistance of H , E being an ideal cipher, and (MR-)PRIV and (MR-)AUTH of Π . Similarly, they proved PRIV and AUTH of $\text{comPACT}[\Pi, E, H]$ under the same assumptions on E and H as for PACT, assuming Π is AE-secure. They proved both PACT and comPACT

$\text{PACT}.\mathcal{E}[\Pi, \mathcal{E}, \mathcal{H}](K, N, A, M)$ 1. $(\tilde{C}, \tilde{T}) \leftarrow \Pi.\mathcal{E}(K, N, \boxed{A}, M)$ 2. $K^* \leftarrow \mathcal{H}(K, N, A)$ 3. $T \leftarrow \mathcal{E}(K^*, \tilde{T})$ 4. return $(\tilde{C}, T)$	$\text{PACT}.\mathcal{D}[\Pi, \mathcal{E}, \mathcal{H}](K, N, A, (\tilde{C}, T))$ 1. $K^* \leftarrow \mathcal{H}(K, N, A)$ 2. $\tilde{T} \leftarrow \mathcal{E}^{-1}(K^*, T)$ 3. $M \leftarrow \Pi.\mathcal{D}(K, N, \boxed{A}, (\tilde{C}, \tilde{T}))$ 4. return M
---	--

Fig. 8: PACT and comPACT transforms. comPACT is obtained by replacing the gray-boxed AD, $\boxed{A}$, with an empty string ε .

are CMT secure with the same assumption as above, with an additional mild assumption on Π , called Ciphertext Collision (CC) [BBD24].

As in the case of chaSIV, PACT and comPACT employ the mixed model for their primitives. We show attacks against any instantiation of them.




Fig. 9: (Left) PACT and (Right) comPACT.

Privacy attack against PACT. Let Π be a variant of SIV. For encrypting a tuple (N, A, M) , Π first performs the MAC computation to determine the tag $\tilde{T}$ as a (MAC) function of (K, N, A, M) . Next, it derives the internal key as $K^* = \mathcal{H}(K, N, A)$. After parsing M into τ -bit blocks, the i -th ciphertext block is obtained as $\tilde{C}[i] = M[i] \oplus \mathcal{E}(K^*, \tilde{T} \oplus \langle i-1 \rangle_\tau)$. As this Π follows the template of SIV, it is provably secure by adequately instantiating the MAC function and $\mathcal{E}$. Note that $\mathcal{E}$'s instantiation must be reasonable for an ideal cipher. We observe that $\tilde{C}[1] = M[1] \oplus \mathcal{E}(K^*, \tilde{T} \oplus \langle 0 \rangle_\tau (= \tilde{T})) = M[1] \oplus T$ holds for any valid encryption transcript $(N, A, M, (\tilde{C}, T))$ of PACT for $|M| \geq \tau$. This observation implies a privacy break and plaintext recovery of $M[1]$ if M is unknown. Full plaintext recovery occurs if $|M| \leq \tau$.

Authenticity attack against PACT. Consider a variant of GCM using Universal hash-then-Encrypt. Let Π be an AE scheme based on a nonce-based encryption scheme $\Gamma : \mathcal{K} \times \mathcal{N}' \times \mathcal{M} \rightarrow \mathcal{M}$ for $\mathcal{N}' = \mathcal{N} \times \mathcal{AD}$, a block cipher $\mathcal{E} : \mathcal{K}_b \times \mathcal{M}_b \rightarrow \mathcal{M}_b$, and GHASH defined over $\text{GF}(2^\tau)$. Let N' denote the tuple (N, A) and let

$\Pi.\mathcal{E}(K, N, A, M) = (\tilde{C}, \tilde{T})$, where

$$\tilde{C} = \Gamma(K, N', M), \quad \tilde{T} = E^{-1}(K^*, V), \quad V = \text{GHASH}(L, (A, \tilde{C})),$$

and $L = \text{KDF}(K)$ for some $\text{KDF} : \mathcal{K} \rightarrow \{0, 1\}^\tau$, $K^* = H(K, N, A)$ as in the case of privacy attack.⁷ The black-box security of Π is easy to prove by adequately instantiating KDF , Γ and E . If we want Π to be a nonce-based AE, Γ can be a CTR mode with nonce N' . If we want Π to be an MRAE, so that it is a meaningful counterexample of PACT, Γ can be realized by a tweakable wide-block cipher $\tilde{E} : \mathcal{K} \times \mathcal{N}' \times \mathcal{M} \rightarrow \mathcal{M}$, such as HCTR(2) [WFW05, CHB21], with tweak N' ; $\Gamma(K, N', M) = \tilde{E}(K, N', M) = \tilde{C}$. Such Π will ensure MR-PRIV as any distinct input (N, A, M) will produce a random $\tilde{C}$. This Π will also ensure MR-AUTH for its Enc-then-MAC structure. After transformation, $\text{PACT}[\Pi, E, H]$ computes T as $E(K^*, \tilde{T}) = E(K^*, E^{-1}(K^*, V)) = V$ thus it reveals V , the output of GHASH. The knowledge of V allows us to derive the GHASH key L . The attack is as follows.

1. Encrypt (N, A, M) and receive $(\tilde{C}, T) = \text{PACT}[\Pi, E, H].\mathcal{E}(K, N, A, M)$.
2. From $T = V = \text{GHASH}(L, (A, \tilde{C}))$, recover L .
3. Choose $\tilde{C}' \neq \tilde{C}$ such that $\text{GHASH}(L, (A, \tilde{C})) = \text{GHASH}(L, (A, \tilde{C}'))$. Decrypt $(N, A, (\tilde{C}', T))$ which will be accepted with probability 1.

Finding $\tilde{C}'$ in Step 3 is straightforward, see *e.g.* [Jou06, FIM20].

Attacking comPACT. The above privacy and authenticity attacks also work for comPACT, with a minor modification to the authenticity attack.

4.7 Key Recovery Attacks

Our attacks shown above exploit AE schemes that appear natural due to their use of strong primitives in a generic composition. We show that if we adopt slightly artificial forms of (authenticated) encryption, some transforms allow even total key recovery. The baseline encryption schemes we used are provably secure on their own; hence, they are theoretically valid counterexamples.

The trick is simple; we use a counter mode with an additional (and unnecessary) Key Blinding (KB) mechanism. See below for details.

Key recovery against AEaH. The most straightforward key-recovery attack is against AEaH. The privacy attack against AEaH uses CTR mode $\tilde{C}[i] \leftarrow M[i] \oplus H(K, N, \langle i \rangle)$. For key recovery, we instead use CTR with key blinding, dubbed CTRKB, such that $\tilde{C}[i] \leftarrow M[i] \oplus H(K, N, \langle i \rangle) \oplus K$. The privacy attack is converted into a key recovery attack; K is recovered by taking the sum

⁷ While we try to mimic GCM in its use of GHASH, feeding A to GHASH may be redundant as K^* already contains information on A . Variations of the attack are possible when GHASH takes only $\tilde{C}$ or takes $\tilde{C}$ and a constant instead of A .

$\tilde{C}[1] \oplus M[1] \oplus T$. It is reasonable to assume that CTRKB maintains standard model security if the inner H is a well-analyzed cryptographic hash function such as SHA-2/3, which is also a reasonable candidate for RO instantiation.⁸

Key recovery against EtH. We modify the PRIV attack so that it uses a variant of CTRKB such that $C[i] = \text{msb}_{|M[i]|}(H(K, \text{cst}, N, \langle i-1 \rangle_c)) \oplus K \oplus M[i]$. The attack is obtained by changing the last step of the PRIV attack, as in the case of AEaH.

Key recovery against AEth. We modify the PRIV attack so that it uses a variant of CTRKB such that $C[i] = M[i] \oplus H(K, N, A, V \oplus \langle i-1 \rangle) \oplus K$.

Key recovery against PACT/comPACT. We modify the PRIV attack so that it uses a variant of CTRKB such that $\tilde{C}[i] = M[i] \oplus E_{K^*}(T \oplus \langle i-1 \rangle) \oplus K$.

Key recovery against chaSIV. Our key recovery attack against chaSIV is more involved than others. The simple key blinding technique seems inapplicable, since the F output is secret and the H algorithm should not have an exploitable structure, as originally assumed to be a monolithic RO. Let π be a κ -bit cryptographic permutation for even κ . Let $r = \kappa$. We assume $\mathcal{M} = \{0, 1\}^{\kappa/2}$ for simplicity. Recall $|K| = |K_e| = \kappa$. Let $\text{cst}_1 \in \mathcal{AD}$ and $\text{cst}_2 \in \mathcal{M}$ be some constant bit strings used by the PRIV attack. Let $F : \mathcal{K} \times \mathcal{N} \rightarrow \mathcal{K}$ and $\Gamma : \mathcal{K} \times \mathcal{N} \times \mathcal{M} \rightarrow \mathcal{M}$ be defined as follows. First, we parse the key as $K = (K[0] \parallel K[1])$ where $|K[0]| = |K[1]| = \kappa/2$. Γ is a variant of CTR mode using π with truncation. The encryption is

$$\Gamma.\mathcal{E}(K_e, N_e, M) = C = \text{msb}_{\kappa/2}(\pi(N_e \oplus K_e)) \oplus M.$$

The decryption of Γ is analogously defined. We next design F so that it behaves badly when combined with the above Γ . In more detail, we define

$$F(N, K) = K_e = \text{msb}_{\kappa}(V_N) \oplus S[\text{msb}_1(N)],$$

where $S[i] = \pi^{-1}(K[i] \parallel 0^{\kappa/2})$ for $i \in \{0, 1\}$ and $V_N = H(K, N, \text{cst}_1, \text{cst}_2)$.

We show this combination indeed leaks the key. For an encryption query $(N, A = \text{cst}_1, M = \text{cst}_2)$ with $\text{msb}_1(N) = 0$, the adversary receives (C, T) and recovers $K[0]$ since

$$\begin{aligned} C \oplus M &= \text{msb}_{\kappa/2}(\pi(N_e \oplus K_e)) \\ &= \text{msb}_{\kappa/2}(\pi(\underbrace{\text{msb}_{\kappa}(V_N)}_{N_e} \oplus \underbrace{S[0] \oplus \text{msb}_{\kappa}(V_N)}_{K_e})) \\ &= \text{msb}_{\kappa/2}(\pi(\pi^{-1}(K[0] \parallel 0^{\kappa/2}))) \\ &= \text{msb}_{\kappa/2}(K[0] \parallel 0^{\kappa/2}) \\ &= K[0]. \end{aligned}$$

⁸ In a general context, it may happen that additional key blinding collapses original security. A typical example is CTR built on an Even-Mansour cipher.

Another encryption query with N' such that $\text{msb}_1(N') = 1$ will reveal $K[1]$. Thus, using two encryption queries with different MSBs, the full K will be recovered.

These “contrived” functions F and Γ themselves would be secure if π and H are reasonable instantiations of a public random permutation and an RO. Let us give a brief analysis. Our F takes the sum of $S[\text{msb}_1(N)]$ and V_N , where the latter is deemed as an output of a PRF $H(K, *, \text{cst}_1, \text{cst}_2)$. Assuming π and H are algorithmically different, adding $S[\text{msb}_1(N)]$ to V_N should cause no problem on PRF security. A provable security bound can be easily obtained if we make idealized assumptions about them. Our Γ is effectively a CTR mode of truncated permutation π with input masking by the key. If π is a public random permutation, Γ has $(\kappa/2 - \log_2 \kappa)$ -bit security [HNS16]. We remark that the design of this counterexample is subtle: we cannot use Even-Mansour for Γ because in that case K_e will remain in the keystream $M \oplus C$, hence it will not be PRIV secure. If we skip truncation for Γ , it will be insecure, as an inverse permutation query will reveal K_e . Moreover, if we use the entire K instead of half to generate $S[i]$, it reveals only part of K , and the rest remains unrevealed.

This example of Γ should look somewhat unnatural. There would be more straightforward, intuitive examples. However, our example works irrespective of the instantiation of H . Thus, we think it is a meaningful attack against an instantiation of `chaSIV`.

5 Provably Secure Generic Transforms

5.1 Fixing CK25 transforms

Our attacks exploit the gap between the mixed model and its instantiation. A natural next step is to address the gap; however, this is inherently difficult if we keep the mixed model.

As mentioned in Sect. 3, any transform solely relying on the standard model is inherently safe. Following this principle, we present how to modify the broken CK25 transforms to obtain standard model proofs.⁹ Note that uninstantiable transforms (`AEaH`, `EtH`, `AEtH`, `PACT`, and `comPACT`) cannot have standard model proofs, hence modifications are necessary for them. To do so, we keep changes minimal and, in particular, consider whether the original scheme requires nonce-based rekeying, since this can be a key factor in determining overall performance, as observed by CK25. Concretely, we present fixed versions for `AEaH`, `AEtH`, `EtH`, and `chaSIV`. For `HtAE`, we show that standard-model security is achievable without changing the algorithm. As mentioned in the Introduction, two of them (`AEaH-KS` and `AEtH-KS`) are the first transforms that are

1. strictly generic,
2. context committing,

⁹ We do not consider `PACT` and `comPACT`, as they are not strictly generic due to their requirement of an additional property for AE and potential revisions mentioned by the designers (private communication).

3. free from nonce-based rekeying, and
4. (entirely) standard-model secure

known in the literature (See also Tables 1 and 2). In other words, we salvage AEaH and AEtH, those equipped with Properties 1 to 3 but uninstantiable, and add Property 4, keeping Property 3.

Random IVs do not help. Our attacks exploit the fact that the broken schemes allow the baseline AE to fully simulate the computation performed by the transform. To thwart these attacks, it looks plausible to generate a random IV, IV , inside the transform and give IV *only* to the hash function. However, this generally does not work, since a random IV is part of the encryption output. The adversary can learn the random IV used by an encryption query and provide it to AE (*e.g.*, as part of AD) via another query. Some broken examples could be shown, for instance, taking AEaH.

5.2 Fixing AEaH

We present how to fix AEaH. For this, we assume H has an explicit key input $\mathcal{K}$ (for simplicity: the key space can be different from that of AE), making $H : \mathcal{K} \times \{0, 1\}^* \rightarrow \{0, 1\}^t$. We introduce a PRG $G : \mathcal{K} \rightarrow \mathcal{K} \times \mathcal{K}$. The fix is consistent and simple: we apply $G(K) = (K_1, K_2)$ and use K_1 for Π and K_2 for H . We call this technique *Key separation* (KS). Since KS is needed only once, it is preferable to nonce-based rekeying. The fixed scheme, AEaH-KS, is presented by Fig. 10 (top). We prove that AEaH-KS is secure assuming G is an RCR-PRG and H is a CR-PRF.

Theorem 1. *Let $\mathcal{A}_{\text{priv}}$ and $\mathcal{A}_{\text{auth}}$ be $(q_e, \sigma_e, \mathbf{tm})$ -PRIV and $(q_e, q_d, \sigma_e, \sigma_d, \mathbf{tm})$ -AUTH adversaries against AEaH-KS $[\Pi, G, H]$. Let $q = q_e + q_d$ and $\sigma = \sigma_e + \sigma_d$. Then*

$$\begin{aligned} \text{Adv}_{\text{AEaH-KS}[\Pi, G, H]}^{\text{priv}}(\mathcal{A}_{\text{priv}}) &\leq \text{Adv}_G^{\text{prg}}(\mathcal{A}_G) + \text{Adv}_H^{\text{prf}}(\mathcal{A}_H) + \text{Adv}_{\Pi}^{\text{priv}}(\mathcal{A}_{\Pi}) \\ \text{Adv}_{\text{AEaH-KS}[\Pi, G, H]}^{\text{auth}}(\mathcal{A}_{\text{auth}}) &\leq \text{Adv}_G^{\text{prg}}(\mathcal{A}'_G) + \text{Adv}_H^{\text{prf}}(\mathcal{A}'_H) + \text{Adv}_{\Pi}^{\text{auth}}(\mathcal{A}'_{\Pi}) + \frac{q_d}{2^t}, \end{aligned}$$

holds, where $\mathcal{A}_G$ and $\mathcal{A}'_G$ are $\mathbf{tm}'$ -PRG adversaries. $\mathcal{A}_H$ is a $(q_e, \sigma_e, \mathbf{tm}')$ -PRF adversary, $\mathcal{A}_{\Pi}$ is a $(q_e, \sigma_e, \mathbf{tm}')$ -PRIV adversary, $\mathcal{A}'_H$ is a $(q, \sigma, \mathbf{tm}')$ -PRF adversary, and $\mathcal{A}'_{\Pi}$ is a $(q_e, q_d, \sigma_e, \sigma_d, \mathbf{tm}')$ -AUTH adversary.

Theorem 2. *Let $\mathcal{A}_{\text{cmt}}$ be $\mathbf{tm}$ -CMT adversary against AEaH-KS $[\Pi, G, H]$.*

$$\text{Adv}_{\text{AEaH-KS}[\Pi, G, H]}^{\text{cmt}}(\mathcal{A}_{\text{cmt}}) \leq \text{Adv}_H^{\text{cr}}(\mathcal{A}_H) + \text{Adv}_G^{\text{rcr}}(\mathcal{A}_G),$$

where $\mathcal{A}_H$ is a $\mathbf{tm}'$ -CR adversary and $\mathcal{A}_G$ is a $\mathbf{tm}'$ -RCR adversary.

AEaH-KS.$\mathcal{E}$[Π, G, H](K, N, A, M) 1. $(K_1, K_2) \leftarrow G(K)$ 2. $C \leftarrow \Pi.\mathcal{E}(K_1, N, \varepsilon, M)$ 3. $T \leftarrow H(K_2, N, A)$ 4. return (C, T)	AEaH-KS.$\mathcal{D}$[Π, G, H]($K, N, A, (C, T)$) 1. $(K_1, K_2) \leftarrow G(K)$ 2. $M \leftarrow \Pi.\mathcal{D}(K_1, N, \varepsilon, C)$ 3. if $M = \perp \vee H(K_2, N, A) \neq T$: 4. return $\perp$ 5. return M
AEtH-KS.$\mathcal{E}$[Π, G, H](K, N, A, M) 1. $(K_1, K_2) \leftarrow G(K)$ 2. $C \leftarrow \Pi.\mathcal{E}(K_1, N, A, M)$ 3. $T \leftarrow H(K_2, N, A, \text{lsb}_r(C))$ 4. return (C, T)	AEtH-KS.$\mathcal{D}$[Π, G, H]($K, N, A, (C, T)$) 1. $(K_1, K_2) \leftarrow G(K)$ 2. $T^* \leftarrow H(K_2, N, A, \text{lsb}_r(C))$ 3. if $T^* \neq T$: return $\perp$ 4. $M \leftarrow \Pi.\mathcal{D}(K_1, N, A, C)$ 5. return M
EtH-KS.$\mathcal{E}$[Γ, G, H](K, N, A, M) 1. $(K_1, K_2) \leftarrow G(K)$ 2. $C \leftarrow \Gamma.\mathcal{E}(K_1, N, M)$ 3. $T \leftarrow H(K_2, N, A, C)$ 4. return (C, T)	EtH-KS.$\mathcal{D}$[Γ, G, H]($K, N, A, (C, T)$) 1. $(K_1, K_2) \leftarrow G(K)$ 2. $T^* \leftarrow H(K_2, N, A, C)$ 3. if $T^* \neq T$: return $\perp$ 4. $M \leftarrow \Gamma.\mathcal{D}(K_1, N, C)$ 5. return M

Fig. 10: Fixed versions of CK25 transforms by means of key separation. (Top) AEaH-KS, (Middle) AEtH-KS, and (Bottom) EtH-KS.

5.3 Proofs of Theorems 1 and 2 (AEaH-KS)

Proof of Theorem 1. Let $\$G$ be a uniform random source over $\mathcal{K} \times \mathcal{K}$ and let $R_H : \{0, 1\}^* \rightarrow \{0, 1\}^t$ be a variable-input-length URF. Let H_K denote the oracle that returns $H(K, N, A)$ with uniform key $K \leftarrow \$\mathcal{K}$ on input $(N, A) \in \mathcal{N} \times \mathcal{AD}$. Let Θ_1 , Θ_2 , and Θ_3 denote AEaH-KS[Π, G, H], AEaH-KS[$\Pi, \$G, H$], and AEaH-KS[$\Pi, \G, R_H] respectively. We have

$$\begin{aligned}
& \mathbf{Adv}_{\text{AEaH-KS}[\Pi, G, H]}^{\text{priv}}(\mathcal{A}_{\text{priv}}) \\
& \leq \mathbf{Adv}_{\Theta_1, \mathcal{E}, \Theta_2, \mathcal{E}}^{\text{ind}}(\mathcal{A}_{\text{priv}}) + \mathbf{Adv}_{\Theta_2, \mathcal{E}, \Theta_3, \mathcal{E}}^{\text{ind}}(\mathcal{A}_{\text{priv}}) + \mathbf{Adv}_{\Theta_3, \mathcal{E}, \$}^{\text{ind}}(\mathcal{A}_{\text{priv}}) \\
& \leq \mathbf{Adv}_{G, \$G}^{\text{ind}}(\mathcal{A}_G) + \mathbf{Adv}_{H_K, R_H}^{\text{ind}}(\mathcal{A}_H) + \mathbf{Adv}_{\Theta_3, \mathcal{E}, \$}^{\text{ind}}(\mathcal{A}_{\text{priv}}) \\
& \leq \underbrace{\mathbf{Adv}_G^{\text{prg}}(\mathcal{A}_G)}_{\text{by definition}} + \underbrace{\mathbf{Adv}_H^{\text{prf}}(\mathcal{A}_H)}_{\text{by definition}} + \mathbf{Adv}_{\Pi}^{\text{priv}}(\mathcal{A}'_{\Pi}).
\end{aligned}$$

Here, $\mathbf{Adv}_{\Theta_3, \$}^{\text{ind}}(\mathcal{A}_{\Theta_3}) \leq \mathbf{Adv}_{\Pi, \mathcal{E}, \$}^{\text{ind}}(\mathcal{A}'_{\Pi})$ holds as the T part in any encryption response is always independently uniform and random for both games (because $T = R_H(N, A)$ in Θ_3 and nonce does not repeat), thus not contributing to the advantage. This proves the privacy claim.

For authenticity bound, we have

$$\begin{aligned}
& \mathbf{Adv}_{\text{AEaH-KS}[\Pi, G, H]}^{\text{auth}}(\mathcal{A}_{\text{auth}}) \\
& \leq \mathbf{Adv}_{\Theta_1, \Theta_2}^{\text{ind}}(\mathcal{A}'_{\text{auth}}) + \mathbf{Adv}_{\Theta_2}^{\text{auth}}(\mathcal{A}_{\text{auth}}) \\
& \leq \mathbf{Adv}_{\Theta_1, \Theta_2}^{\text{ind}}(\mathcal{A}'_{\text{auth}}) + \mathbf{Adv}_{\Theta_2, \Theta_3}^{\text{ind}}(\mathcal{A}''_{\text{auth}}) + \mathbf{Adv}_{\Theta_3}^{\text{auth}}(\mathcal{A}_{\text{auth}}) \\
& \leq \mathbf{Adv}_{G, \mathbb{S}_G}^{\text{ind}}(\mathcal{A}_G) + \mathbf{Adv}_{H_K, R_H}^{\text{ind}}(\mathcal{A}_H) + \mathbf{Adv}_{\Theta_3}^{\text{auth}}(\mathcal{A}_{\text{auth}}) \\
& = \mathbf{Adv}_G^{\text{prg}}(\mathcal{A}_G) + \mathbf{Adv}_H^{\text{prf}}(\mathcal{A}_H) + \mathbf{Adv}_{\Theta_3}^{\text{auth}}(\mathcal{A}_{\text{auth}})
\end{aligned} \tag{1}$$

The last term $\mathbf{Adv}_{\Theta_3}^{\text{auth}}(\mathcal{A}_{\text{auth}}) := \mathbf{Adv}_{\Theta_3, (\Theta_3, \mathcal{E}, \perp)}^{\text{ind}}(\mathcal{A}_{\text{auth}})$ needs further hybrids. For clarity, we introduce the following attributes of decryption queries. A decryption query (N', A', C', T') is said to be Π -trivial if there exists an encryption query-response tuple (N, A, M, C, T) such that $(N' = N, A' \neq A, C' = C)$, and H -trivial if there exists an encryption query-response tuple (N, A, M, C, T) such that $(N' = N, A' = A, C' \neq C)$. We remark that a trivial decryption query is prohibited in its entirety: thus, if $(N' = N, A' = A, C' = C)$ holds, it means $T' \neq T$. Such a query will never be successful (*i.e.*, the response is always $\perp$). Therefore, we can exclude such decryption queries. A decryption query which is not Π -trivial is called Π -nontrivial, including *e.g.*, the case that encryption queries do not use N' . Clearly, the sets of Π -trivial queries and H -trivial queries are exclusive, and the latter is a subset of Π -nontrivial queries.

For any keyed function (or oracle) $F_K : \mathcal{X} \rightarrow \mathcal{Y}$, let $\text{vrf}[F_K] : \mathcal{X} \times \mathcal{Y} \rightarrow \{\top, \perp\}$ be the verification oracle defined as $\text{vrf}[F_K](X, Y) = \top$ if $F_K(X) = Y$ and $\text{vrf}[F_K](X, Y) = \perp$ if $F_K(X) \neq Y$.

We decompose the decryption routine into two sub-oracles, O_1^D and O_2^D , where O_1^D performs AE decryption and O_2^D performs verification regarding H , namely $\text{vrf}[H_K]$. The response to a decryption query is determined by calling the sub-oracles, and is a tuple rather than a single variable (either the decrypted message or $\perp$). Let $\Pi.\mathcal{D}'_K$ denote a decryption oracle with AD fixed to empty, *i.e.*, $\Pi.\mathcal{D}'_K(N, C) = \Pi.\mathcal{D}_K(N, \varepsilon, C)$. The original decryption routine receiving a query (N, A, C, T) is reinterpreted as follows, assuming $O_1^D = \Pi.\mathcal{D}'_{K_1}$ and $O_2^D = \text{vrf}[H_{K_2}]$.

1. $O_2^D(N, A, T) \rightarrow Y_2$
2. If $Y_2 = \perp$, return $\perp$
3. $O_1^D(N, C) \rightarrow Y_1$
4. Return Y_1

Observe that the adversary needs to break both O_1^D and O_2^D . Keeping this feature, we modify the routine as follows.

1. If (N, A, C, T) is Π -trivial, $O_2^D(N, A, T) \rightarrow Y_2$ and return (ε, Y_2)
2. Else $O_1^D(N, C) \rightarrow Y_1$
3. If $Y_1 \neq \perp$, $O_2^D(N, A, T) \rightarrow Y_2$ and return (Y_1, Y_2)
4. Else return $(\perp, \perp)$

It is straightforward to see that this modification does not reduce the chance of win: the output is a tuple (Y_1, Y_2) in the modified routine, and the original response is always recovered from the output tuple, and as in the original routine, the adversary is successful only if it breaks both O_1^D and O_2^D , including the trivial win against O_i^D . Note that a decryption query leading to a trivial win against both sub-oracles is impossible by definition. With this modification, any Π -nontrivial query will be first processed by O_1^D . We write (O_1^D, O_2^D) to mean the decryption oracle. Let $\Theta_3^1 = (\Theta_3.\mathcal{E}, \Theta_3^1.\mathcal{D})$, where $\Theta_3^1.\mathcal{D} = (\Pi.\mathcal{D}', \text{vrf}[\mathbf{R}_H])$. Similarly, let $\Theta_3^2 = (\Theta_3.\mathcal{E}, \Theta_3^2.\mathcal{D})$, where $\Theta_3^2.\mathcal{D} = (\perp, \text{vrf}[\mathbf{R}_H])$ and let $\Theta_3^3 = (\Theta_3.\mathcal{E}, \Theta_3^3.\mathcal{D})$, where $\Theta_3^3.\mathcal{D} = (\perp, \perp)$. Clearly, Θ_3^1 and Θ_3^3 (in the modified routine) are the counterparts of Θ_3 and $(\Theta_3.\mathcal{E}, \perp)$ (in the original routine). Since the modified routine augments the original decryption oracle response,

$$\mathbf{Adv}_{\Theta_3, (\Theta_3.\mathcal{E}, \perp)}^{\text{ind}}(\mathcal{A}_{\text{auth}}) \leq \mathbf{Adv}_{\Theta_3^1, \Theta_3^2}^{\text{ind}}(\mathcal{A}_{\text{auth}}) + \mathbf{Adv}_{\Theta_3^2, \Theta_3^3}^{\text{ind}}(\mathcal{A}_{\text{auth}}) \quad (2)$$

holds. For bounding $\mathbf{Adv}_{\Theta_3^1, \Theta_3^2}^{\text{ind}}(\mathcal{A}_{\text{auth}})$, we observe that distinguishing is possible only by a Π -nontrivial query, since a Π -trivial query is processed by $\mathbf{R}_H$ for both worlds. Moreover, once the adversary receives $Y_1 \neq \perp$, this deterministically indicates that the world is Θ_3^1 . Once the adversary receives $Y_1 = \perp$ the response is identical, $(\perp, \perp)$. That is, Y_2 does not contribute to the advantage. Thus,

$$\mathbf{Adv}_{\Theta_3^1, \Theta_3^2}^{\text{ind}}(\mathcal{A}_{\text{auth}}) \leq \mathbf{Adv}_{\Pi}^{\text{auth}}(\mathcal{A}'_{\Pi}). \quad (3)$$

For bounding $\mathbf{Adv}_{\Theta_3^2, \Theta_3^3}^{\text{ind}}(\mathcal{A}_{\text{auth}})$, we observe that a Π -nontrivial query will not help as it results in $(\perp, \perp)$ response for both worlds. Thus, Π -trivial queries, which do not contain H -trivial queries, only help. This implies that the game reduces to a MAC unforgeability game involving $\mathbf{R}_H$ with (N, A) as the message. Since $\mathbf{R}_H$ is completely random, the forgery probability is $q_d/2^t$: thus,

$$\mathbf{Adv}_{\Theta_3^2, \Theta_3^3}^{\text{ind}}(\mathcal{A}_{\text{auth}}) \leq \frac{q_d}{2^t}. \quad (4)$$

Combining Eqs. (1), (2), (3), and (4), we conclude the proof of Theorem 1.

Proof of Theorem 2. Suppose $Z = ((K, N, A, M), (K', N', A', M'), (C, T))$ be the final output of $\mathcal{A}_{\text{cmt}}$ with the pair of corresponding ciphertext and tag (C, T) . Let $(K_1, K_2) = \mathbf{G}(K)$ and $(K'_1, K'_2) = \mathbf{G}(K')$.

Assuming Z is valid, we consider two cases. First, if $K = K'$, we have $(K_1, K_2) = (K'_1, K'_2)$. When $(N, A) = (N', A')$ holds, it implies $M = M'$, hence Z is not valid. For Z to be valid, $\mathbf{H}(K_2, N, A) = \mathbf{H}(K'_2, N', A') (= T)$ must hold for some $(N, A) \neq (N', A')$, implying a collision of $\mathbf{H}$ outputs. Second, suppose $K \neq K'$. The subcase $[K_2 = K'_2]$ reduces to a right-half collision on $\mathbf{G}$. The subcase $[K_2 \neq K'_2]$ implies a collision on $\mathbf{H}$. Overall, for any case, the adversary needs to find either a collision on $\mathbf{H}$ or a right-half collision on $\mathbf{G}$, or more concretely,

$$\begin{aligned} \mathbf{Adv}_{\text{AEaH-KS}[\Pi, \mathbf{G}, \mathbf{H}]}^{\text{cmt}}(\mathcal{A}_{\text{cmt}}) &\leq \max\{\mathbf{Adv}_{\mathbf{H}}^{\text{cr}}(\mathcal{A}_{\mathbf{H}}), \mathbf{Adv}_{\mathbf{G}}^{\text{rcr}}(\mathcal{A}_{\mathbf{G}})\}, \\ &\leq \mathbf{Adv}_{\mathbf{H}}^{\text{cr}}(\mathcal{A}_{\mathbf{H}}) + \mathbf{Adv}_{\mathbf{G}}^{\text{rcr}}(\mathcal{A}_{\mathbf{G}}), \end{aligned}$$

which concludes the proof.

5.4 Fixing AEtH

We can fix AEtH in the same manner as AEaH. We apply KS to AEtH and obtain AEtH-KS as shown in Fig. 10 (middle). We prove its security assuming G is an RCR-PRG and H is a CR-PRF.

Theorem 3. *Let $\mathcal{A}_{\text{priv}}$ and $\mathcal{A}_{\text{auth}}$ be $(q_e, \sigma_e, \text{tm})$ -MR-PRIV and $(q_e, q_d, \sigma_e, \sigma_d, \text{tm})$ -MR-AUTH adversaries against AEtH-KS $[\Pi, G, H]$. Let $q = q_e + q_d$ and $\sigma = \sigma_e + \sigma_d$. We have*

$$\text{Adv}_{\text{AEtH-KS}[\Pi, G, H]}^{\text{mr-priv}}(\mathcal{A}_{\text{priv}}) \leq \text{Adv}_G^{\text{prg}}(\mathcal{A}_G) + \text{Adv}_H^{\text{prf}}(\mathcal{A}_H) + \text{Adv}_{\Pi}^{\text{mr-priv}}(\mathcal{A}_{\Pi}) + \frac{q_e^2}{2^{\tau+1}},$$

$$\text{Adv}_{\text{AEtH-KS}[\Pi, G, H]}^{\text{mr-auth}}(\mathcal{A}_{\text{auth}}) \leq \text{Adv}_G^{\text{prg}}(\mathcal{A}'_G) + \text{Adv}_{\Pi}^{\text{mr-auth}}(\mathcal{A}'_{\Pi}),$$

where $\mathcal{A}_G$ and $\mathcal{A}'_G$ are tm' -PRG adversaries, $\mathcal{A}_H$ is a $(q_e, \sigma_e, \text{tm}')$ -PRF adversary, $\mathcal{A}_{\Pi}$ is a $(q_e, \sigma_e, \text{tm}')$ -MR-PRIV adversary, and $\mathcal{A}'_{\Pi}$ is a $(q_e, q_d, \sigma_e, \sigma_d, \text{tm}')$ -MR-AUTH adversary.

Theorem 4. *Let $\mathcal{A}_{\text{cmt}}$ be a tm -CMT adversary against AEtH-KS $[\Pi, G, H]$.*

$$\text{Adv}_{\text{AEtH-KS}[\Pi, G, H]}^{\text{cmt}}(\mathcal{A}_{\text{cmt}}) \leq \text{Adv}_H^{\text{cr}}(\mathcal{A}_H) + \text{Adv}_G^{\text{rcr}}(\mathcal{A}_G),$$

where $\mathcal{A}_H$ is a tm' -CR adversary and $\mathcal{A}_G$ is a tm' -RCR adversary.

5.5 Proofs of Theorems 3 and 4 (AEtH-KS)

Proof of Theorem 3. As in Theorem 1, the hybrid arguments that replace G with a random source $\$G$ and H with a URF R_H are trivial. We derive the PRIV bound for Θ_1 , which denotes AEtH $[\Pi, \$G, R_H]$. Let $\Theta_2 = \text{AEtH}[\Pi, \$G, R_H]$ denote a slight variant of Θ_1 whose encryption generates C in a completely random way irrespective of the input. For MR-PRIV, we have

$$\begin{aligned} \text{Adv}_{\Theta_1}^{\text{mr-priv}}(\mathcal{A}_{\text{priv}}) &\leq \text{Adv}_{\Theta_1, \mathcal{E}, \Theta_2, \mathcal{E}}^{\text{ind}}(\mathcal{A}_{\text{priv}}) + \text{Adv}_{\Theta_2, \mathcal{E}, \$}^{\text{ind}}(\mathcal{A}_{\text{priv}}) \\ &\leq \text{Adv}_{\Pi}^{\text{mr-priv}}(\mathcal{A}_{\Pi}) + \frac{q_e^2}{2^{\tau+1}}, \end{aligned} \tag{5}$$

where the first term of Eq. (5) holds since $T \leftarrow R_H(K_2, N, A, \text{lsb}_{\tau}(C))$ for $K_2 \leftarrow \$K$ holds for both worlds (thus the advantage will not increase if T is omitted from the response). The second term holds since, unless the $\text{lsb}_{\tau}(C)$ part does not collide, the outputs of both $\Theta_2, \mathcal{E}$ and $\$$ are equivalent, and C is random for both worlds. By adding computational terms with respect to G and H, this proves the MR-PRIV bound.

For the MR-AUTH bound, we define Θ_1 as AEtH $[\Pi, \$G, H]$. We also define Θ_2 as AEtH $[(\Pi, \mathcal{E}, \perp), \$G, H]$. Note that $\Theta_2, \mathcal{D}$ is equivalent to $\perp$ oracle as it always

returns $\perp$ irrespective of the computation of H during decryption, because Π is modified so that it always returns $\perp$ when decrypting. To distinguish Θ_1 from Θ_2 , the adversary needs to invoke $M \neq \perp$ at the last line of the decryption routine (Fig. 5), implying the authenticity break of Π . In addition, a decryption query implying a trivial win against Π , *i.e.*, (N, A, C, T) such that $\Pi.\mathcal{E}(K, N, A, M) \rightarrow C$ is already computed in an encryption query, will not win by definition. Thus we have

$$\mathbf{Adv}_{\Theta_1}^{\text{mr-auth}}(\mathcal{A}_{\text{auth}}) \leq \mathbf{Adv}_{\Theta_1, \Theta_2}^{\text{ind}}(\mathcal{A}'_{\text{auth}}) \leq \mathbf{Adv}_{\Pi}^{\text{mr-auth}}(\mathcal{A}'_{\Pi}).$$

Adding PRG and PRF terms to the above bound concludes the proof.

Proof of Theorem 4. Following the proof of Theorem 2, suppose $Z = ((K, N, A, M), (K', N', A', M'), (C, T))$ be a valid input tuple determined by the adversary, and the corresponding ciphertext-tag tuple (C, T) . Let $(K_1, K_2) = G(K)$ and $(K'_1, K'_2) = G(K')$. When $K = K'$, $(N, A) = (N', A')$ implies $M = M'$, thus Z is not valid (attack never succeeds). When $K = K'$ and $(N, A) \neq (N', A')$ hold, it implies H 's collision. When $K \neq K'$, if $K_2 \neq K'_2$ occurs we need a collision of H , and if $K_2 = K'_2$ occurs it implies a right-half collision of G . Thus, a valid Z either requires a collision on H or a right-half collision of G . This concludes the proof.

5.6 Fixing EtH

As for EtH, we apply the key separation. The resulting scheme, denoted by EtH-KS, is almost identical to the generic composition, namely Enc-then-MAC. Our analysis is almost identical to the existing security analysis of Enc-then-MAC: for black-box security by Namprempre *et al.* [NRS14], and for committing security by Struck and Weishäupl [SW24] and Bhaumik *et al.* [BCC⁺24].

Theorem 5. *Let $\mathcal{A}_{\text{priv}}$ and $\mathcal{A}_{\text{auth}}$ be $(q_e, \sigma_e, \text{tm})$ -PRIV and $(q_e, q_d, \sigma_e, \sigma_d, \text{tm})$ -AUTH adversaries against EtH-KS[Π, G, H]. Let $q = q_e + q_d$ and $\sigma = \sigma_e + \sigma_d$. We have*

$$\begin{aligned} \mathbf{Adv}_{\text{EtH-KS}[\Gamma, G, H]}^{\text{priv}}(\mathcal{A}_{\text{priv}}) &\leq \mathbf{Adv}_G^{\text{prg}}(\mathcal{A}_G) + \mathbf{Adv}_H^{\text{prf}}(\mathcal{A}_H) + \mathbf{Adv}_{\Gamma}^{\text{priv}}(\mathcal{A}_{\Gamma}), \\ \mathbf{Adv}_{\text{EtH-KS}[\Gamma, G, H]}^{\text{auth}}(\mathcal{A}_{\text{auth}}) &\leq \mathbf{Adv}_G^{\text{prg}}(\mathcal{A}'_G) + \mathbf{Adv}_H^{\text{prf}}(\mathcal{A}'_H) + \frac{q_d}{2^t}, \end{aligned}$$

where $\mathcal{A}_G$ and $\mathcal{A}'_G$ are tm' -PRG adversaries, $\mathcal{A}_H$ is a $(q_e, \sigma_e, \text{tm}')$ -PRF adversary, $\mathcal{A}'_H$ is a (q, σ, tm') -PRF adversary, and $\mathcal{A}_{\Gamma}$ is a $(q_e, \sigma_e, \text{tm}')$ -PRIV adversary.

Theorem 6. *Let $\mathcal{A}_{\text{cmt}}$ be a tm -CMT adversary against EtH-KS[Γ, G, H]. Then,*

$$\mathbf{Adv}_{\text{EtH-KS}[\Gamma, G, H]}^{\text{cmt}}(\mathcal{A}_{\text{cmt}}) \leq \mathbf{Adv}_H^{\text{cr}}(\mathcal{A}_H) + \mathbf{Adv}_G^{\text{rcr}}(\mathcal{A}_G)$$

holds, where $\mathcal{A}_H$ is a tm' -CR adversary and $\mathcal{A}_G$ is a tm' -RCR adversary.

Appendix A shows the proofs of Theorems 5 and 6.

5.7 Fixing chaSIV

Several options to fix chaSIV could be considered. Adding KS on top is probably the most intuitive option, though we propose a simpler, more efficient alternative. This reuses the original components and roughly preserves the original complexity; the nonce is processed by F and H in the original, while the fixed version processes it twice by F. Our fixed version, chaSIV-NKS (for Nonce-based Key Separation), is shown in Fig. 11.

Theorem 7. *Let $\mathcal{A}_{\text{priv}}$ and $\mathcal{A}_{\text{auth}}$ be $(q_e, \sigma_e, \text{tm})$ -MR-PRIV and $(q_e, q_d, \sigma_e, \sigma_d, \text{tm})$ -MR-AUTH adversaries against chaSIV-NKS $[\Gamma, F, H]$. Let $q = q_e + q_d$ and $\sigma = \sigma_e + \sigma_d$. Throughout the MR-PRIV game, let Q_e denote the number of distinct nonces among the encryption queries. We assume that i -th nonce is repeated R_i times for $i \in [1..Q_e]$. Let $R_e^* := \max_{i \in [1..Q_e]} R_i$. For the MR-AUTH game, we assume Γ is tidy. Let Q_{ed} denote the number of distinct nonces among the union of encryption and decryption queries. We define R_i for $i \in [1..Q_{ed}]$ as the number of queries using i -th nonce and define $R_{ed}^* := \max_{i \in [1..Q_{ed}]} R_i$.*

$$\begin{aligned} \text{Adv}_{\text{chaSIV-NKS}[\Gamma, F, H]}^{\text{mr-priv}}(\mathcal{A}_{\text{priv}}) &\leq \text{Adv}_F^{\text{prf}}(\mathcal{A}_F) + Q_e \cdot \text{Adv}_H^{\text{prf}}(\mathcal{A}_H) \\ &\quad + Q_e \cdot \text{Adv}_\Gamma^{\text{priv}}(\mathcal{A}_\Gamma) + \frac{\sum_{i \in [1..Q_e]} R_i^2}{2^r}, \\ \text{Adv}_{\text{chaSIV-NKS}[\Gamma, F, H]}^{\text{mr-auth}}(\mathcal{A}_{\text{auth}}) &\leq \text{Adv}_F^{\text{prf}}(\mathcal{A}'_F) + Q_{ed} \cdot \text{Adv}_H^{\text{prf}}(\mathcal{A}'_H) + \frac{q_d}{2^t}, \end{aligned}$$

where $\mathcal{A}_F$ is a $(2Q_e, \text{tm}')$ -PRF adversary, $\mathcal{A}'_F$ is a $(2Q_{ed}, \text{tm}')$ -PRF adversary, $\mathcal{A}_H$ is a $(R_e^*, \sigma_e, \text{tm}')$ -PRF adversary, $\mathcal{A}'_H$ is a $(R_{ed}^*, \sigma, \text{tm}')$ -PRF adversary, and $\mathcal{A}_\Gamma$ is a $(R_e^*, \sigma_e, \text{tm}')$ -PRIV adversary.

Theorem 8. *Let $\mathcal{A}_{\text{cmt}}$ be a tm -CMT adversary against chaSIV-NKS $[\Gamma, F, H]$. Then*

$$\text{Adv}_{\text{chaSIV-NKS}[\Gamma, F, H]}^{\text{cmt}}(\mathcal{A}_{\text{cmt}}) \leq \text{Adv}_H^{\text{cr}}(\mathcal{A}_H) + \text{Adv}_F^{\text{rcr}}(\mathcal{A}_F)$$

holds, where $\mathcal{A}_H$ is a tm' -CR adversary and $\mathcal{A}_F$ is a tm' -RCR adversary.

Appendix B shows the proofs of Theorems 7 and 8.

5.8 HtAE under the standard model

HtAE uses a hash function $H : \{0, 1\}^* \rightarrow \mathcal{K} \times \{0, 1\}^t$. For encryption input (N, A, M) , $\text{HtAE}.\mathcal{E}_K[\Pi, H]$ first computes $(K_e, T) \leftarrow H(K, N, A)$ and $C \leftarrow \Pi.\mathcal{E}_{K_e}(N, \varepsilon, M)$, and outputs (C, T) . The decryption is analogously defined, see Fig. 1. As described, we did not find attacks against the original HtAE, even though it still relies on the mixed model, *i.e.*, H is treated as an RO. We prove the standard model security of HtAE without changing the algorithm, but changing the assumption on H. In more detail, we assume $H : \mathcal{K} \times \mathcal{N} \times \mathcal{A} \rightarrow \mathcal{K} \times \{0, 1\}^t$ is a PRF with key $K \leftarrow \$\mathcal{K}$ and is right-half collision resistant. For clarity, we call this transform HtAE-ST.

chaSIV-NKS. $\mathcal{E}[\Gamma, F, H](K, N, A, M)$	chaSIV-NKS. $\mathcal{D}[\Gamma, F, H](K, N, A, (C, T))$
1. $K_e \leftarrow F(K, 0 \parallel N)$ 2. $K_h \leftarrow F(K, 1 \parallel N)$ 3. $T \leftarrow H(K_h, A, M)$ 4. $N_e \leftarrow \text{lsb}_r(T)$ 5. $C \leftarrow \Gamma.\mathcal{E}(K_e, N_e, M)$ 6. return (C, T)	1. $K_e \leftarrow F(K, 0 \parallel N)$ 2. $K_h \leftarrow F(K, 1 \parallel N)$ 3. $N_e \leftarrow \text{lsb}_r(T)$ 4. $M \leftarrow \Gamma.\mathcal{D}(K_e, N_e, C)$ 5. if $M = \perp$: return $\perp$ 6. $T^* \leftarrow H(K_h, A, M)$ 7. if $T^* \neq T$: return $\perp$ 8. return M

Fig. 11: A fixed version of chaSIV, chaSIV-NKS.

Theorem 9. Let $\mathcal{A}_{\text{priv}}$ and $\mathcal{A}_{\text{auth}}$ be $(q_e, \sigma_e, \text{tm})$ -PRIV and $(q_e, q_d, \sigma_e, \sigma_d, \text{tm})$ -AUTH adversaries against $\text{HtAE-ST}[\Pi, H]$. Let Q_d denote the number of distinct nonces in decryption queries. Let $q = q_e + q_d$ and $\sigma = \sigma_e + \sigma_d$. Then

$$\begin{aligned} \text{Adv}_{\text{HtAE-ST}[\Pi, H]}^{\text{priv}}(\mathcal{A}_{\text{priv}}) &\leq \text{Adv}_{\text{H}}^{\text{prf}}(\mathcal{A}_{\text{H}}) + q_e \text{Adv}_{\Pi}^{\text{priv}}(\mathcal{A}_{\Pi}), \\ \text{Adv}_{\text{HtAE-ST}[\Pi, H]}^{\text{auth}}(\mathcal{A}_{\text{auth}}) &\leq \text{Adv}_{\text{H}}^{\text{prf}}(\mathcal{A}'_{\text{H}}) + Q_d \text{Adv}_{\Pi}^{\text{auth}}(\mathcal{A}'_{\Pi}) + \frac{q_d}{2^t} \end{aligned}$$

holds, where $\mathcal{A}_{\text{H}}$ is a $(q_e, \sigma_e, \text{tm}')$ -PRF adversary, $\mathcal{A}_{\Pi}$ is a $(1, \sigma_e, \text{tm}')$ -PRIV adversary, $\mathcal{A}'_{\text{H}}$ is a (q, σ, tm') -PRF adversary, and $\mathcal{A}'_{\Pi}$ is a $(1, q_d, \sigma_e, \sigma_d, \text{tm}')$ -AUTH adversary.

Theorem 10. Let $\mathcal{A}_{\text{cmt}}$ be a tm -CMT adversary against $\text{HtAE-ST}[\Pi, H]$. Then

$$\text{Adv}_{\text{HtAE-ST}[\Pi, H]}^{\text{cmt}}(\mathcal{A}_{\text{cmt}}) \leq \text{Adv}_{\text{H}}^{\text{rcf}}(\mathcal{A}_{\text{H}})$$

holds, where $\mathcal{A}_{\text{H}}$ is a tm' -RCR adversary.

Theorem 10 is identical to the original CMT proof [CK25, Theorem 3]. Appendix C shows the proofs of Theorems 9 and 10.

5.9 On CTX and CTY transforms

Despite their use of a mixed model, we do not find instantiation-based attacks against CTX and CTY. Since they are outside our main scope, we briefly outline what could be shown for them. Let Π be a tag-based AE whose encryption consists of a plain encryption $\Pi.\tilde{\mathcal{E}} : \mathcal{K} \times \mathcal{N} \times \mathcal{AD} \times \mathcal{M} \rightarrow \mathcal{M}$ and a tag generation $\Pi.\mathcal{T} : \mathcal{K} \times \mathcal{N} \times \mathcal{AD} \times \mathcal{M} \rightarrow \mathcal{M}$. We have $\Pi.\mathcal{E}(K, N, A, M) = (\tilde{C}, \tilde{T}) = (\Pi.\tilde{\mathcal{E}}(K, N, A, M), \Pi.\mathcal{T}(K, N, A, M))$. The inverse of $\Pi.\tilde{\mathcal{E}}$ is denoted by $\Pi.\tilde{\mathcal{D}} : \mathcal{K} \times \mathcal{N} \times \mathcal{AD} \times \mathcal{M} \rightarrow \mathcal{M}$, where $\Pi.\tilde{\mathcal{D}}(K, N, A, \Pi.\tilde{\mathcal{E}}(K, N, A, M)) = M$ holds for any valid (N, A, M) . See Fig. 12 for their algorithms. First, we can apply KS to CTX and CTY to obtain their standard-model-secure variants. As with (e.g.) AEaH-KS, we use RCR-PRG $G(K) \rightarrow (K_1, K_2)$ for KS, where K_1 is given to Π and K_2 is given to H . We assume H_{K_2} is a CR-PRF. Let $\text{CTX-KS}[\Pi, G, H]$ and $\text{CTY-KS}[\Pi, G, H]$ denote our fixed variants as shown by Fig. 12. We observe

$\text{CTX-KS.}\mathcal{E}[\Pi, G, H](K, N, A, M)$ 1. $(K_1, K_2) \leftarrow G(K)$ 2. $\tilde{C} \leftarrow \Pi.\tilde{\mathcal{E}}(K_1, N, A, M)$ 3. $\tilde{T} \leftarrow \Pi.\mathcal{T}(K_1, N, A, M)$ 4. $T \leftarrow H(K_2, N, A, \tilde{T})$ 5. return $(\tilde{C}, T)$	$\text{CTX-KS.}\mathcal{D}[\Pi, G, H](K, N, A, (\tilde{C}, T))$ 1. $(K_1, K_2) \leftarrow G(K)$ 2. $M \leftarrow \Pi.\tilde{\mathcal{D}}(K_1, N, A, \tilde{C})$ 3. $\tilde{T}^* \leftarrow \Pi.\mathcal{T}(K_1, N, A, M)$ 4. if $T \neq H(K_2, N, A, \tilde{T}^*)$: return $\perp$ 5. return M
$\text{CTY-KS.}\mathcal{E}[\Pi, G, H](K, N, A, M)$ 1. $(K_1, K_2) \leftarrow G(K)$ 2. $\tilde{C} \leftarrow \Pi.\tilde{\mathcal{E}}(K_1, N, \varepsilon, M)$ 3. $\tilde{T} \leftarrow \Pi.\mathcal{T}(K_1, N, \varepsilon, M)$ 4. $T \leftarrow H(K_2, N, A, \tilde{T})$ 5. return $(\tilde{C}, T)$	$\text{CTY-KS.}\mathcal{D}[\Pi, G, H](K, N, A, (\tilde{C}, T))$ 1. $(K_1, K_2) \leftarrow G(K)$ 2. $M \leftarrow \Pi.\tilde{\mathcal{D}}(K_1, N, \varepsilon, \tilde{C})$ 3. $\tilde{T}^* \leftarrow \Pi.\mathcal{T}(K_1, N, \varepsilon, M)$ 4. if $T \neq H(K_2, N, A, \tilde{T}^*)$: return $\perp$ 5. return M

Fig. 12: Standard-model secure versions of CTX and CTY transforms. (Top) CTX-KS, (Bottom) CTY-KS. The original $\text{CTX}[\Pi, H]$ and $\text{CTY}[\Pi, H]$ are obtained by omitting the line $[(K_1, K_2) \leftarrow G(K)]$ and using K instead of K_1 and K_2 .

that $\text{CTX}[\Pi, H]$ assuming H as an RO is equivalent to $\text{CTX}[\Pi, R_H]$ where R_H is a URF compatible with H if the adversary has no primitive queries to H inside $\text{CTX}[\Pi, H]$. Therefore, AE security bound for $\text{CTX-KS}[\Pi, G, H]$ will be obtained as a sum of PRG and PRF terms, and that obtained by the original bound of [CR22] with the number of primitive queries set to zero.¹⁰ For CMT security, the proof and the bound, the sum of RCR bound of G and CR bound of H , are mostly identical to those of AEtH-KS (Theorem 2). The same analysis applies to CTY: the AE security bound of $\text{CTY-KS}[\Pi, G, H]$ is obtained by the sum of PRG and PRF terms and the original bound of (RO-based) CTY [BH24] with omission of primitive queries. CMT security bound is similarly obtained as CTX-KS.

Dedicated analyses might yield better, finer AE security bounds, and there may be options beyond KS for achieving standard model security. In particular, proving/disproving the standard-model security of CTX/CTY without changing the algorithm, as we did for HtAE , remains open. They would be interesting future topics.

6 Conclusions

We have revisited state-of-the-art generic transforms that rely on a mixed security model and demonstrated practical attacks against them once instantiated. Considering the principal purpose of generic transforms, our findings warn against using them without applying adequate fixes, including ours. The idealized model is useful. However, our work suggests a (to our knowledge, new) type of pitfall

¹⁰ This observation in principle holds for CK25 transforms that we fixed; however, we provided their proofs from scratch as they are in our main scope, and for completeness and cleanness. Consequently, we simplified some proofs and improved some bounds relative to those obtained by a direct application of CK25 bounds.

for constructions based on the mixed security model. There might be more cases in other applications. Further investigation is an interesting direction.

Acknowledgements

We would like to thank Shan Chen and Vukašin Karadžić (CK25) and Arghya Bhattacharjee and Ritam Bhaumik and Chandranan Dhar (BBD24) for fruitful discussions and constructive comments.

References

- ADG⁺22. Ange Albertini, Thai Duong, Shay Gueron, Stefan Kölbl, Atul Luykx, and Sophie Schmieg. How to abuse and fix authenticated encryption without key commitment. In Kevin R. B. Butler and Kurt Thomas, editors, *USENIX Security 2022*, pages 3291–3308. USENIX Association, August 2022.
- BBD24. Arghya Bhattacharjee, Ritam Bhaumik, and Chandranan Dhar. Universal context commitment without ciphertext expansion. *IACR Cryptol. ePrint Arch.*, page 1382, 2024.
- BBP04. Mihir Bellare, Alexandra Boldyreva, and Adriana Palacio. An uninstantiable random-oracle-model scheme for a hybrid-encryption problem. In Christian Cachin and Jan Camenisch, editors, *EUROCRYPT 2004*, volume 3027 of *LNCS*, pages 171–188. Springer, Berlin, Heidelberg, May 2004.
- BCC⁺24. Ritam Bhaumik, Bishwajit Chakraborty, Wonseok Choi, Avijit Dutta, Jérôme Govinden, and Yaobin Shen. The committing security of MACs with applications to generic composition. In Leonid Reyzin and Douglas Stebila, editors, *CRYPTO 2024, Part IV*, volume 14923 of *LNCS*, pages 425–462. Springer, Cham, August 2024.
- BDPV08. Guido Bertoni, Joan Daemen, Michaël Peeters, and Gilles Van Assche. On the indistinguishability of the Sponge construction. In Nigel P. Smart, editor, *EUROCRYPT 2008*, volume 4965 of *LNCS*, pages 181–197. Springer, Berlin, Heidelberg, April 2008.
- BDPV12. Guido Bertoni, Joan Daemen, Michaël Peeters, and Gilles Van Assche. Duplexing the Sponge: Single-pass authenticated encryption and other applications. In Ali Miri and Serge Vaudenay, editors, *SAC 2011*, volume 7118 of *LNCS*, pages 320–337. Springer, Berlin, Heidelberg, August 2012.
- BFM15. Christina Brzuska, Pooya Farshim, and Arno Mittelbach. Random-oracle uninstantiability from indistinguishability obfuscation. In Yevgeniy Dodis and Jesper Buus Nielsen, editors, *TCC 2015, Part II*, volume 9015 of *LNCS*, pages 428–455. Springer, Berlin, Heidelberg, March 2015.
- BH22. Mihir Bellare and Viet Tung Hoang. Efficient schemes for committing authenticated encryption. In Orr Dunkelman and Stefan Dziembowski, editors, *EUROCRYPT 2022, Part II*, volume 13276 of *LNCS*, pages 845–875. Springer, Cham, May / June 2022.
- BH24. Mihir Bellare and Viet Tung Hoang. Succinctly-committing authenticated encryption. In Leonid Reyzin and Douglas Stebila, editors, *CRYPTO 2024, Part IV*, volume 14923 of *LNCS*, pages 305–339. Springer, Cham, August 2024.

- Bla06. John Black. The ideal-cipher model, revisited: An uninstantiable blockcipher-based hash function. In Matthew J. B. Robshaw, editor, *FSE 2006*, volume 4047 of *LNCS*, pages 328–340. Springer, Berlin, Heidelberg, March 2006.
- BN00. Mihir Bellare and Chanathip Namprempre. Authenticated encryption: Relations among notions and analysis of the generic composition paradigm. In Tatsuaki Okamoto, editor, *ASIACRYPT 2000*, volume 1976 of *LNCS*, pages 531–545. Springer, Berlin, Heidelberg, December 2000.
- BR93. Mihir Bellare and Phillip Rogaway. Random oracles are practical: A paradigm for designing efficient protocols. In Dorothy E. Denning, Raymond Pyle, Ravi Ganesan, Ravi S. Sandhu, and Victoria Ashby, editors, *ACM CCS 93*, pages 62–73. ACM Press, November 1993.
- BR00. Mihir Bellare and Phillip Rogaway. Encode-then-encipher encryption: How to exploit nonces or redundancy in plaintexts for efficient cryptography. In Tatsuaki Okamoto, editor, *ASIACRYPT 2000*, volume 1976 of *LNCS*, pages 317–330. Springer, Berlin, Heidelberg, December 2000.
- CGH04. Ran Canetti, Oded Goldreich, and Shai Halevi. The random oracle methodology, revisited. *J. ACM*, 51(4):557–594, 2004.
- CHB21. Paul Crowley, Nathan Huckleberry, and Eric Biggers. Length-preserving encryption with HCTR2. *IACR Cryptol. ePrint Arch.*, page 1441, 2021.
- CK25. Shan Chen and Vukašin Karadžić. Committing authenticated encryption: Generic transforms with hash functions. In Serge Fehr and Pierre-Alain Fouque, editors, *EUROCRYPT 2025, Part I*, volume 15601 of *LNCS*, pages 93–123. Springer, Cham, May 2025.
- CR22. John Chan and Phillip Rogaway. On committing authenticated-encryption. In Vijayalakshmi Atluri, Roberto Di Pietro, Christian Damsgaard Jensen, and Weizhi Meng, editors, *ESORICS 2022, Part II*, volume 13555 of *LNCS*, pages 275–294. Springer, Cham, September 2022.
- DFG23. Jean Paul Degabriele, Marc Fischlin, and Jérôme Govinden. The indistinguishability of the duplex and its practical applications. In Jian Guo and Ron Steinfeld, editors, *ASIACRYPT 2023, Part VIII*, volume 14445 of *LNCS*, pages 237–269. Springer, Singapore, December 2023.
- DGRW18. Yevgeniy Dodis, Paul Grubbs, Thomas Ristenpart, and Joanne Woodage. Fast message franking: From invisible salamanders to encryptment. In Hovav Shacham and Alexandra Boldyreva, editors, *CRYPTO 2018, Part I*, volume 10991 of *LNCS*, pages 155–186. Springer, Cham, August 2018.
- DHM⁺25. Joan Daemen, Seth Hoffert, Silvia Mella, Gilles Van Assche, and Ronny Van Keer. Shaking up authenticated encryption. In *EuroS&P*, pages 861–882. IEEE, 2025.
- Dwo07a. Morris Dworkin. *Recommendation for Block Cipher Modes of Operation: Galois/Counter Mode (GCM) and GMAC*. NIST-SP 800-38D, 2007.
- Dwo07b. Morris Dworkin. *Recommendation for Block Cipher Modes of Operation: the CCM Mode for Authentication and Confidentiality*. NIST-SP 800-38C, 2007.
- Dwo15. Morris J. Dworkin. Sha-3 standard: Permutation-based hash and extendable-output functions:, 2015-07-01 04:07:00 2015.
- FIM20. Rintaro Fujita, Takanori Isobe, and Kazuhiko Minematsu. ACE in chains: How risky is CBC encryption of binary executable files? In Mauro Conti, Jianying Zhou, Emiliano Casalicchio, and Angelo Spognardi, editors, *ACNS 2020, Part I*, volume 12146 of *LNCS*, pages 187–207. Springer, Cham, October 2020.

- FOR17. Pooya Farshim, Claudio Orlandi, and Răzvan Roşie. Security of symmetric primitives under incorrect usage of keys. *IACR Trans. Symm. Cryptol.*, 2017(1):449–473, 2017.
- GLR17. Paul Grubbs, Jiahui Lu, and Thomas Ristenpart. Message franking via committing authenticated encryption. In Jonathan Katz and Hovav Shacham, editors, *CRYPTO 2017, Part III*, volume 10403 of *LNCS*, pages 66–97. Springer, Cham, August 2017.
- HKR01. Helena Handschuh, Lars R. Knudsen, and Matthew J. B. Robshaw. Analysis of SHA-1 in encryption mode. In David Naccache, editor, *CT-RSA 2001*, volume 2020 of *LNCS*, pages 70–83. Springer, Berlin, Heidelberg, April 2001.
- HM24a. Shoichi Hirose and Kazuhiko Minematsu. Compactly committing authenticated encryption using encryptment and tweakable block cipher. In Claude Carlet, Kalikinkar Mandal, and Vincent Rijmen, editors, *SAC 2023*, volume 14201 of *LNCS*, pages 233–252. Springer, Cham, August 2024.
- HM24b. Viet Tung Hoang and Sanketh Menda. Robust AE with committing security. In Kai-Min Chung and Yu Sasaki, editors, *ASIACRYPT 2024, Part IX*, volume 15492 of *LNCS*, pages 312–342. Springer, Singapore, December 2024.
- HNS16. Shoichi Hirose, Yusuke Naito, and Takeshi Sugawara. Output masking of tweakable Even-Mansour can be eliminated for message authentication code. In Roberto Avanzi and Howard M. Heys, editors, *SAC 2016*, volume 10532 of *LNCS*, pages 341–359. Springer, Cham, August 2016.
- IIM21. Takanori Isobe, Ryoma Ito, and Kazuhiko Minematsu. Security analysis of SFrame. In Elisa Bertino, Haya Shulman, and Michael Waidner, editors, *ESORICS 2021, Part II*, volume 12973 of *LNCS*, pages 127–146. Springer, Cham, October 2021.
- IOM12. Tetsu Iwata, Keisuke Ohashi, and Kazuhiko Minematsu. Breaking and repairing GCM security proofs. In Reihaneh Safavi-Naini and Ran Canetti, editors, *CRYPTO 2012*, volume 7417 of *LNCS*, pages 31–49. Springer, Berlin, Heidelberg, August 2012.
- Jou06. Antoine Joux. Authentication Failures in NIST version of GCM. Public Comments to NIST, 2006. <http://csrc.nist.gov/groups/ST/toolkit/BKM/comments.html>.
- KBC97. Hugo Krawczyk, Mihir Bellare, and Ran Canetti. HMAC: keyed-hashing for message authentication. *RFC*, 2104:1–11, 1997.
- KR11. Ted Krovetz and Phillip Rogaway. The software performance of authenticated-encryption modes. In Antoine Joux, editor, *FSE 2011*, volume 6733 of *LNCS*, pages 306–327. Springer, Berlin, Heidelberg, February 2011.
- KRS25. Dmitry Khovratovich, Ron D. Rothblum, and Lev Soukhanov. How to prove false statements: Practical attacks on fiat-shamir. In Yael Tauman Kalai and Seny F. Kamara, editors, *CRYPTO 2025, Part VI*, volume 16005 of *LNCS*, pages 3–26. Springer, Cham, August 2025.
- KSW24. Juliane Krämer, Patrick Struck, and Maximiliane Weishäupl. Committing AE from Sponges security analysis of the NIST LWC finalists. *IACR Trans. Symm. Cryptol.*, 2024(4):191–248, 2024.
- LGR21. Julia Len, Paul Grubbs, and Thomas Ristenpart. Partitioning oracle attacks. In Michael Bailey and Rachel Greenstadt, editors, *USENIX Security 2021*, pages 195–212. USENIX Association, August 2021.
- MLGR23. Sanketh Menda, Julia Len, Paul Grubbs, and Thomas Ristenpart. Context discovery and commitment attacks - how to break CCM, EAX, SIV, and more. In Carmit Hazay and Martijn Stam, editors, *EUROCRYPT 2023*,

- Part IV*, volume 14007 of *LNCS*, pages 379–407. Springer, Cham, April 2023.
- NIS. Nist lightweight cryptography. <https://csrc.nist.gov/projects/lightweight-cryptography>.
- NRS14. Chanathip Namprempre, Phillip Rogaway, and Thomas Shrimpton. Reconsidering generic composition. In Phong Q. Nguyen and Elisabeth Oswald, editors, *EUROCRYPT 2014*, volume 8441 of *LNCS*, pages 257–274. Springer, Berlin, Heidelberg, May 2014.
- NSS23. Yusuke Naito, Yu Sasaki, and Takeshi Sugawara. Committing security of Ascon: Cryptanalysis on primitive and proof on mode. *IACR Trans. Symm. Cryptol.*, 2023(4):420–451, 2023.
- NSS24. Yusuke Naito, Yu Sasaki, and Takeshi Sugawara. KIVR: Committing authenticated encryption using redundancy and application to GCM, CCM, and more. In Christina Pöpper and Lejla Batina, editors, *ACNS 2024, Part I*, volume 14583 of *LNCS*, pages 318–347. Springer, Cham, March 2024.
- NSS25. Yusuke Naito, Yu Sasaki, and Takeshi Sugawara. Committing wide encryption mode with minimum ciphertext expansion. *IACR Trans. Symmetric Cryptol.*, 2025(1):44–69, 2025.
- PRZ99. Mohammad Peyravian, Allen Roginsky, and Nevenko Zunic. Hash-based encryption system. *Comput. Secur.*, 18(4):345–350, 1999.
- Rog02. Phillip Rogaway. Authenticated-encryption with associated-data. In Vijayalakshmi Atluri, editor, *ACM CCS 2002*, pages 98–107. ACM Press, November 2002.
- Rog04. Phillip Rogaway. Nonce-based symmetric encryption. In Bimal K. Roy and Willi Meier, editors, *FSE 2004*, volume 3017 of *LNCS*, pages 348–359. Springer, Berlin, Heidelberg, February 2004.
- RS06. Phillip Rogaway and Thomas Shrimpton. A provable-security treatment of the key-wrap problem. In Serge Vaudenay, editor, *EUROCRYPT 2006*, volume 4004 of *LNCS*, pages 373–390. Springer, Berlin, Heidelberg, May / June 2006.
- RVV15. Reza Reyhanitabar, Serge Vaudenay, and Damian Vizár. Boosting OMD for almost free authentication of associated data. In Gregor Leander, editor, *FSE 2015*, volume 9054 of *LNCS*, pages 411–427. Springer, Berlin, Heidelberg, March 2015.
- Sar10. Palash Sarkar. A simple and generic construction of authenticated encryption with associated data. *ACM Trans. Inf. Syst. Secur.*, 13(4):33:1–33:16, 2010.
- STEA. Construct an encryption algorithm using hash function. Crypto.StackExchange. <https://crypto.stackexchange.com/questions/15188/construct-an-encryption-algorithm-using-hash-function>.
- STEB. Is SHA-256 secure as a CTR block cipher? Crypto.StackExchange. <https://crypto.stackexchange.com/questions/1656/is-sha-256-secure-as-a-ctr-block-cipher>.
- STEC. Is it feasible to build a stream cipher from a cryptographic hash function? Crypto.StackExchange. <https://crypto.stackexchange.com/questions/48/is-it-feasible-to-build-a-stream-cipher-from-a-cryptographic-hash-function>.
- STED. Can I use HMAC-SHA1 in counter mode to make a stream cipher? Crypto.StackExchange. <https://crypto.stackexchange.com/questions/12810/can-i-use-hmac-sha1-in-counter-mode-to-make-a-stream-cipher>.

- SW24. Patrick Struck and Maximiliane Weishäupl. Constructing committing and leakage-resilient authenticated encryption. *IACR Trans. Symm. Cryptol.*, 2024(1):497–528, 2024.
- TMK⁺25. Meltem Sönmez Turan, Kerry McKay, Jinkeon Kang, John Kelsey, and Donghoon Chang. *Ascon-Based Lightweight Cryptography Standards for Constrained Devices: Authenticated Encryption, Hash, and Extendable Output Functions*. NIST SP 800-232, 2025.
- WFW05. Peng Wang, Dengguo Feng, and Wenling Wu. HCTR: A variable-input-length enciphering mode. In *CISC*, volume 3822 of *Lecture Notes in Computer Science*, pages 175–188. Springer, 2005.

A Proofs of Theorems 5 and 6 (EtH-KS)

A.1 Proof of Theorem 5

As with other proofs, we skip the hybrids regarding the replacement of $\mathbf{G}$ and $\mathbf{H}$. Let Θ_1 denote $\text{EtH-KS}[\Gamma, \$_{\mathbf{G}}, \mathbf{R}_{\mathbf{H}}]$. We have

$$\begin{aligned} \text{Adv}_{\text{EtH-KS}[\Gamma, \mathbf{G}, \mathbf{H}]}^{\text{priv}}(\mathcal{A}_{\text{priv}}) &\leq \text{Adv}_{\text{EtH-KS}[\Gamma, \mathbf{G}, \mathbf{H}].\mathcal{E}, \Theta_1.\mathcal{E}}^{\text{ind}}(\mathcal{A}_{\text{priv}}) + \text{Adv}_{\Theta_1.\mathcal{E}, \$}^{\text{ind}}(\mathcal{A}_{\text{priv}}) \\ &\leq \text{Adv}_{\mathbf{G}}^{\text{prg}}(\mathcal{A}_{\mathbf{G}}) + \text{Adv}_{\mathbf{H}}^{\text{prf}}(\mathcal{A}_{\mathbf{H}}) + \text{Adv}_{\Gamma}^{\text{priv}}(\mathcal{A}_{\Gamma}). \end{aligned}$$

As in the case of the privacy bound proof of AEaH-KS, $\text{Adv}_{\Theta_1.\mathcal{E}, \$}^{\text{ind}}(\mathcal{A}_{\text{priv}}) \leq \text{Adv}_{\Gamma}^{\text{priv}}(\mathcal{A}_{\Gamma})$ holds as $\mathcal{A}_{\text{priv}}$ is nonce-respecting and the T part is always independently uniform and random for both worlds. For authenticity bound, we have

$$\begin{aligned} \text{Adv}_{\text{EtH-KS}[\Gamma, \mathbf{G}, \mathbf{H}]}^{\text{auth}}(\mathcal{A}_{\text{auth}}) &\leq \text{Adv}_{\text{EtH-KS}[\Gamma, \mathbf{G}, \mathbf{H}], \Theta_1}^{\text{ind}}(\mathcal{A}'_{\text{auth}}) + \text{Adv}_{\Theta_1}^{\text{auth}}(\mathcal{A}_{\text{auth}}) \\ &\leq \text{Adv}_{\mathbf{G}}^{\text{prg}}(\mathcal{A}'_{\mathbf{G}}) + \text{Adv}_{\mathbf{H}}^{\text{prf}}(\mathcal{A}'_{\mathbf{H}}) + \text{Adv}_{\Theta_1}^{\text{auth}}(\mathcal{A}_{\text{auth}}). \end{aligned}$$

We then have $\text{Adv}_{\Theta_1}^{\text{auth}}(\mathcal{A}_{\text{auth}}) \leq q_d/2^t$. This is because, for a decryption query (N', A', C', T') , the probability that $\mathbf{R}_{\mathbf{H}}(K_2, N', A', C') = T'$ holds is at most $1/2^t$, which means the probability that $\mathcal{A}_{\text{auth}}$ obtains some outputs other than $\perp$ from $\Theta_1.\mathcal{D}$ in the game is at most $q_d/2^t$. This concludes the proof.

A.2 Proof of Theorem 6

For a tuple $Z = ((K, N, A, M), (K', N', A', M'), (C, T))$, let $(K_1, K_2) = \mathbf{G}(K)$ and $(K'_1, K'_2) = \mathbf{G}(K')$. When $K = K'$, $(N, A) = (N', A')$ implies $M = M'$, thus Z is not valid (the attack never succeeds). If $K = K'$ and $(N, A) \neq (N', A')$ hold, it implies $\mathbf{H}$'s collision. When $K \neq K'$, if $K_2 \neq K'_2$ occurs it implies a collision of $\mathbf{H}$, and if $K_2 = K'_2$ occurs it implies right-half collision of $\mathbf{G}$. Thus, a valid Z either requires a collision on $\mathbf{H}$ or a right-half collision of $\mathbf{G}$. This concludes the proof.

B Proofs of Theorems 7 and 8 (chaSIV-NKS)

B.1 Proof of Theorem 7

The target scheme is $\text{chaSIV-NKS}[\Gamma, F, H]$. For proving the MR-PRIV bound, we first apply the hybrid argument to the original by replacing F with a URF R_F . This incurs additional $\text{Adv}_F^{\text{prf}}(\mathcal{A}_F)$ for $\mathcal{A}_F$ using $2Q_e$ queries. After the hybrid argument, the scheme is transformed into the composition of Γ and H as in the original chaSIV . The difference is that, for each query taking nonce N , the scheme first uniformly samples $(K_e^N, K_h^N) \leftarrow \mathcal{K}^2$ independently for each N , and Γ takes the (nonce-dependent) key K_e^N . The PRF H takes (A, M) as input with the key K_h^N . We define this scheme as $\text{chaSIV-NKS}'[\Gamma, H]$ and denote it by Θ_1 . Moreover, let Θ_2 denote $\text{chaSIV-NKS}'[\Pi, R_H]$ which computes $T = R_H(N, A, M)$ for a URF R_H . The rest is identical to Θ_1 ; Θ_2 ignores K_h^N and randomly generates K_e^N per N and uses it for Γ .

For the MR-PRIV bound, we have

$$\text{Adv}_{\text{chaSIV-NKS}[\Gamma, F, H]}^{\text{mr-priv}}(\mathcal{A}_{\text{priv}}) \leq \text{Adv}_{\text{chaSIV-NKS}'[\Gamma, H]}^{\text{mr-priv}}(\mathcal{A}_{\text{priv}}) + \text{Adv}_F^{\text{prf}}(\mathcal{A}_F), \quad (6)$$

$$\begin{aligned} \text{Adv}_{\text{chaSIV-NKS}'[\Gamma, H]}^{\text{mr-priv}}(\mathcal{A}_{\text{priv}}) &\leq \text{Adv}_{\Theta_1, \mathcal{E}, \Theta_2, \mathcal{E}}^{\text{ind}}(\mathcal{A}_{\text{priv}}) + \text{Adv}_{\Theta_2, \mathcal{E}, \$}^{\text{ind}}(\mathcal{A}_{\text{priv}}), \\ \text{Adv}_{\Theta_1, \mathcal{E}, \Theta_2, \mathcal{E}}^{\text{ind}}(\mathcal{A}_{\text{priv}}) &\leq Q_e \cdot \text{Adv}_H^{\text{prf}}(\mathcal{A}_H), \end{aligned} \quad (7)$$

$$\text{Adv}_{\Theta_2, \mathcal{E}, \$}^{\text{ind}}(\mathcal{A}_{\text{priv}}) \leq Q_e \cdot \text{Adv}_\Gamma^{\text{priv}}(\mathcal{A}_\Gamma) + \frac{\sum_{i \in [1..Q_e]} R_i^2}{2^r} \quad (8)$$

holds. Eq. (7) holds from a simple multi-instance analysis of PRF: there are Q_e instances for H each taking at most R_e^* queries in Θ_1 , while R_H takes N as a part of input. To see Eq. (8) holds, we observe that the T part is uniform for both worlds, thus the distinguishing game reduces to the multi-instance distinguishing game for Γ , where each instance is indexed by N , and there are Q_e instances. The bad event is a collision of $N_e = \text{msb}_r(T)$ to Γ under the same instance, *i.e.*, under the same K_e^N , in Θ_2 . Since T is completely random, the bad event occurs with probability at most $\sum_{i \in [1..Q_e]} R_i^2 / 2^r$. If this bad event does not occur, the game reduces to the standard multi-instance PRIV game for Γ under the condition that N_e s do not collide among the same instance. The advantage of this game is at most $Q_e \cdot \text{Adv}_\Gamma^{\text{priv}}(\mathcal{A}_\Gamma)$, as the PRIV game allows arbitrarily distinct choice of N_e . From Eqs. (6) to (8), we conclude the MR-PRIV proof.

For the MR-AUTH bound, we define Θ_1 and Θ_2 as in the MR-PRIV proof. Thus, we have

$$\begin{aligned} \text{Adv}_{\text{chaSIV-NKS}[\Gamma, F, H]}^{\text{mr-auth}}(\mathcal{A}_{\text{auth}}) &\leq \text{Adv}_{\text{chaSIV-NKS}'[\Gamma, H]}^{\text{mr-auth}}(\mathcal{A}_{\text{auth}}) + \text{Adv}_F^{\text{prf}}(\mathcal{A}'_F), \\ \text{Adv}_{\text{chaSIV-NKS}'[\Gamma, H]}^{\text{mr-auth}}(\mathcal{A}_{\text{auth}}) &\leq \text{Adv}_{\Theta_1, \Theta_2}^{\text{ind}}(\mathcal{A}'_{\text{auth}}) + \text{Adv}_{\Theta_2}^{\text{auth}}(\mathcal{A}_{\text{auth}}). \end{aligned} \quad (9)$$

The term $\text{Adv}_{\Theta_1, \Theta_2}^{\text{ind}}(\mathcal{A}'_{\text{auth}})$ in Eq. (9) is at most $Q_{ed} \cdot \text{Adv}_H^{\text{prf}}(\mathcal{A}'_H)$. For the last term in Eq. (9), we relax the game so that K_e^N is given to the adversary on any query involving N . The tidiness assumption on Γ ensures that (K_e^N, C) for a valid ciphertext C uniquely determines the corresponding plaintext M . Thanks to the

tidiness, the game effectively reduces to the deterministic MAC unforgeability game for $R_{\tilde{H}}$ on the message (N, A, M) ; note that N is part of the message, hence it can repeat. Since $R_{\tilde{H}}$ is a URF, the forgery probability (which is equal to the last term of Eq. (9)) is $q_d/2^t$. This proves MR-AUTH bound.

B.2 Proof of Theorem 8

The proof is mostly identical to the proof of Theorem 2. For a tuple $Z = ((K, N, A, M), (K', N', A', M'), (C, T))$, let $(K_e, K_h) = F(K, N)$ and $(K'_e, K'_h) = F(K', N')$. If $(N, K) = (N', K')$, we have $(K_e, K_h) = (K'_e, K'_h)$, which implies $M = M'$ and $A \neq A'$. This implies a collision on H . If $(N, K) \neq (N', K')$, the subcase $(K_h = K'_h)$ implies a right-half collision on F . Another subcase $(K_h \neq K'_h)$ implies a collision on H . Thus, a valid Z implies a collision on H or a right-half collision on G , which concludes the proof.

C Proofs of Theorems 9 and 10 (HtAE-ST)

C.1 Proof of Theorem 9

As with other proofs, let Θ_1 denote $\text{HtAE-ST}[\Pi, R_H]$. For PRIV bound, we have

$$\begin{aligned} \text{Adv}_{\text{HtAE-ST}[\Pi, H]}^{\text{priv}}(\mathcal{A}_{\text{priv}}) &\leq \text{Adv}_{\text{HtAE-ST}[\Pi, H].\mathcal{E}, \Theta_1.\mathcal{E}}^{\text{ind}}(\mathcal{A}_{\text{priv}}) + \text{Adv}_{\Theta_1.\mathcal{E}, \$}^{\text{ind}}(\mathcal{A}_{\text{priv}}) \\ &\leq \text{Adv}_H^{\text{prf}}(\mathcal{A}_H) + q_e \text{Adv}_{\Pi}^{\text{priv}}(\mathcal{A}_{\Pi}), \end{aligned}$$

where $\mathcal{A}_{\Pi}$ is a PRIV adversary against Π involving a randomly chosen key, and it performs one encryption query. $\text{Adv}_{\Theta_1.\mathcal{E}, \$}^{\text{ind}}(\mathcal{A}_{\text{priv}}) \leq q_e \text{Adv}_{\Pi}^{\text{priv}}(\mathcal{A}_{\Pi})$ holds as $\mathcal{A}_{\text{priv}}$ is nonce-respecting, and the T part is always independently uniform and random for both worlds; moreover, a secret key of Π is chosen uniformly at random every time $\mathcal{A}_{\text{priv}}$ performs an encryption query.

For authenticity, let $\Pi_{\perp} = (\Pi.\mathcal{E}, \perp)$ and $\Theta_2 = \text{HtAE-ST}[\Pi_{\perp}, R_H]$. We have

$$\begin{aligned} \text{Adv}_{\text{HtAE-ST}[\Pi, H]}^{\text{auth}}(\mathcal{A}_{\text{auth}}) &\leq \text{Adv}_H^{\text{prf}}(\mathcal{A}'_H) + \text{Adv}_{\Theta_1}^{\text{auth}}(\mathcal{A}_{\text{auth}}) \\ &\leq \text{Adv}_H^{\text{prf}}(\mathcal{A}'_H) + \text{Adv}_{\Theta_1, \Theta_2}^{\text{ind}}(\mathcal{A}'_{\text{auth}}) + \text{Adv}_{\Theta_2}^{\text{auth}}(\mathcal{A}_{\text{auth}}) \\ &\leq \text{Adv}_H^{\text{prf}}(\mathcal{A}'_H) + Q_d \cdot \text{Adv}_{\Pi}^{\text{auth}}(\mathcal{A}'_{\Pi}) + \frac{q_d}{2^t}, \end{aligned}$$

where $\mathcal{A}'_{\Pi}$ is an AUTH adversary against Π involving a randomly chosen key, and it performs at most one encryption query. In the above evaluations, we have $\text{Adv}_{\Theta_1, \Theta_2}^{\text{ind}}(\mathcal{A}'_{\text{auth}}) \leq Q_d \text{Adv}_{\Pi}^{\text{auth}}(\mathcal{A}'_{\Pi})$ similarly to the PRIV case; observe that there may be instances (indexed by N) taking no decryption queries, however, they do not contribute to the advantage by definition. Thus, only Q_d instances can contribute to the advantage. The reason why $\text{Adv}_{\Theta_2}^{\text{auth}}(\mathcal{A}_{\text{auth}}) := \text{Adv}_{\Theta_2, (\Theta_2.\mathcal{E}, \perp)}^{\text{ind}}(\mathcal{A}_{\text{auth}}) \leq q_d/2^t$ holds is as follows. We say the tuple (N', A', T') in a decryption query (N', A', C', T') is *new* if $\mathcal{A}_{\text{auth}}$ has not queried (N', A', M) to $\Theta_2.\mathcal{E}$ and received (C, T') from $\Theta_2.\mathcal{E}$ for some M and C . When a decryption

query (N', A', C', T') contains a new tuple (N', A', T') , the distinguishing probability between Θ_2 and $(\Theta_2.\mathcal{E}, \perp)$ is at most $1/2^t$ as R_H is a random function. When (N', A', T') is not new, there exists an encryption query and its response $(N', A', M, (C, T'))$ and $C' \neq C$ holds. Thus, the invoked decryption query to $\Pi_\perp.\mathcal{D}$, i.e. (N', A', C') , is not a forbidden query in the authenticity game against $\Pi_\perp$, and $\Pi_\perp.\mathcal{D}$ will return $\perp$ for that query. Therefore, the distinguishing probability between Θ_2 and $(\Theta_2.\mathcal{E}, \perp)$ is zero with such a query. From the above discussions, we have $\mathbf{Adv}_{\Theta_2}^{\text{auth}}(\mathcal{A}_{\text{auth}}) \leq q_d/2^t$.

C.2 Proof of Theorem 10

Suppose $Z = ((K, N, A, M), (K', N', A', M'), (C, T))$ be the final output of $\mathcal{A}_{\text{cmt}}$ with the corresponding tuple (C, T) . For any (K, N, A) , let $H^{\text{left}}(K, N, A) = K_e$ where $H(K, N, A) = (K_e, T)$. Let $K_e = H^{\text{left}}(K, N, A)$ and $K'_e = H^{\text{left}}(K', N', A')$. When $(K, N, A) = (K', N', A')$, we have $K_e = K'_e$ and $M = M'$ since

$$\Pi.\mathcal{D}(K_e, N, \varepsilon, C) = \Pi.\mathcal{D}(K'_e, N', \varepsilon, C)$$

holds, which means Z is not valid. When $(K, N, A) \neq (K', N', A')$, the probability that $H^{\text{right}}(K, N, A) = H^{\text{right}}(K', N', A')$ holds is equal to $\mathbf{Adv}_H^{\text{rcr}}(\mathcal{A}_H)$ for some $\mathcal{A}_H$ by definition. Therefore, we have

$$\mathbf{Adv}_{\text{HtAE-ST}[\Pi, H]}^{\text{cmt}}(\mathcal{A}_{\text{cmt}}) \leq \mathbf{Adv}_H^{\text{rcr}}(\mathcal{A}_H),$$

which concludes the proof.