

NI-DKG: Non-Interactive Distributed Key Generation Using Blockchain and Zero-Knowledge Proofs

Alex Kampa¹, Pau Escrich¹, Marta Bellés-Muñoz², and Roger Baig³

¹ Vocdoni Association, 6300 Zug, Switzerland

² Pompeu Fabra University, 08018 Barcelona, Spain

³ Polytechnic University of Catalonia, 08034 Barcelona, Spain

{alex,pau}@vocdoni.org, marta.belles@upf.edu, roger.baig@upc.edu

Abstract. We present a non-interactive DKG protocol that eliminates complaint procedures through the systematic use of ZK proofs. The protocol requires a timed public bulletin board with adjoined computing capacity as its coordination layer, a capability realized in practice by smart-contract-enabled blockchains. Existing smart-contract DKGs detect invalid contributions only after submission, requiring dispute phases that introduce timing constraints and attack surfaces. In our protocol, each participant submits a single zk-SNARK proving the correctness of their contribution: polynomial commitment consistency, Feldman verification equations, and correct share encryption. The smart contract rejects any invalid proof, simplifying the protocol to non-interactive phases delimited by block numbers. The protocol relies on standard primitives: Shamir’s secret sharing, Feldman commitments, hashed ElGamal encryption, and Chaum-Pedersen discrete-log equality proofs, with on-chain verification via zk-SNARKs. We provide an implementation outline for EVM-compatible chains, including circuit specifications for key generation, threshold decryption, and optional secret key disclosure. We also discuss EVM verifier constraints on the number of public inputs and sketch approaches to address them.

Keywords: Distributed key generation · Threshold cryptography · Verifiable secret sharing · Zero-knowledge proofs · Blockchain.

1 Introduction

Distributed key generation (DKG) enables a group of participants to jointly generate cryptographic keys without relying on a trusted dealer. This primitive underpins threshold signatures, randomness beacons, and privacy-preserving protocols. A compromised key generation phase can undermine all subsequent security guarantees, making the correctness and robustness of DKG a critical concern.

1.1 From DKG to ADKG

Classical DKG protocols, beginning with Pedersen’s 1991 construction [15], assume synchronous communication: messages arrive within a known bounded delay. This assumption can fail in practice, as real networks exhibit unbounded delays and partitions, and can be subject to adversarial scheduling.

Asynchronous DKG (ADKG) removes the synchrony assumption: the protocol must terminate and produce correct output regardless of message scheduling, provided that fewer than one-third of participants are Byzantine. The first ADKG construction was proposed by Kokoris-Kogias et al. [14] in 2020. Since then, progress has been rapid. Communication complexity has dropped from $O(\kappa n^4)$ to $O(\kappa n^3)$ [4, 8, 9, 23], and recent breakthroughs have achieved nearly-quadratic $\tilde{O}(n^2)$ communication [3, 11]. Adaptive security is now achievable with minimal or silent setup [2, 11]. A recent systematization of knowledge by Bacho and Kavousi [6] provides a comprehensive analysis of the DKG landscape.

However, a significant gap remains between theory and practice. The new constructions rely on sophisticated primitives such as packed asynchronous verifiable secret sharing (VSS), multi-valued validated Byzantine agreement and asynchronous common subset protocols. These primitives have yet to see widespread implementation or sustained cryptanalytic scrutiny.

1.2 Blockchain-assisted DKG

On the practical side, several works have explored using blockchains as a coordination layer for DKG. Schindler et al. [16] introduced EthDKG, implementing Joint-Feldman DKG on Ethereum. EthDKG requires interactive dispute resolution: a party that receives an invalid share posts a complaint on-chain, and the dealer must respond within a timeout. Sober et al. [18] improved on this by using zero-knowledge (ZK) non-interactive arguments of knowledge (zk-SNARKs) for dispute resolution, but the protocol still retains a complaint phase in which the ZK proof is generated reactively, in response to a dispute.

In the public bulletin board model, Applebaum and Pinkas [5] showed that per-party communication can be made independent of the number of DKG participants using low-density parity-check (LDPC) codes, though their construction uses techniques that are far from standard practice. Zhang et al. [24] explored decentralized ciphertext policy attribute-based encryption (CP-ABE) for single-round sharing with $O(n)$ reconstruction complexity.

The only major production deployment, Shutter Network [19], has operated on Gnosis chain since October 2022, using Feldman VSS with interactive complaints. While Shutter demonstrates the viability of on-chain DKG, it has known limitations: $O(n^2)$ communication, a synchrony assumption, and interactive dispute resolution that increases gas costs and protocol duration.

There remains a gap for protocols that are practical and rely on well-known, battle-tested cryptographic primitives, while still achieving public verifiability, eliminating interactive complaint procedures, and being compatible with Ethereum virtual machine (EVM)-based smart contract verification. The EVM is the execution environment of Ethereum and a wide range of compatible chains, includ-

ing Polygon, BNB Chain, and Avalanche. It is the most widely deployed smart contract platform, making EVM compatibility a practical deployment target.

1.3 Our Contribution

We present a blockchain-assisted DKG protocol for EVM-compatible chains whose central design principle is the *complete elimination of interactive complaint procedures* through systematic use of ZK proofs.

In all prior smart-contract DKG protocols [16, 18, 19], correctness of a participant’s contribution is verified after the fact: other participants detect invalid shares and file complaints, triggering an interactive dispute resolution phase. Even in the zk-SNARK-enhanced protocol of Sober et al. [18], ZK proofs are used reactively: a dealer generates a proof only when challenged. The complaint mechanism introduces an entire protocol phase with its own timing constraints, and opens a class of attacks that exploit the interactive structure: griefing (filing false complaints to waste gas), timing attacks (filing complaints near deadlines), and exclusion of honest participants whose dispute responses are delayed.

In our protocol, every participant proactively proves the correctness of their entire contribution in a single zk-SNARK submitted alongside the contribution data. The smart contract accepts a submission only if the proof is valid. There is no complaint phase because there is nothing to complain about: invalid contributions are rejected at submission time.

This has several consequences. First, the protocol structure simplifies: each participant’s actual DKG contribution obligation reduces to a single blockchain transaction, and phases are delimited by block numbers with no interactive back-and-forth. Second, the entire DKG transcript becomes publicly verifiable. Any observer can confirm correct execution by inspecting the on-chain record. Third, the attack surface associated with interactive dispute resolution is eliminated by construction, not by parameter tuning.

The protocol uses standard primitives: Shamir secret sharing, Feldman-style polynomial commitments [10], hashed ElGamal encryption, Chaum-Pedersen discrete-log equality proofs, and zk-SNARKs. We also discuss EVM verifier constraints on the number of public inputs and sketch approaches to address them. The protocol supports threshold decryption and optional secret key disclosure as post-DKG operations, with corresponding ZK proofs for each phase.

1.4 Motivating application

The protocol is motivated by Vocdoni’s DAVINCI end-to-end verifiable voting system [22, 20], in which voter ballots are encrypted under a threshold ElGamal public key generated via DKG. The corresponding secret key is never reconstructed during the election; decryption of the tally is performed collaboratively by the keyholding nodes only after polls close, preserving tally privacy throughout the voting period.

1.5 Asynchrony model

Our protocol uses the blockchain as a timed public bulletin board, inheriting its discrete global clock given by the block number, ordering, availability, and persistence guarantees. Phases are delimited by block numbers, and participants

have a known window to submit contributions. A participant who fails to submit is simply excluded. This is not fully asynchronous in the theoretical sense, and the protocol inherits the underlying chain’s partial synchrony⁴. However, it matches the operational model of blockchain-native applications and avoids the complexity of asynchronous agreement primitives.

Furthermore, the blockchain’s broadcast model is more resilient to targeted message suppression than point-to-point asynchronous networks, where an adversary can selectively delay or drop messages between specific pairs of participants. Submitting a transaction instead involves broadcasting to a wide network of nodes, making selective suppression significantly harder⁵.

1.6 Outline

Section 2 provides background on secret sharing, VSS, and DKG. Section 3 discusses node selection, while Section 4 presents the DKG protocol and its applications in threshold decryption and secret key disclosure. Section 5 revisits the ZK implementation, providing circuit specifications for each phase. Finally, Section 6 discusses EVM verifier constraints on public inputs and sketches approaches to address them.

2 Background

2.1 Secret Sharing

Secret sharing refers to methods for distributing a secret among a group of participants, each of whom is given a share of the secret. The secret can be reconstructed only when predetermined groups of participants combine their shares. Of particular importance is the case when there are n participants, and we require that any set of t or more shares can be used to reconstruct the secret, while any set of fewer shares reveals nothing about the secret. This is called a (t, n) *threshold scheme*.

The fundamental construction is due to Shamir [17]. A trusted dealer holds a secret $\sigma \in \mathbb{F}_q$ and generates a random polynomial $f(x) \in \mathbb{F}_q[X]$ of degree $t - 1$ such that $f(0) = \sigma$. The shares $s(j) := f(j)$ for $j \in [n]$ are then privately distributed to n parties. Any set of t parties can recover the secret by combining their shares using Lagrange interpolation. However, $t - 1$ (or fewer) parties can learn nothing about σ . Shamir secret sharing is used, in some form, by almost all practical secret sharing and DKG protocols.

⁴ Fully asynchronous consensus has been achieved by protocols such as HoneyBadgerBFT and BEAT, but all major smart-contract-enabled production blockchains assume partial synchrony, including Ethereum, Solana, Cosmos SDK chains, Sui, and Aptos. Sui comes closest to asynchrony, as its DAG mempool layer has async properties, but its total ordering layer still assumes partial synchrony.

⁵ As to mempool-level censorship, it would require controlling a large number of blockchain nodes throughout the entire submission window. This capability would threaten consensus integrity itself and falls outside our threat model.

2.2 Verifiable Secret Sharing (VSS)

With Shamir secret sharing, a faulty dealer could distribute spurious or random shares to participants. In that case, different subsets of t participants could reconstruct different secrets, rendering the protocol useless. Verifiable secret sharing (VSS) overcomes this limitation by allowing participants to verify that their shares are consistent.

Feldman [10] proposed the first non-interactive VSS protocol. The dealer publishes commitments to the polynomial coefficients:

$$C(k) := a(k) G \quad \text{for } k \in \{0, \dots, t-1\}$$

where G is a generator of a cyclic group $\mathcal{G}$ of prime order q in which the discrete logarithm problem is hard. Any participant j holding share $s(j) := f(j)$ can then verify its correctness by checking the Feldman verification equation:

$$s(j) G \stackrel{?}{=} \sum_{k=0}^{t-1} j^k C(k) \tag{1}$$

This equation is central to our protocol: proving that it holds for all participants is the main task of the ZK proof in the key generation phase.

Note that Feldman's scheme reveals $C(0) = \sigma G$, the public key corresponding to the secret. For our application this is acceptable and indeed required.

2.3 Distributed Key Generation (DKG)

A distributed key generation (DKG) protocol allows a group of n participants to jointly generate a shared public key and the corresponding secret shares of a private key, without any trusted dealer. Participants do not know the private key, but any subset of t or more of them is able to reconstruct it or to collaboratively perform cryptographic operations such as signing or decrypting messages.

The first DKG protocol was proposed by Pedersen [15]. It consists of the parallel invocation of n Feldman VSS instances: every participant generates a random polynomial of degree $t-1$, publishes commitments to its coefficients, and privately sends shares to all other participants. The shared public key can then be computed from the published commitments to the constant terms.

Pedersen's protocol assumes synchronous communication and uses an interactive complaint mechanism to handle faulty dealers. If a participant receives an invalid share, it broadcasts a complaint; the accused dealer must respond within a timeout or be disqualified. This interactive dispute resolution introduces timing assumptions and is the source of several practical complications, as discussed in the introduction.

3 Node Selection

The security of any DKG protocol ultimately depends on the assumption that fewer than some threshold of participants are malicious. The cryptographic protocol enforces this guarantee given an honest majority, but it cannot create one.

Node selection is a deployment concern that sits outside the protocol specification but is essential to its security in practice.

While the specific node selection strategy will depend on the application, collusion and coercion resistance will always be critical challenges. We outline several complementary approaches that can be combined to strengthen both collusion and coercion resistance.

Vetting. Node operators must undergo identity verification to prevent Sybil attacks.

Reputation. Choose well-known entities with significant reputations at stake, creating strong disincentives for malicious behavior.

Organization size. Large organizations are generally considered to be less coercion-resistant than smaller organizations or individuals.

Diversity. Nodes should be diverse both geographically and sector-wise (e.g., academic institutions, NGOs, private companies, independent operators). This provides resilience against regional internet outages, political or legal pressure, and reduces the likelihood of collusion.

Randomness. Nodes participating in a DKG round should be selected using verifiable random functions (VRFs) or blockchain-based randomness beacons to prevent adversaries from predicting or influencing selection.

Economic stake. Requiring nodes to lock collateral makes collusion economically costly and enables penalties for misbehavior through slashing mechanisms.

If we want a more diverse node ecosystem, we could define a two-tier structure: a core group of highly trusted nodes from established organizations, and a second tier of vetted ordinary nodes. For any DKG execution, we could require that a minimum fraction of participants come from the highly trusted tier. This will provide a baseline security guarantee while maintaining openness.

4 Our DKG and Decryption Protocol

In this section we describe the DKG protocol and its subsequent usage for threshold decryption and (optionally) secret key disclosure. The main output of the DKG is a public key PK whose corresponding secret key is not known to anyone. Diagrams summarizing these protocols can be found in Appendix 1 (DKG process), Appendix 2 (threshold decryption process) and Appendix 3 (secret key disclosure process).

4.1 Setting

Let $\mathcal{G}$ be a cyclic group of prime order q , written additively, with generator $G \in \mathcal{G}$. Scalars are elements of $\mathbb{F}_q$. We assume a smart-contract-enabled blockchain on which group operations in $\mathcal{G}$ (point addition and scalar multiplication) can be efficiently executed by smart contracts. Additionally, we assume that other operations, including the computation of hash function values, can also be performed

efficiently⁶. We further assume a set of eligible nodes is known, and for each eligible node i there exists a public key pub_i for an encryption scheme used to encrypt shares. We denote this encryption function by $\text{Enc}_{\text{share}}$. In the following, we assume that this encryption function is Hashed ElGamal, which is described in Appendix 4. Suitable smart contracts have been deployed, and the public keys of eligible nodes have been recorded.

4.2 Key Generation Process

Phase 1: Initiation

When initiating a DKG process, the organizer submits the protocol parameters (t, n) as well as policy parameters. Such policy parameters could include:

- Minimum participation criteria;
- Duration of phases 2 and 4;
- Minimum number of valid contributions in phase 4;
- Whether the secret key should remain secret “forever”, or whether it will need to be disclosed at some stage (e.g., either at a future time or block number, or via a request of the organizer).

The smart contract should ensure that parameters are within reasonable limits. For every DKG process, a unique pseudo-random round identifier rid is also generated. This identifier can be used as a domain separator.

Phase 2: Show of Hands

This phase lasts for a predetermined number of blocks. Before the DKG starts, the protocol determines which nodes are ready to participate. Eligible nodes monitor the blockchain and record their readiness to participate in round rid . This prevents the protocol from stalling due to offline nodes. Optionally, these candidate nodes could be required to deposit collateral.

Phase 3: Determine Participating Nodes (or Abort)

At the beginning of this phase we know the number and composition of the set of candidate nodes. If some predetermined policy criteria are not met, e.g. if there are not enough candidate nodes, the process aborts. Aborting is expected to be rare, as there should be some easy off-chain way to determine how many eligible nodes are available at a given time before initiating a DKG process. If the criteria are met, then a set of n selected nodes is determined. Some randomness available on the blockchain should be used to make this selection unpredictable.

While some blockchains support automatic execution at a specified block height via native scheduling mechanisms⁷, selecting participating nodes gener-

⁶ Chains such as Solana, Sui, and Aptos largely satisfy these assumptions through rich native cryptographic primitives and, on Sui and Aptos, native Poseidon hash support. EVM-compatible chains are less suitable: while precompiles exist for BN254 and BLS12-381, they are limited in scope and remain costly in practice, and there is no precompile for Poseidon.

⁷ Some Cosmos SDK-based chains allow block-height-triggered execution, by exposing Cosmos SDK’s block-based execution hooks as a user-facing scheduling primitive. Neutron is a prominent example.

ally requires submitting an explicit transaction. On certain chains, this can be delegated to decentralized automation networks⁸.

Phase 4: Main DKG Process

This phase lasts for a predetermined number of blocks. Each participant $i \in [n]$ generates a random polynomial in $\mathbb{F}_q[X]$ of degree $t - 1$:

$$f_i(x) := \sum_{k=0}^{t-1} a_{i,k} x^k,$$

where $a_i(k)$ are the coefficients randomly chosen from $\mathbb{F}_q$. Each participant also samples a set of random scalars to be used for encrypting messages for other participants:

$$\{r_{i,j}\} \quad \text{for } j \neq i, \text{ with } i, j \in [n]$$

Each participant then publishes the following data on the blockchain:

1. Public commitments of the polynomial coefficients:

$$C_i(k) := a_{i,k} G \quad \text{for } k \in \{0, \dots, t-1\}$$

2. Encrypted shares for each of the other participants. The share sent by participant i to participant j is $s_i(j) \triangleq f_i(j)$. Instead of sending the share via a private channel, an encrypted version of the share is published on-chain:

$$\begin{aligned} E_i(j) &:= \text{Enc}_{\text{share}}[\text{pub}_j, s_i(j)] \quad \text{for } j \neq i \\ &:= (R_i(j), \sigma_i(j)) \end{aligned}$$

3. ZK proof π_i for the following two statements:
 - For all $j \neq i$, the Feldman verification equation holds (see Equation (1)):

$$s_i(j) G \stackrel{?}{=} \sum_{k=0}^{t-1} j^k C_i(k)$$

This proves that the shares are consistent with the committed polynomial coefficients.

- For all $j \neq i$, $E_i(j)$ is a correct encryption of $s_i(j)$.

The smart contract will only accept a submission if π_i is correct. If the proof is not correct, the transaction fails and no data is written to the blockchain.

Commitments to polynomial coefficients must be stored on-chain because they will be used in the next stage. Encrypted shares and ZK proofs do not need to be stored as long as they remain accessible as calldata.

⁸ Widely used examples for EVM-compatible chains include Chainlink Automation and Gelato Network, which provide trust-minimized but not fully trustless transaction scheduling.

Phase 5: Compute Public Key and Shares of the Secret

At the beginning of this phase, the number of successful contributions is known. Let $\mathcal{I} \subseteq [n]$ be the set of indices of these nodes. The process aborts if an insufficient number of contributions have been submitted. Let's assume that a sufficient number of contributions were submitted. The shared secret polynomial $F(x)$ is then:

$$F(x) := \sum_{\ell \in \mathcal{I}} f_{\ell}(x)$$

We have:

$$F(x) = \sum_{\ell \in \mathcal{I}} a_{\ell,0} + x \sum_{\ell \in \mathcal{I}} a_{\ell,1} + x^2 \sum_{\ell \in \mathcal{I}} a_{\ell,2} + \dots + x^{t-1} \sum_{\ell \in \mathcal{I}} a_{\ell,t-1}$$

And therefore:

$$\begin{aligned} F(x)G &= \sum_{\ell \in \mathcal{I}} a_{\ell,0}G + x \sum_{\ell \in \mathcal{I}} a_{\ell,1}G + x^2 \sum_{\ell \in \mathcal{I}} a_{\ell,2}G + \dots + x^{t-1} \sum_{\ell \in \mathcal{I}} a_{\ell,t-1}G \\ &= \sum_{\ell \in \mathcal{I}} C_{\ell}(0) + x \sum_{\ell \in \mathcal{I}} C_{\ell}(1) + x^2 \sum_{\ell \in \mathcal{I}} C_{\ell}(2) + \dots + x^{t-1} \sum_{\ell \in \mathcal{I}} C_{\ell}(t-1) \end{aligned}$$

We define aggregated commitments as follows:

$$\overline{C}(k) := \sum_{\ell \in \mathcal{I}} C_{\ell}(k)$$

Then $F(x)G$ can be written as:

$$F(x)G = \sum_{k=0}^{t-1} x^k \overline{C}(k) \tag{2}$$

Privately compute shares of the secret key. Each of the participating nodes can compute their share of the secret as follows:

$$d_i := \sum_{\ell \in \mathcal{I}} s_{\ell}(i)$$

Recall that $s_{\ell}(i) = f_{\ell}(i)$, so that $d_i = F(i)$. Note that this computation is performed completely off-chain, and that these shares must be kept secret.

Publish the public key. $F(0)$ is the secret key and $F(0)G$ is the public key, which can be computed by the smart contract as follows:

$$\text{PK} = \overline{C}(0)$$

Publish commitments to secret shares. It will also be useful to publish the values of $D_i := d_i G = F(i) G$. These are essentially commitments to the secrets d_i , but they can be computed publicly. Using Equation (2) we have:

$$D_i = \sum_{k=0}^{t-1} i^k \overline{C}(k)$$

These commitments $\{D_i\}_{i \in \mathcal{I}}$ can be computed by the smart contract.

4.3 Threshold Decryption

An ElGamal ciphertext of message M under public key PK has the form:

$$\text{Enc}_{\text{msg}}[\text{PK}, M] = (C_1, C_2) = (rG, M + r\text{PK}),$$

and decryption with secret key sk satisfying $\text{PK} = \text{sk} G$ yields:

$$M = C_2 - \text{sk} C_1.$$

Here $\text{sk} = F(0)$ is not known to anyone. The goal is to compute $\Delta := \text{sk} C_1$ without disclosing sk .

Phase 1: Publication of Ciphertext

A ciphertext (C_1, C_2) is published together with the round identifier rid . The nodes that participated in that DKG process monitor the blockchain and know when such ciphertexts are published.

Phase 2: Publication of Partial Decryptions

Participating nodes publish a partial decryption of C_1 :

$$\delta_i := d_i C_1$$

together with a proof of discrete log equality:

$$(A_i, B_i, z_i)$$

Details of the Chaum-Pedersen proof. We assume that commitments D_i to the secret shares d_i have been published. Node i then needs to prove that

$$(Y_1, Y_2) := (D_i, \delta_i)$$

have the same discrete log with respect to G and C_1 . Node i samples a random r and computes:

$$\begin{aligned} A_i &= r G \\ B_i &= r C_1 \\ e &= \text{Hash}(\text{rid} \parallel Y_1 \parallel Y_2 \parallel A_i \parallel B_i) \\ z_i &= r + e d_i \end{aligned}$$

The proof is the tuple (A_i, B_i, z_i) , and its verification requires:

$$\begin{aligned} e' &:= \text{Hash}(\text{rid} \parallel Y_1 \parallel Y_2 \parallel A_i \parallel B_i) \\ z G &\stackrel{?}{=} A_i + e' Y_1 \\ z C_1 &\stackrel{?}{=} B_i + e' Y_2 \end{aligned}$$

This verification can be done directly by the smart contract.

Phase 3: Decryption

Once t or more partial decryptions have been published, the smart contract can be called to recover the message. Let $\mathcal{Q} = \{x_1, \dots, x_t\} \subseteq \mathcal{I}$ be a qualifying set of nodes that have published partial decryptions $\{\delta_{x_k}\}_{k \in [t]}$ and define $\mathcal{Q}_k$ as follows:

$$\mathcal{Q}_k = \{x_1, \dots, x_t\} \setminus \{x_k\}$$

We know that:

$$\text{sk} = F(0) = \sum_{k \in [t]} \lambda_k d_{x_k} \quad \text{where} \quad \lambda_k = \prod_{u \in \mathcal{Q}_k} \frac{u}{u - x_k}$$

We therefore have:

$$\Delta := \text{sk} C_1 = \sum_{k \in [t]} \lambda_k d_{x_k} C_1 = \sum_{k \in [t]} \lambda_k \delta_{x_k}$$

The decrypted message is then:

$$M = C_2 - \Delta$$

These computations can also be done directly by the smart contract.

4.4 Secret Key Disclosure

In some use cases, the secret key needs to be published. The recovery of the secret key can either be triggered by an external event monitored by the nodes or after a predetermined block number. Every participant knows the value of $F(x)$ at a single point. Because F is of degree $t - 1$, any subset of t participants is therefore able to reconstruct this polynomial.

Phase 1: Request Secret Key Disclosure

The request to disclose the secret key can be sent by the organizer or triggered automatically, assuming this is allowed by the policy parameters.

Phase 2: Publication of Secret Key Shares

Every participating node $i \in \mathcal{I}$ publishes their share of the secret d_i . The smart contract makes sure that the submitted value is correct by verifying:

$$d_i G \stackrel{?}{=} D_i$$

Phase 3: Computation of the Secret Key

Once t or more values have been submitted, anyone can compute the secret key $F(0)$. The smart contract should also be called to reconstruct the private key $F(0)$ in order to have a public record. Let $\mathcal{Q}' = \{x_1, \dots, x_t\} \subseteq \mathcal{I}$ be a qualifying set of nodes that have published their shares of the secret. We have, by Lagrange interpolation, and with λ_k defined as above:

$$F(0) = \sum_{k \in [t]} F(x_k) \lambda_k$$

The secret key can therefore be computed by the smart contract as:

$$\text{sk} = \sum_{k \in [t]} d_{x_k} \lambda_k$$

5 Implementation Outline for EVM chains

In Section 4, we assumed that smart contracts can perform group operations in $\mathcal{G}$. On many smart-contract platforms (e.g. EVM chains), this is not the case: the chain cannot efficiently add points or multiply by scalars in $\mathcal{G}$. It is efficient to offload all computations and checks involving group operations in $\mathcal{G}$ to zero-knowledge proofs. By employing zk-SNARKs over a curve supported by the execution environment (such as BN254 or BLS12-381 on EVM chains), the smart contract's role is reduced to: (i) bookkeeping, (ii) proof verification, and optionally (iii) hashing or other operations to compute commitments.

To this end, we revisit only the phases of Section 4 relevant to the ZK implementation, highlighting key architectural details and providing formal circuit specifications where applicable.

5.1 Data Model and Circuit Specifications

The smart contract treats group elements in $\mathcal{G}$ as opaque encodings, i.e., byte strings. The meaning of these encodings is enforced only through ZK proofs. We assume that the set of eligible nodes, along with their public encryption keys, is stored on-chain. We define the vector of these public keys as:

$$\mathbf{pub} \triangleq \{\text{pub}_j\}_{j \in [n]}$$

We further define the following vectors for each $i \in [n]$:

$$\begin{aligned} \mathbf{a}_i &\triangleq \{a_{i,k}\}_{k \in \{0, \dots, t-1\}} \\ \mathbf{C}_i &\triangleq \{C_i(k)\}_{k \in \{0, \dots, t-1\}} \\ \mathbf{s}_i &\triangleq \{s_i(j)\}_{j \in [n] \setminus \{i\}} \\ \mathbf{r}_i &\triangleq \{r_{i,j}\}_{j \in [n] \setminus \{i\}} \\ \mathbf{E}_i &\triangleq \{(R_i(j), \sigma_i(j))\}_{j \in [n] \setminus \{i\}} \end{aligned}$$

The vector of commitments to the values d_j will be denoted:

$$\mathbf{D} \triangleq \{D_j\}_{j \in \mathcal{I}}$$

Note also that for every interaction with the blockchain, nodes will need to identify themselves. A node can be identified by its blockchain address, in which case simply submitting a transaction from that address identifies the node. A more flexible alternative is to identify nodes by a signing public key, in which case a transaction can be submitted from any address, but must be accompanied by a valid signature.

5.2 ZK Proofs - DKG

Phase 1: Initiation

After a new DKG process is initiated, its parameters are stored on-chain. The three parameters that are relevant for ZK proofs are:

$$\mathbf{params} \triangleq (\text{rid}, n, t)$$

Phase 4: Main DKG Process

During this phase, every selected node i samples a random vector $\mathbf{a}_i$ and $\mathbf{r}_i$ and computes the vectors $\mathbf{C}_i$, $\mathbf{s}_i$ and $\mathbf{E}_i$. The computation makes use of **params** and **pub**, which are public information.

The node then computes the proof π_i with private inputs $\mathbf{a}_i$ and $\mathbf{r}_i$ and public inputs **params**, **pub**, $\mathbf{C}_i$ and $\mathbf{E}_i$. Note that the number of public inputs can be reduced by computing vector commitments. This will be discussed below, in Section 6.

The proof π_i together with $\mathbf{C}_i$ and $\mathbf{E}_i$ is sent to the smart contract, which stores $\mathbf{C}_i$ on-chain if the proof is correct. The following describes the circuit for which the proof is generated. Note that the public inputs that are sent as calldata are marked with an asterisk.

PRIVATE INPUTS

$\mathbf{a}_i$	(polynomial coefficients)
$\mathbf{r}_i$	(encryption randomness)

PUBLIC INPUTS

params	(protocol parameters)
pub	(public keys)
* $\mathbf{C}_i$	(polynomial commitments)
* $\mathbf{E}_i$	(encrypted shares)

IN-CIRCUIT COMPUTATIONS/VALIDATIONS

Compute shares for all $j \neq i$:

$$s_i(j) := \sum_{k=0}^{t-1} a_i(k) j^k \pmod{q}$$

Verify commitments to polynomial coefficients for all $k \in \{0, \dots, t-1\}$:

$$C_i(k) \stackrel{?}{=} a_i(k) G$$

Verify Feldman equations for all $j \neq i$:

$$s_i(j) G \stackrel{?}{=} \sum_{k=0}^{t-1} j^k C_i(k)$$

Verify knowledge of encryption randomness for all $j \neq i$:

$$R_i(j) \stackrel{?}{=} r_{i,j} G$$

Verify correct encryption for all $j \neq i$:

$$\sigma_i(j) \stackrel{?}{=} s_i(j) + \text{Hash}[r_{i,j} \text{ pub}_j] \pmod{q}$$

Phase 5: Compute Public Key and Shares of the Secret

Once the set $\mathcal{I}$ of participating nodes is known, the public key PK as well as commitments to private shares can be computed and published by anyone. However, to have a public reference, it is important to publish it on-chain as well. The corresponding ZK proof can be described as follows.

PRIVATE INPUTS

-

PUBLIC INPUTS

params	(protocol parameters)
$\mathcal{I} \subseteq [n]$	(indices of participating nodes)
$\{\mathbf{C}_i\}_{i \in \mathcal{I}}$	(all polynomial commitments)
* $\mathbf{D}$	(commitments to private shares)
* PK	(public key)

IN-CIRCUIT COMPUTATIONS/VALIDATIONS

Compute aggregated commitments for $k \in \{0, \dots, t-1\}$:

$$\overline{C}(k) := \sum_{\ell \in \mathcal{I}} C_\ell(k)$$

Verify Public Key:

$$\text{PK} \stackrel{?}{=} \overline{C}(0)$$

Verify commitments to private shares, for $i \in \mathcal{I}$:

$$D_i \stackrel{?}{=} \sum_{k=0}^{t-1} i^k \overline{C}(k)$$

5.3 ZK Proofs - Threshold Decryption

Phase 2: Publication of Partial Decryptions

In this phase, every node i provides a partial decryption δ_i of the ciphertext.

PRIVATE INPUTS

d_i	(private share)
-------	-----------------

PUBLIC INPUTS

params	(protocol parameters)
C_1	(first part of ciphertext)
D_i	(commitment to secret share)
$* \delta_i$	(partial decryption of C_1)
$* (A_i, B_i, z_i)$	(proof of DLEQ)

IN-CIRCUIT COMPUTATIONS/VALIDATIONS

The circuit computes:

$$e' := \text{Hash}(\text{rid} \parallel D_i \parallel \delta_i \parallel A_i \parallel B_i)$$

The two constraints are then:

$$\begin{aligned} z_i G &\stackrel{?}{=} A_i + e' D_i \\ z_i C_1 &\stackrel{?}{=} B_i + e' \delta_i \end{aligned}$$

Phase 3: Decryption

Once at least t partial decryptions have been published in phase 2, anyone can decrypt the ciphertext.

PRIVATE INPUTS

-

PUBLIC INPUTS

(C_1, C_2)	(ciphertext)
$\{\delta_i\}$	(partial decryptions)
$* \mathcal{Q}$	(qualifying set)
$* M$	(decrypted message)

IN-CIRCUIT COMPUTATIONS/VALIDATIONS

Compute Lagrange coefficients for $k \in [t]$:

$$\lambda_k := \prod_{u \in \mathcal{Q}_k} \frac{u}{u - x_k}$$

Verify decryption:

$$M \stackrel{?}{=} C_2 - \sum_{k \in [t]} \lambda_k \delta_{x_k} \pmod{q}$$

5.4 ZK Proof - Secret Key Disclosure**Phase 2: Publication of Secret Key Shares**

Every participating node $i \in \mathcal{I}$ computes a ZK proof for the following circuit.

PRIVATE INPUTS

-

PUBLIC INPUTS

$* d_i$	(secret key share)
D_i	(commitment to secret key)

IN-CIRCUIT COMPUTATIONS/VALIDATIONS

Verify that the secret key share is correct:

$$d_i G \stackrel{?}{=} D_i$$

Phase 3: Computation of the Secret Key

Any participant can compute a proof for the following circuit.

PRIVATE INPUTS

-

PUBLIC INPUTS

$* \text{sk}$	(secret key)
$* \mathcal{Q}$	(qualifying set)
$\{d_j\}_{j \in \mathcal{Q}}$	(shares of secret key)

IN-CIRCUIT COMPUTATIONS/VALIDATIONS

Compute Lagrange coefficients for $k \in [t]$:

$$\lambda_k := \prod_{u \in \mathcal{Q}_k} \frac{u}{u - x_k} \quad \text{where} \quad \mathcal{Q}_k = \mathcal{Q} \setminus \{x_k\}$$

Verify secret key reconstruction:

$$\text{sk} \stackrel{?}{=} \sum_{k \in [t]} d_{x_k} \lambda_k \pmod{q}$$

6 EVM Constraints and Public Input Scalability

In the preceding section, to make the protocol explicit, we took the naïve approach of sending all public inputs directly to the verifier. Many of these public inputs are actually elliptic curve points. For a BN254 verifier, these will be Baby Jubjub points, which are represented as two 32-byte scalars. If we were using a BLS12-381 verifier, the curve points would be on Jubjub and require four scalars. In addition, every encrypted share requires one curve point and one scalar.

From now on, let’s assume that the circuit is over BN254. The two critical phases are DKG phases 4 and 5 from Section 5.2, as they require the most public inputs. Assuming that $t \approx n/2$, an estimate of the number of inputs counted as 32-byte scalars is given in Table 1.

Table 1. Estimated number of public inputs (32-byte scalars)

Phase	Inputs n=10 n=20 n=40			
DKG phase 4 (per participant)	$\approx 6n$	≈ 120	≈ 240	≈ 480
DKG phase 5 (organizer)	$\approx n^2/2$	≈ 50	≈ 200	≈ 800

Now let’s look at two widely used verifiers: Groth16 [13] and FFLONK [12]. Their basic characteristics in terms of cost and public inputs are shown in Table 2. The data is taken from [21].

Table 2. BN254 Verifier Comparison

Property	Groth16 FFLONK	
Fixed gas cost	207,000	200,000
Per-input gas cost	6,650	400
+ calldata (sparse)	~ 140	
+ calldata (dense)	~ 512	
+ storage access (cold)	2,100	
Max. inputs	~ 700	~ 38

This shows that FFLONK, which is limited to about 38 public input scalars, cannot be used even if there are only $n = 10$ participants in the DKG. Groth16 could be used for up to $n \approx 35$ but at the cost of an overall very high gas expenditure: approximately 3.5m gas per participant in DKG phase 4, and 5m gas for the organizer in phase 5, for an overall total of almost 130m gas.

Several approaches exist to reduce public inputs substantially: hashing a vector of inputs to a single commitment via a zk-friendly function such as Posei-

don, at the cost of additional in-circuit constraints and on-chain computation; or Binding Random Linear Combinations (BRLC), which commit to a vector via a post-commitment random linear combination with minimal circuit overhead. A detailed analysis of these strategies is left to a future extended treatment of this work.

7 Conclusion

We have presented a blockchain-assisted DKG protocol that eliminates interactive complaint procedures entirely through systematic use of ZK proofs. Every participant proves the correctness of their contribution in a single zk-SNARK submitted alongside their data. The smart contract rejects invalid contributions at submission time, removing the need for dispute resolution phases and the class of attacks they enable.

The protocol is built from standard primitives and targets deployment on EVM-compatible chains. As discussed in Section 6, naïvely passing all public inputs to the verifier is impractical for committees of any realistic size, and public input reduction strategies will be necessary for practical deployment. A full analysis of these strategies is left to a future extended treatment of this work.

Limitations. The protocol inherits the timing model of the underlying blockchain, typically partial synchrony rather than full asynchrony. Participants who miss their submission window are excluded, which is appropriate for blockchain-native applications but falls short of fully asynchronous ADKG guarantees. The protocol does not provide adaptive security: the set of participants is fixed before the protocol begins, and security relies on a static corruption assumption. Finally, deployment on Layer 2 scaling networks, which offer significantly lower execution costs, may be more practical for larger values of n .

Future work. The most immediate next steps are concrete implementation and benchmarking for realistic parameter choices ($n \in \{10, 20, 40\}$), and formal security analysis characterizing the assumptions under which the ZK-ification preserves the security of the underlying Pedersen DKG.

The gap between sophisticated but undeployed ADKG protocols and simple but limited production systems remains wide. This work contributes a protocol that occupies a practical middle ground: not optimal in any single theoretical dimension, but deployable, auditable, and built on primitives whose security properties are well understood.

Appendix 1 DKG Process Diagram




Fig. 1. DKG protocol sequence diagram

Appendix 2 Threshold Decryption Process Diagram




Fig. 2. Threshold decryption sequence diagram

Appendix 3 Secret Key Disclosure Process Diagram




Fig. 3. Secret key disclosure sequence diagram.

Appendix 4 Hashed ElGamal Encryption

For the encryption of the shares, a straightforward choice is the standard Hashed ElGamal construction, which is the Diffie–Hellman KEM–DEM core underlying DHIES / ECIES [7, 1]. Since our plaintext space is $\mathbb{F}_q$, we instantiate the DEM as additive masking $c = m + \text{HashToField}(S) \bmod q$, where the hash output is reduced modulo q .

Participant i must encrypt the share $s_i(j) \in \mathbb{F}_q$ in a ZK-friendly manner. Let participant j have the BabyJubjub keypair:

$$\text{sk}_j \in \mathbb{F}_q, \quad \text{pub}_j = \text{sk}_j G$$

To encrypt a message $m \in \mathbb{F}_q$ for participant j , pick $r \in \mathbb{F}_q$ (with $r \neq 0$) and compute:

$$\begin{aligned} R &= r G && \text{(ephemeral public key)} \\ S &= r \text{pub}_j && \text{(ephemeral secret point)} \\ s &= \text{Hash}(S) && \text{(shared secret scalar)} \\ c &= m + s \pmod{q} && \text{(masked message scalar)} \end{aligned}$$

Note that **Hash** is a hash function that maps BabyJubjub elliptic curve points to scalars in $\mathbb{F}_q$. As is common in modern presentations, one may instead derive the pad as $s := \text{Hash}(R \parallel S)$ in order to bind the derived scalar explicitly to the ephemeral public key.

The ciphertext is then:

$$\text{Enc}_{\text{share}}[\text{pub}_j, s_i(j)] = (R, c)$$

The decryption works as follows:

$$\begin{aligned} S' &= \text{sk}_j R && \text{(recover secret point)} \\ s' &= \text{Hash}(S') && \text{(recover secret scalar)} \\ m &= c - s' \pmod{q} && \text{(unmask message scalar)} \end{aligned}$$

To verify correctness of (R, c) , the prover supplies a zero-knowledge proof of knowledge of r such that:

$$\begin{aligned} R &\stackrel{?}{=} r G \\ c &\stackrel{?}{=} m + \text{Hash}(r \text{pub}_j) \bmod (q) \end{aligned}$$

All operations remain inside BabyJubjub and $\mathbb{F}_q$.

Note that Hashed ElGamal is IND-CPA secure but malleable. In our protocol, the contract accepts only ciphertexts accompanied by valid ZK proofs of correct formation. This prevents adversaries from submitting malformed or adaptively manipulated ciphertexts. Within this restricted protocol model, the malleability of Hashed ElGamal does not yield practical attacks.

References

- [1] Michel Abdalla, Mihir Bellare, and Phillip Rogaway. *DHIES: An Encryption Scheme Based on the Diffie–Hellman Problem*. Tech. rep. Full version dated September 18, 2001. Cryptology ePrint / Full version on the authors’ webpages, Sept. 2001. URL: <https://web.cs.ucdavis.edu/~rogaway/papers/dhies.pdf>.
- [2] Ittai Abraham et al. “Bingo: Adaptivity and Asynchrony in Verifiable Secret Sharing and Distributed Key Generation”. In: *Advances in Cryptology – CRYPTO 2023*. Springer, 2023, pp. 39–70.
- [3] Ittai Abraham et al. *Nearly Quadratic Asynchronous Distributed Key Generation*. IACR ePrint 2025/006. 2025.
- [4] Ittai Abraham et al. “Reaching consensus for asynchronous distributed key generation”. In: *Proceedings of the 2021 ACM Symposium on Principles of Distributed Computing*. 2021, pp. 363–373.
- [5] Benny Applebaum and Benny Pinkas. “Distributing Keys and Random Secrets with Constant Complexity”. In: *Theory of Cryptography Conference (TCC 2024)*. Vol. 15366. LNCS. Springer, 2024, pp. 478–510.
- [6] Renas Bacho and Alireza Kavousi. “SoK: Dlog-based Distributed Key Generation”. In: *IEEE Symposium on Security and Privacy (S&P)*. 2025, pp. 614–632.
- [7] David Cash, Eike Kiltz, and Victor Shoup. “The Twin Diffie–Hellman Problem and Applications”. In: *Advances in Cryptology – EUROCRYPT 2008*. Vol. 4965. Lecture Notes in Computer Science. Springer, 2008, pp. 127–145. DOI: 10.1007/978-3-540-78967-3_8.
- [8] Sourav Das et al. “Practical asynchronous distributed key generation”. In: *2022 IEEE Symposium on Security and Privacy (SP)*. IEEE. 2022, pp. 2518–2534.
- [9] Sourav Das et al. “Practical asynchronous high-threshold distributed key generation and distributed polynomial sampling”. In: *32nd USENIX Security Symposium*. 2023, pp. 5359–5376.
- [10] Paul Feldman. “A practical scheme for non-interactive verifiable secret sharing”. In: *28th Annual Symposium on Foundations of Computer Science*. IEEE. 1987, pp. 427–438.
- [11] Hao Feng and Qiang Tang. “Asymptotically Optimal Adaptive Asynchronous Common Coin and DKG with Silent Setup”. In: *Advances in Cryptology – CRYPTO 2025*. Springer, 2025, pp. 647–680.
- [12] Ariel Gabizon and Zachary J. Williamson. *fflonk: a Fast-Fourier Inspired Verifier Efficient Version of PlonK*. Cryptology ePrint Archive, Report 2021/1167. 2021. URL: <https://eprint.iacr.org/2021/1167>.
- [13] Jens Groth. “On the Size of Pairing-Based Non-interactive Arguments”. In: *Advances in Cryptology – EUROCRYPT 2016*. Vol. 9666. Lecture Notes in Computer Science. Springer, 2016, pp. 305–326. DOI: 10.1007/978-3-662-49896-5_11.
- [14] Eleftherios Kogkris-Kogias, Dahlia Malkhi, and Alexander Spiegelman. “Asynchronous distributed key generation for computationally-secure ran-

- domness, consensus, and threshold signatures”. In: *Proceedings of the 2020 ACM SIGSAC Conference on Computer and Communications Security*. 2020, pp. 1751–1767.
- [15] Torben Pryds Pedersen. “A threshold cryptosystem without a trusted party”. In: *Advances in Cryptology—EUROCRYPT’91*. Springer, 1991, pp. 522–526.
 - [16] Philipp Schindler et al. *ETHDKG: Distributed Key Generation with Ethereum Smart Contracts*. Cryptology ePrint Archive, Report 2019/985. 2019.
 - [17] Adi Shamir. “How to Share a Secret”. In: *Communications of the ACM* 22.11 (1979), pp. 612–613. DOI: 10.1145/359168.359176.
 - [18] Michael Sober et al. “Distributed Key Generation with Smart Contracts using zk-SNARKs”. In: *Proceedings of the 38th ACM/SIGAPP Symposium on Applied Computing (SAC)*. Best Paper Award, Distributed Systems Track. ACM. 2023, pp. 229–238. DOI: 10.1145/3555776.3577677.
 - [19] Shutter Network Team. *Shutter: Distributed Key Generation with Applications to Front-Running Protection*. Cryptology ePrint Archive, Report 2024/1981. Deployed on Gnosis Chain since October 2022. 2024.
 - [20] Vocdoni. *DAVINCI - The universal voting protocol*. <https://davinci.vote/>. 2026.
 - [21] Vocdoni. “Gas Cost Analysis for ZK Proof Verification on Ethereum”. Internal technical report. 2026.
 - [22] Vocdoni. *Vocdoni - Blockchain Voting Technology*. <https://vocdoni.io/>. 2026.
 - [23] Haibin Zhang, Sisi Duan, et al. “Practical Asynchronous Distributed Key Generation: Improved Efficiency, Weaker Assumption, and Standard Model”. In: *IEEE/IFIP International Conference on Dependable Systems and Networks (DSN)*. 2023, pp. 568–581.
 - [24] Liang Zhang et al. “1-Round Distributed Key Generation with Efficient Reconstruction Using Decentralized CP-ABE”. In: *IEEE Transactions on Information Forensics and Security* 17 (2022), pp. 894–907. DOI: 10.1109/TIFS.2022.3152356.